

AXONA

Applications

*What runs on Axona today, what gets better if it
migrates, and what we are building next.*

David A. Smith
Axona.net

An Axona Applications Note v0.4 2026-06-01

The protocol's defaults are the product's features.

How to Read This

THIS DOCUMENT HAS TWO PARTS.

Part I — Existing Applications: A Deep Dive surveys the applications and services that would benefit from running on the Axona protocol. It leads with the highest-impact, lowest-effort case — the managed pub/sub services an application can replace with the protocol’s own primitive — then the decentralized protocols and centralized realtime services Axona can carry, and finally the established applications built on a DHT today. The Axona protocol is a live, browser-native network (@axona/protocol v2.10.0) with an authenticated handshake, signed and freshness-bounded messages, and a full pub/sub lifecycle; each entry maps an application’s need onto that surface.

Part II — Future Applications: Axona Powered covers the applications we are building *on* Axona rather than migrating *to* it: `civildefense.io` (the first end-user application), `SYZL` (the adaptive social feed), and a deliberately open pipeline for what comes next.

I Existing Applications: A Deep Dive

EACH ENTRY FOLLOWS the same shape: what the application or service does, where it falls short today, and what changes on the Axona protocol. We lead with the case that is both highest-impact and lowest-effort for the developer — replacing a managed pub/sub service — because it asks nothing of the application but a change of transport.¹

What the Axona Protocol Provides

The analysis throughout rests on a deployed substrate, not a plan. Four properties carry every case:

- **Browser-native and live.** Axona was designed for the WebRTC environment from the start. The reference network runs at `axona.net` via the `axona-peer` browser node and the `axona-bridge` signaling broker.²
- **Authenticated, not anonymous-by-accident.** Every peer link is established through the `axona/4` authenticated handshake, with the mesh channel-binding value bound to the DTLS certificate fingerprints — so a relaying bridge cannot transparently intercept “direct” peer traffic.

On the numbers. Latency figures in Part I are back-of-envelope projections from the Axona simulator’s 25,000-node benchmark, applied to each application’s typical workload — the simulator is the scalable testbed where protocol changes are measured before deployment. They are projections, not field measurements of the migrated application. What is *not* a projection is the protocol surface itself — the API verbs, the security properties, the lifecycle semantics — which is live and testable today.

¹ Streaming-class applications — those that become possible only with the Axona protocol’s latency budget but do not exist today — are the subject of Part II, not this part.

² The bridge carries WebRTC offer/answer payloads only and then drops out — it sees no application traffic, and any operator can run one. A federated bridge mesh is on the roadmap.

- **Signed and fresh.** Every published message is an Ed25519-signed envelope (format v2) over a domain-tagged core carrying a per-publisher monotonic sequence number and timestamp, with a freshness window enforced at ingress. Replay-to-fresh-subscribers is closed.
- **A real lifecycle, not just fan-out.** Beyond pub/sub, the kernel now exposes `unsub`, `pull` (by id *or* latest), `kill` (creator-only delete with tombstones), `unpub` (owner-only topic destroy), bounded per-topic queues with deterministic eviction, and a hold-time TTL.³ Membership-gated topics are supported via shared-encryption group keys at the application layer.

³ 24-hour default, 48-hour ceiling, sliding on pull.

These four properties recur throughout the catalogue as the reason each application benefits.

Pub/Sub-as-a-Service — the Managed Realtime Tier

THE HIGHEST-IMPACT CASE asks the least of the developer. A class of companies sells *pub/sub itself* as a service — their product is the WebSocket fan-out an application would otherwise build and operate, because doing it at scale (connection management, presence, history, global points of presence) is hard. The Axona protocol changes the economics directly: fan-out becomes a property of the mesh, the per-message meter disappears, and so does the operator who could read the traffic. An application replaces the service with the protocol's own pub/sub primitive — a change of transport, not of architecture. These services sort into four tiers by how cleanly Axona displaces them.

Tier 1 — Realtime client pub/sub (the bullseye). Vendors: Pusher Channels, Ably, PubNub, Firebase Realtime Database, Supabase Realtime, Azure Web PubSub / SignalR Service, Stream (getstream.io). A client subscribes to a channel, a publisher sends, every subscriber receives it in real time — plus presence, short-term history, and reconnect handling. The meter is per-message *and* per connection-minute *and* per channel-minute,⁴ or per monthly-active-user. This is the closest match in the entire document, because Axona's pub/sub primitive *is* this product — minus the operator and the meter.

⁴ Ably is roughly \$2.50 per million messages plus \$1 per million connection-minutes; PubNub starts near \$98/mo for 1,000 monthly-active users. The product is the fan-out infrastructure you would otherwise run yourself.

Managed-service feature	Axona equivalent
Channel / topic	<code>peer.sub(topic)</code> over <code>deriveTopicId</code> (public or owned)
Publish	<code>peer.pub</code> — Ed25519-signed, freshness-bounded envelope
Presence (“who’s online”)	<code>peer.peers</code> + <code>onPeerJoin</code> / <code>onPeerLeave</code> — native
Message history / “rewind”	replay cache + <code>peer.pull</code> (by id <i>or</i> latest) + hold-time TTL
Channel teardown / revoke	<code>unsub</code> / <code>kill</code> (creator delete + tombstone) / <code>unpub</code>
Private channels	Model 3 shared-encryption topic (group key at the app layer)
Reach analytics	<code>peer.metrics</code> (publishes / subscribers / <code>reshare_count</code>)

Tier-1 realtime pub/sub features and their Axona equivalents.

The pitch is economic and architectural at once: peers carry the fan-out, so marginal cost per message trends to zero and there is no per-connection bill; and because there is no central relay, Pusher/Ably/PubNub’s structural ability to read every channel’s traffic simply does not exist — messages are signed and, where wanted, end-to-end encrypted. Geographic locality (the S2 prefix) is a bonus none of these vendors offer natively.

Tier 2 — Push notifications (complement, not replace).

Vendors: Firebase Cloud Messaging (FCM), Apple Push Notification service (APNs), OneSignal, Airship, Pusher Beams. Axona delivers “came-online” and event pushes to *connected* peers natively, but it *cannot wake a closed application* — only the OS push channel can originate that, and only the platform vendor controls it. So Axona replaces the in-app realtime layer and *pairs with* FCM/APNs for the wake-the-device layer; it does not displace them.

Tier 3 — Durable server-side event streaming (NOT a replacement). Vendors: Apache Kafka / Confluent Cloud, Google Cloud Pub/Sub, AWS SNS + SQS / Kinesis / EventBridge, Azure Event Hubs / Service Bus, CloudAMQP (RabbitMQ), Solace Pub-Sub+, IBM MQ. These are durable commit logs and queues: long retention, partitioning, replay-from-offset, exactly-once or strict FIFO delivery, transactions, backend consumers. Axona is deliberately *not* this — its retention is a bounded queue (≤ 256 messages) under a 24–48 h hold-time TTL, its ordering is per-publisher rather than a global total order, and its delivery is high but best-effort under churn, not the exactly-once SLA Kafka and Ably market. Positioning Axona as a Kafka replacement would be the overclaim that discredits the rest of this document; it is a realtime fan-out layer, not an event-sourcing

backbone.

Tier 4 — MQTT / IoT brokers (partial). Vendors: HiveMQ Cloud, EMQX Cloud, Azure Event Grid (MQTT), AWS IoT Core. Conceptually adjacent — topics, retained messages ($\approx$ Axona’s replay-latest), QoS tiers — and Axona’s geo-locality is genuinely attractive for regional device fleets. But these brokers win on backend data-plane bridging (dozens of connectors into Kafka, Kinesis, databases) and device-constrained QoS that Axona does not provide. Axona fits *peer-to-peer* device coordination, not the industrial-broker role with enterprise integrations.

The honest gap list. To genuinely displace Ably/Pusher/PubNub rather than merely resemble them, four gaps must be stated plainly:

1. **No exactly-once / no global ordering SLA** — per-publisher sequence only; high delivery, not guaranteed delivery.
2. **Short, bounded history** — minutes-to-days, not the months of stored messages PubNub sells.
3. **No managed support / uptime SLA** — Axona is infrastructure you adopt, not a vendor you page at 3 a.m.
4. **Presence and fan-out at PaaS scale are unproven** — the claim needs the friction-realistic benchmark before it is field-credible.

The trade is explicit: you give up the SLA, the durability, and the managed support; you get no per-message fee, no operator able to read your traffic, and geographic locality for free. For the large class of realtime features that do not need a durability guarantee — chat, presence, live cursors, notifications, feeds — that is a favorable trade, and exactly the class Tier 1 vendors serve.

Decentralized Social & Messaging Protocols

A COHORT OF PROTOCOLS already run the right *shape* — signed messages fanned out through a relay layer — but pay for it with operator-run relays or homeservers that are the centralization and availability weak point. For these, Axona is not a rewrite: it is a better transport that drops the relay as a single point of failure. In each case Axona supplies the *fan-out and discovery layer*, not the protocol’s own identity or data model — those ride as opaque, app-signed payloads.

Nostr. “Notes and Other Stuff Transmitted by Relays” — clients publish secp256k1-signed events to a handful of independently-operated relays and read from several at once for coverage. Relay availability is fragile, popular relays face moderation and centralization pressure,

and a client must redundantly POST each event to N relays for reach. *What Axona changes:* a Nostr event (its own signature intact) becomes the opaque payload of an Axona publish, and relays become axonal-tree topics. One publish fans out through the mesh instead of N redundant uploads, relay operators stop being load-bearing, and geographic locality routes events toward the people near them. Axona replaces the *relay layer*, not Nostr’s **npub** identity.

Waku / WalletConnect. Waku (libp2p gossipsub, successor to Ethereum’s Whisper) carries privacy-preserving messages for the Status/Logos stack; WalletConnect relays end-to-end-encrypted session messages between a wallet and a dApp. Both still lean on a relay tier the user does not control. *What Axona changes:* encrypted session and messaging payloads map onto a Model 3 shared-encryption topic — the protocol carries opaque ciphertext and the relay operator disappears. WalletConnect’s “pair wallet and dApp, exchange encrypted requests” is a two-party encrypted topic, which is natively what Axona does.

Farcaster. A “sufficiently decentralized” social network: hubs store and propagate messages (casts, reactions) and gossip them via libp2p gossipsub, while identity and the username registry live on-chain. The hub-to-hub gossip layer is the operational weight. *What Axona changes:* Axona becomes the hub gossip transport — propagation rides axonal trees with bounded per-relay cost and geographic locality — while the on-chain registry stays exactly where it is. The decentralization story improves without touching the identity layer.

Matrix. Federated real-time communication: homeservers federate over the server-to-server API, and each room is replicated across the participating homeservers. A homeserver is a single point of failure and a load/centralization point for all of its users; the long-running P2P Matrix experiments exist precisely to remove it. *What Axona changes:* Axona offers the transport and discovery those experiments need — room events become a topic, handle/room discovery becomes a lookup, and a user’s reachability stops depending on one homeserver’s uptime.⁵

Bluesky / AT Protocol. Personal Data Servers (PDSes) hold user repos; a central **Relay** crawls them and emits a global “firehose” of repo commits that App Views index. The Relay is a massive centralized aggregation-and-fan-out point — the firehose is, structurally, one enormous pub/sub channel run by one party. *What Axona changes:* the firehose rebroadcast becomes an axonal-tree topic (or a set of them, partitioned by region or collection), and PDS discovery becomes a DHT lookup. This is the largest lift in the cohort, but also the clearest case of a centralized pub/sub firehose that Axona’s primitive is shaped to carry.

⁵ Durable room history stays an app-layer concern — a peer or service that retains it — since Axona’s hold-time is bounded.

Realtime Apps Currently Renting Centralized Infrastructure

THE PREVIOUS COHORT already thinks in pub/sub. This one does not — these applications buy a realtime layer (a CDN, a WebSocket PaaS, a game backend) because building one is hard. Axona’s axonal trees plus S2 geographic locality are the thing they would otherwise rent.

Peer-assisted CDN & P2P video. Peer5 (now part of Microsoft), Streamroot (Lumen), CDNBye, and THEO’s P2P module form a WebRTC mesh among viewers to share HLS/DASH segments, offloading the origin CDN and cutting bandwidth cost — especially for live events where everyone wants the same segment at the same instant. *What Axona changes:* a near-ideal fit — a live stream’s segment topic fans out through an axonal tree, and the S2 prefix clusters the swarm by region so peers pull from neighbors, not across oceans. Axona supplies the mesh formation, locality, and fan-out these products hand-build, with signed segments and no central coordinator tracking who watched what.

Collaborative editing & multiplayer. CRDT libraries (Yjs, Automerge) need a sync transport plus “awareness” (live cursors, presence); teams either run a `y-websocket` server or buy a multiplayer backend — Liveblocks, PartyKit, Cloudflare Durable Objects, Convex, Ably Spaces. *What Axona changes:* a document becomes a topic — CRDT update messages publish to it, and awareness/presence is native (`peer.peers` + join/leave). Because CRDTs converge from any order of opaque updates, Axona’s carry-opaque-bytes model fits cleanly, and per-publisher sequence numbers give each editor’s stream a stable order.⁶

Live data feeds. Sports scores, financial tickers, live blogs, election and flight trackers — one publisher, a high-frequency update stream, and a large read-only audience — served today by CDN-plus-WebSocket stacks or a Tier 1 pub/sub vendor. *What Axona changes:* the pub/sub sweet spot with nothing working against it — the data is inherently ephemeral (the TTL is a feature, not a limitation), the audience is read-only (no membership complexity), and fan-out to a geo-distributed crowd is exactly what axonal trees do. Of every category in this document, live feeds need the least from Axona to work well.

Partial fits. *Multiplayer game netcode & matchmaking* (Photon, Colyseus, Nakama, Hathora, Playroom): matchmaking discovery maps onto a DHT lookup and peer/relay state sync onto the mesh — good for non-authoritative or lockstep designs, though authoritative twitch shooters still want a server tick and the hardest cases are latency-

⁶ Durable document persistence remains app-layer — a peer or service that retains the doc — since Axona’s queue and hold-time are bounded. The honest boundary of the fit.

bound. *Feature-flag / config streaming* (LaunchDarkly, ConfigCat): pushing flag updates to clients is a small, clean pub/sub topic. *Software & game-patch distribution* (npm/apt/container registries, game patchers): DHT discovery plus swarm delivery is BitTorrent’s pattern applied to updates — strong on bandwidth economics, bounded by the need for an authoritative signed manifest the app supplies.

Where Axona does not fit. To keep the catalogue honest, three categories are explicit non-targets. *Distributed coordination and config stores* (etcd, ZooKeeper, Consul) require strong consensus (a CP system); Axona is not a consensus protocol and should not pretend to be. *Durable event-sourcing and data pipelines* are the Tier 3 streaming case already excluded. And *strong-consistency databases and search indices* need guarantees a best-effort fan-out mesh does not provide. Axona is a realtime discovery-and-fan-out layer; where the requirement is consensus or durable total order, the right answer is a different tool.

Applications Built on a DHT Today

THE REMAINING APPLICATIONS already run a DHT — for peer discovery, content addressing, or descriptor lookup — and would gain from swapping it for the Axona protocol’s locality-aware routing and pub/sub fan-out. These are higher-effort migrations than the pub/sub services above — they touch an application’s existing overlay — but the latency and architecture wins are substantial.

BitTorrent (Mainline DHT)

What it does. Mainline DHT is BitTorrent’s trackerless peer-discovery layer: a torrent’s `info_hash` maps into the keypace, swarm peers register under that key, and new peers query the DHT to find the swarm.

Current implementation. Kademlia, 160-bit IDs, $K=8$ buckets, $\alpha=3$ parallel queries.⁷

Performance & shortfall. Cold-start `get_peers` latency is typically 5–30 s, dominated by the iterative α -parallel walk through globally-scattered random-ID peers. Modern clients hide this behind a centralized tracker for the first discovery round; the DHT is the tracker-resistance *backup*, not the primary path — precisely because the cold-start latency is unacceptable for a primary path.

What Axona changes. Mainline lookups translate directly (same operation, same data model). Projected lookup latency drops to ~250–300 ms at 25K nodes — a 20–100× improvement that collapses the user-visible delay below the threshold where tracker fallback adds value. The qualitative shift: the DHT becomes *the* discovery layer, not a fallback.

⁷ The largest production DHT in the world by node count — an estimated 10–20 million nodes at peak.

IPFS / Filecoin / libp2p

What it does. Content-addressable storage: content is referenced by a CID, and the libp2p Kad-DHT routes lookups to peers that announced they hold it.

Current implementation. libp2p Kad-DHT (Kademlia variant), default $K=20$, $\alpha=3$, tens of thousands of nodes.

Performance & shortfall. CID lookup averages 1–5 s; provider-record retrieval adds 1–3 s; end-to-end first-byte from cold cache is 5–15 s — 25–300× slower than a centralized HTTP CDN (50–200 ms). The IPFS gateway tier exists to paper over this with centralized servers. Separately, `libp2p-pubsub` (gossipsub) for IPNS updates carries its own mesh-scaling challenges.

What Axona changes. DHT-layer lookups drop to ~250–300 ms (projection), making end-to-end first-byte ~500 ms–1 s — a 10–20× improvement that turns CDN-replacement viability from hypothetical into a real question. Axona’s axonal pub/sub trees would replace gossipsub for IPNS updates with bounded per-relay fan-out cost. The combination makes IPFS deployable on browser-class infrastructure for both halves, which it currently is not.

Ethereum (discv5)

What it does. discv5 is Ethereum’s peer-discovery layer — how a new node finds peers — using a Kademlia variant with verifiable-identity ENRs.

Performance & shortfall. Cold-start peer acquisition takes 30 s to several minutes (Kademlia walk plus per-peer ENR verification round-trips). For an on-demand join (a fresh validator, a new searcher) that is operational friction; mempool gossip via gossipsub adds a latency-sensitive scaling dimension the static Kademlia table cannot adapt to.

What Axona changes. Discovery latency drops 10–50× for the DHT layer (projection). The locality-aware bootstrap gives new nodes regionally-sensible initial peer sets — which matters more for Ethereum than BitTorrent because mempool propagation is latency-sensitive. Axona’s axonal trees are a natural fit for “broadcast a transaction to all interested validators in a region,” addressing the gossip-sub limitation directly.⁸

⁸ Axona’s authenticated handshake maps cleanly onto discv5’s identity requirements — the migration does not trade away ENR-style verifiable identity.

Bitcoin Block / Transaction Propagation

What it does. Bitcoin propagates blocks and transactions through a peer-to-peer mesh: DNS-seed bootstrap, `addr`-message peer discovery, gossip propagation.

Performance & shortfall. Full-block propagation averages 1–3 s (compact blocks reduce the announce to a few hundred ms plus a re-request); mempool transaction latency is 50–200 ms. The flooded gossip is *not* locality-aware — a US-east → Asia transaction may route through Europe — and MEV concentrates around peers with privileged low-latency access to miners.

What Axona changes. Locality-aware discovery solves regional clustering; nodes find nearby peers by default and transactions propagate along short-path edges. Compact-block propagation becomes routed pub/sub: every node subscribes to a block-announcement topic, the miner publishes, the axonal tree fans out at ~100 ms at 50K nodes (projection). The tree is built once and reused rather than re-running inventory exchange per peer per round, and peer-to-miner latency becomes a property of geography rather than peering relationships — the privileged-peer problem partially dissolves.

Tor Onion Services v3

What it does. Self-authenticating, location-hidden endpoints; descriptor lookup is mediated by the HSDir hash ring.

Performance & shortfall. Rendezvous takes 5–15 s end-to-end, dominated by Tor circuit setup; the descriptor lookup itself adds 1–3 s. HSDir nodes are infrastructure-tier (a centralization concession), and descriptor caching is short-lived — popular services see fetch storms.

What Axona changes. The descriptor lookup itself drops to ~300 ms (projection), though the Tor circuit still dominates total cost. The structural benefit is decentralizing the HSDir tier: Axona’s hop-caching deposits descriptors at any node along the lookup path, widening the descriptor-fetch tier beyond designated relays. The directory-authority trust model is unchanged — Axona widens the cache tier, it does not replace Tor’s anonymity design.

Storj / Sia (Decentralized Storage)

What it does. Files are sharded, encrypted, and replicated across provider nodes; the DHT mediates shard placement and retrieval.

Performance & shortfall. Storj shard retrieval is 200–800 ms for a 4 KB shard (dominated by the connection to a remote provider, not the lookup); provider discovery at edge is ~1 s. Placement is *region-blind* — a Berlin user may hold shards in São Paulo.

What Axona changes. Locality-aware placement: shards land on providers in the user’s S2 cell with high probability, dropping retrieval to the regional budget (~100–200 ms, projection). Repair coordination (replacing a shard when a provider drops) becomes an axonal-tree topic with high delivery under churn — and the new lifecycle verbs

give it explicit semantics: a provider going offline is a **kill**/tombstone event, not an inference from silence.

Tox / Briar (Peer-to-Peer Messengers)

What it does. Serverless peer-to-peer messengers. Tox uses a Kademlia DHT for contact discovery; Briar uses hybrid local + overlay routing.

Performance & shortfall. Tox contact-availability checks take 2–10 s; NAT makes initial establishment slower. *Group chat is expensive* — every message goes to every member, making large groups infeasible — and “contact came online” requires polling because there is no real pub/sub.

What Axona changes. Contact discovery drops to ~ 250 ms (projection), comparable to a centralized messenger’s offline-status check. Group chat becomes a pub/sub topic: bounded per-hop fan-out scaling to thousands of members; “came online” becomes an automatic event. Crucially, the membership and privacy story is now real, not aspirational: a private group is a **shared-encryption topic** (the group key is distributed at the app layer; the protocol carries opaque ciphertext and gates nothing it shouldn’t), and **kill/unsub/unpub** give first-class leave/remove/teardown. This is the first peer-to-peer messenger architecture that scales to large groups with a centralized service’s latency profile *and* a credible membership model.

YaCy (Decentralized Search)

What it does. A peer-to-peer search engine — each peer crawls part of the web; a custom DHT mediates shared indexing and query distribution.

Performance & shortfall. Queries take 2–30 s (index lookup plus query distribution, both $\sim \log N$ on a high ~ 200 –500 ms per-peer constant). Against Google’s ~ 200 ms median this is unusable interactively — the UX gap is why YaCy stayed a research curiosity.

What Axona changes. DHT-layer query distribution drops to ~ 250 ms (projection); shard-level result caching via hop caching makes an interactive YaCy-like engine plausible. Ranked aggregation across shards is a separate algorithmic problem, and beating Google needs decades of ranking infrastructure — but the *latency floor* stops being the disqualifying factor.

Radicle (Decentralized Git Forge)

What it does. Peer-to-peer git collaboration; repositories addressed by peer-ID, a DHT layer for repository discovery, custom routing for pull requests.

Performance & shortfall. Repository discovery is 1–3 s; change-set propagation is 5–30 s — making “find a project, browse it” feel slow versus GitHub and “watch this repo” poll-based rather than push-based.

What Axona changes. Repository discovery drops to ~300 ms (projection); change-set propagation becomes a `commits-on-this-repo` axonal topic with ~200 ms fan-out, making the forge interactive in the GitHub sense, not just browsable.⁹

⁹ The signed-envelope sequence numbers give watchers a tamper-evident, gap-detectable commit feed for free.

Handshake / ENS (Decentralized DNS)

What it does. Decentralized name resolution; ownership recorded on-chain, a DHT-like layer caches recent lookups.

Performance & shortfall. Hot cached lookups are <100 ms; cold lookups hit the chain and take seconds to minutes, so cache hit-rate dominates. The cache tier is operated by infrastructure peers (a centralization tradeoff) because random-Kademlia caching does not scale.

What Axona changes. Hop caching makes the cache tier *implicit*: every routed lookup deposits its result at intermediate nodes, popular names accrue shortcuts naturally, and the explicit caching infrastructure becomes unnecessary. Cold-cache latency drops because the routing itself is faster.

WebTorrent / WebRTC P2P

What it does. BitTorrent for browsers over WebRTC data channels, bridged to Mainline DHT for peer discovery (most browsers cannot speak Mainline directly).

Performance & shortfall. Discovery is 5–30 s (Mainline latency plus bridge overhead). The *Mainline-DHT bridge is a centralization concession driven entirely by the protocol’s inability to run in browsers* — the bridge operators are a single point of failure and surveillance.

What Axona changes. Axona *is* a browser-native DHT — the cleanest fit in the catalogue. WebTorrent on Axona eliminates the *discovery* bridge: peer discovery runs in-browser at the regional budget (100–300 ms), and the browser becomes a first-class peer-to-peer participant.¹⁰ The connection-cap analysis shows this is sustainable on browser-class infrastructure.

¹⁰ Axona still uses a lightweight *signaling* bridge for WebRTC NAT traversal, but it carries no application or discovery traffic and drops out once peers are connected — a categorically smaller role than the Mainline-DHT bridge it replaces.

Mastodon-Style Federations (No DHT Today)

What it does. A federation of independent ActivityPub servers; users follow accounts on other servers, and the home server fetches and serves federated content.

Current implementation. No DHT. Discovery via WebFinger (centralized per-domain handle resolution); cross-server delivery via direct HTTP push from the publisher’s home server to each subscriber’s home server.

Performance & shortfall. Home servers (often small VPS) show 200–2000 ms under load; popular-post federation can lag by minutes due to per-subscriber-server fan-out. The *home-server-as-bottleneck* pattern means a viral post on a small instance can take the instance down, because every other instance’s followers fetch from the publisher’s home.

What Axona changes. This is the application that becomes *architecturally different*. Handle resolution becomes an Axona lookup (independent of home-server load). Status delivery becomes an axonal-tree topic per follower-graph cluster: a viral post fans out via the tree at bounded per-relay cost, never saturating any single home server. Home servers keep durable storage and WebFinger origin authority, but the fan-out load moves to the DHT layer — making small servers stop being single points of failure for their users’ reach.

Summary of Improvements

Application	Current bottleneck	Axona improvement
BitTorrent (Mainline)	5–30 s cold lookup	~300 ms; primary path possible
IPFS / Filecoin	5–15 s first-byte	~1 s; CDN-replacement viable
Ethereum discv5	Multi-minute join, mempool latency	~1 s join, locality-aware mempool
Bitcoin propagation	Region-blind, MEV-prone	Locality-aware, axonal-tree announce
Tor onion services	DHT minor; HSDir centralization	Wider HSDir tier, faster lookups
Storj / Sia	Region-blind shard placement	Locality-aware shards + repair pub/sub
Tox / Briar messengers	Slow group chat, polling status	Pub/sub group chat, push status, real membership
YaCy	5–30 s queries	~300 ms DHT layer; interactive plausible
Radicle	5–30 s changeset propagation	~200 ms via axonal pub/sub
Handshake / ENS	Cold-cache slow	Implicit caching via hop caching
WebTorrent	Mainline-DHT discovery bridge	Discovery bridge eliminated; browser-native
Mastodon-style federation	Home-server bottleneck	Decentralized fan-out via axonal trees

Cross-application summary. Latency figures are back-of-envelope projections from the Axona simulator’s 25K-node benchmark, to be confirmed on the friction-realistic model.

THE PATTERN. The DHT layer is rarely the *only* bottleneck, but it is consistently *a* bottleneck, and moving to Axona consistently moves it out of the user-visible critical path. The biggest beneficiaries are the interactive applications (messengers, search, dev forges) where DHT latency sits in the user’s path; the least are those where the DHT is already a small fraction of total latency (Tor circuits, Storj’s remote-provider connection). The application that becomes *architecturally different* is the federation pattern, because Axona’s pub/sub primitive replaces the home-server fan-out those patterns suffer under.

Honest framing. We keep the federation claim modest: the architecture is *enabled*, not necessarily preferable in every dimension. And the latency figures are simulator projections, to be confirmed against the friction-realistic model before they are quoted as field results. What is not in question is the protocol surface they run on — authenticated, signed-and-fresh, with a real pub/sub lifecycle, live today.

II Future Applications: Axona Powered

PART I ASKED “what gets better if an existing app migrates?”

Part II asks the opposite: “what do we build that *only* Axona makes practical?” These are applications designed Axona-native from the first line — they consume the pub/sub lifecycle, the geographic locality, the signed-and-fresh envelopes, and the no-central-operator property as load-bearing features, not conveniences.

civildefense.io — the flagship

What it is. A tap-to-report incident map. A user taps a location to register a concern; reports propagate over anonymous peer-to-peer with a 24-hour expiry, surfacing on a shared regional map without any central server holding the data.

Why it is Axona-native. Every primitive the app needs is a protocol primitive:

- **Geographic locality is the product.** The S2 cell prefix in every `nodeId` means a report is relevant to, and routes toward, the people physically near it — at zero coordination cost. An incident map is, structurally, a geographically-keyed pub/sub feed, which is exactly what Axona routes.
- **Expiry is a first-class semantic now, not a side effect.** The original brief leaned on implicit expiry through the replay-cache LRU. With the v2.10.0 lifecycle, the 24-hour window is the **hold-time TTL** — an explicit, principled bound, so reports age out deterministically rather than whenever the cache happens to evict them.
- **Signed, fresh, and tamper-evident.** Reports are Ed25519-signed envelopes with per-publisher sequence numbers and a freshness window — so a stale or replayed report cannot be re-injected at a fresh subscriber, which matters acutely for an incident feed.
- **Anonymous but authenticated transport.** The `axona/4` handshake plus mesh channel-binding means reports travel end-to-end without a central operator and without a relaying bridge being able to read or rewrite them.

What it still needs from us. A membership-gated variant (e.g. a verified-responder channel) wants **Model 3 shared-encryption group keys** — the group key distributed at the application layer, the protocol carrying opaque ciphertext.¹¹

¹¹ `civildefense.io` is one of the two worked references — alongside the open-source reference app — driving that app-layer group-key design.

Status. First end-user application; the primitive fit is why it came together in weeks rather than quarters.

SYZL — the adaptive social feed

What it is. A swipe-based decentralized social feed where every gesture trains your network: *swipe right* to SYZL a post (it reshapes to your followers and your connection to the publisher strengthens), *swipe left* to FYZL it (it stops, and the connection weakens). Ranking runs transparently on your own device from your own swipe history — there is no algorithm operator and no server-side ranking.¹²

¹² Full product brief: SYZL-Brief.md in this folder.

Why it is Axona-native.

- **The mechanism mirrors the protocol.** Axona’s neuromorphic routing strengthens useful connections and prunes the rest; SYZL does the same one layer up, with user attention as the learning signal. Same algorithm, different layer — the protocol’s own story retold for end users.
- **No central operator, no engagement-maximization incentive.** No ads, no rage-bait amplification, no shadowbanning; ranking runs only on-device.
- **Verifiable reach without surveillance.** Publishers see reach (`publishes + reshare_count`) via `peer.metrics` without learning *who* reshared — the privacy property Axona was built for, surfaced at the app layer.
- **Connection-sharing as a feature, not a leak.** Because subscriptions are first-class protocol objects, “share my graph” is a native operation; recipients adopt connections selectively, with provenance.

Lifecycle fit. SYZL is almost entirely application-layer over `@axona/protocol`: `pub/sub` for posts, reshapes, and connection-list cards; `pull` for referenced bodies; `metrics` for the reach dashboard; `unsub` for dropping a rotated-out connection cleanly. The hill-climb plus simulated-annealing exploration that keeps the feed alive is ~200 lines on top of `lookup/peers`.

Syzl Anywhere. A companion Manifest-V3 browser extension that lets a user SYZL any page, selection, or image in one click, running the user’s Axona peer in a background service worker — which doubles as keeping the network warm while the user browses. Roughly six weeks of work, ~70% code-shared with the PWA composer.

The Pipeline — others as we invent them

This section is deliberately open. Part I is, in effect, a menu: several of those “what Axona changes” analyses describe applications that are more interesting built *fresh* on Axona than migrated. Natural near-term candidates, none yet committed:

- **A browser-native swarm transport.** WebTorrent’s discovery bridge eliminated — large-file or live-media distribution among browser peers with no Mainline-DHT intermediary.
- **A large-group peer-to-peer messenger.** Tox/Briar’s group-chat ceiling lifted by pub/sub fan-out plus Model 3 shared-encryption membership.
- **A federated-feed fan-out layer.** The Mastodon-style “viral post on a small instance” failure removed by moving fan-out to axonal trees while home servers keep storage and identity.

As each of these graduates from idea to product, it gets its own brief in this folder and a row in Part II.

What the Axona-Powered Applications Share

Across `civildefense.io`, SYZL, and the pipeline, the same four properties keep recurring as the reason the app is *possible*, not merely *cheaper*:

1. **Geographic locality for free** — the S2 prefix puts relevance and routing in the same place at zero coordination cost.
2. **Pub/sub with a real lifecycle** — pub/sub/unsub/pull/kill/unpub, bounded queues, and hold-time TTL turn “fan a message out” into a managed, expiring, revocable feed.
3. **Signed, fresh, end-to-end-secure messages** — authenticity, freshness, and MITM-resistance are protocol defaults, so the application never has to bolt them on.
4. **No central operator** — ranking, membership, and reach live on-device and in the mesh, which removes the surveillance-and-incentive failure mode common to their centralized counterparts.

These are the same properties Part I shows existing applications *reaching toward* through bridges, gateways, and caching tiers. The forward-looking bet of Part II is that designing for them from the first line produces applications the migration path cannot — because the protocol’s defaults are the product’s features.