

AXONA

Architecture

David A. Smith
Axona.net

An Axona Architecture Note v0.7.13 2026-06-12

A protocol is the contract between layers that don't trust each other.

What I cannot create, I do not understand.

RICHARD FEYNMAN · 1988

I The Stack

AXONA SHIPS AS THREE REPOSITORIES that compose into one running system. Each layer has a single responsibility,¹ a single canonical version identifier (`KERNEL_VERSION`), and a contract surface narrow enough to fit on a postcard.

@axona/protocol (kernel). The pure-JS library. Owns the DHT routing, the pub/sub overlay, the cryptography, and the transport contracts. No browser DOM, no Node-specific I/O, no network code outside abstract transport classes. Runs identically in browsers, in Node, in dht-sim’s simulated network.

axona-peer (browser application). The reference end-user app. Provides identity persistence, region selection, WebRTC mesh formation, a chat-style UI for pub/sub. Imports the kernel as a vendored copy under `vendor/axona-protocol/` so the version that ran when the page was built is the version that runs in the browser.

axona-bridge (deployed server). The signaling + introducer service at `bridge.axona.net`. Runs an embedded `AxonaPeer` of its own (so it’s a full DHT participant, not just a relay), gates connecting peers behind a version handshake, mints short-lived TURN credentials, relays WebRTC signaling between browsers, and forwards Axona wire frames addressed to its own embedded peer.

axona-relay (headless supernode). A Node process that joins the mesh over *real* WebRTC (a spec `RTCPeerConnection` from `node-datachannel`) and runs the same kernel stack a browser peer runs — routing lookups, rooting pub/sub topics, relaying signaling for others — but with no DOM, no public server, and no TURN minting. It is a strict *subset* of the bridge: it dials a bridge once for bootstrap, then is a first-class peer. Running several relays strengthens routing and pub/sub durability without any of them being an inbound server.

Each component carries its own semver track because each ships to a different audience: kernel changes need API-level discipline; peer changes are user-visible UI; bridge and relay changes are operational deploys. But they all share the same kernel version, visible at runtime through `/healthz` (bridge), the `welcome` frame (any peer connecting), and the version row in the peer UI.

¹ The pattern is intentional. A peer’s business logic should never need to know the wire format; a bridge should never need to know what topic a subscriber cares about; the kernel should never need to know which application is using it. Every interface in this document is a place where one layer was given the right to be wrong about the layer below.

A fourth component, `dht-sim`, runs the kernel against a simulated network for protocol research and regression benchmarks. It’s not part of the production stack but shares the same `@axona/protocol` library; every change that ships to production also has to pass the simulator’s 25,000-node battery.




Figure 1: *

Five layers in one peer. Every arrow is a contract; every contract has a smoke-test suite.

II Identity and Addressing

EVERY AXONA PARTICIPANT has a 264-bit identifier. The top 8 bits are a Google S2 cell prefix² computed from the peer’s latitude and longitude; the bottom 256 bits are the SHA-256 of the peer’s Ed25519 public key. A peer in London has an ID starting with `0x44`; a peer in Virginia starts with `0x89`. Two peers in the same region share a leading byte; two random peers share only the bits their keys happen to share.

The address space is treated as `BigInt` in every line of running code. The wire format is 66-character lowercase hex — the **ONLY** representation that ever appears in JSON. The conversion happens at exactly two places: when the transport serialises an outgoing frame, and when the dispatcher parses an incoming one.

Topic identifiers share the address space. A topic is named by the application (a string like `us-east/hello-world`) and anchored at a publisher: `topicId = sha256(publisher || topic-name)`, so its top byte is the publisher’s S2 cell. A regional synthetic publisher (top byte = S2 cell, bottom bits zero) anchors the topic in that region — which is what makes regional pub/sub route naturally: peers in the same region are XOR-closest to topics anchored there. A public topic (no publisher) is `00 || sha256(topic-name)`: the `00` prefix is simply the S2 cell at `00`, an ordinary cell with no special “global” meaning. Public topics are anchored there by convention, so every peer derives the same ID for a given name — and, like any topic, that ID is XOR-closest to the peers whose IDs fall in cell `00`.

The Activation Potential (AP) routing formula, the K-closest se-

² S2 projects the sphere onto the six faces of a cube and partitions each face along a Hilbert curve. We bin to level 3 (8×8 per face, 384 cells globally), then truncate to 8 bits — equivalent to taking the top 8 bits of Google’s 64-bit level-3 cell ID. Result: 192 valid cells, with byte values 192–255 reserved for future system use. The cube projection keeps cell areas within $\sim 1.4\times$ of each other globally — no polar shrinkage — and any standard S2 library can recompute our prefixes by truncation.

Each of the 192 cells carries exactly *one* human-readable name, so a region always presents the same label (no location-dependent flip-flop). Where a cell straddles an ocean and a landmass the land name wins; a multi-country cell takes its dominant city; homogeneous cells (one country, open ocean) keep their name. A name is usually unique to one code, though an area larger than one cell may span adjacent codes. The kernel ships the canonical table and the `regionName/regionCode/resolveRegion` helpers; the `s2-region-visualizer` example renders all 192 — one name and the hex code each — on a globe.

lector, and the synaptome’s stratum buckets all operate on this one space. There is no separate “topic ID space” or “peer ID space”; they share addresses because they share routing.

III The Kernel: @axona/protocol

THE KERNEL IS THE PROTOCOL. Everything else is its clients. The directory structure mirrors the contract surface:

src/contracts/. Abstract base classes: **Transport**, **DHT**, **DHTNode**, **BootstrapService**, **PersistenceAdapter**. These are the seams where alternate implementations plug in. A custom transport (libp2p, QUIC, embedded) extends **Transport**; a database-backed persistence (IndexedDB, SQLite) extends **PersistenceAdapter**.

src/dht/. The active routing layer: **NeuronNode**, **Synapse**, **AxonaPeer**, **AxonaDomain**, **Subscription**. This is where the Axona routing logic lives, and where the developer-facing API surface (**peer.pub**, **peer.sub**, **peer.lookup**) is declared.

src/pubsub/. The pub/sub overlay: **AxonaManager**, **post.js** (topic ID derivation), **envelope.js** (signed publish envelopes), **ed25519.js** (Ed25519 binding for **crypto.subtle**). **AxonaManager** owns the K-closest tree, the replay caches, the dedup LRU, the metrics counters.

src/transport/. Concrete transports: **transport/web/** (browser: BridgeTransport over WebSocket + WebRTCTransport over DataChannels), **transport/node/** (server-side WebSocket), **transport/sim/** (in-process simulation for dht-sim and tests). All implement the abstract **Transport** contract identically.

src/utils/. Pure helpers: **hexid.js** (BigInt↔hex conversions, XOR distance, leading-zero count), **s2.js** (S2 cell computation), JSON codec (**transport/wire.js**).

src/identity/. Ed25519 keypair derivation + persistence envelope. The single entry point **deriveIdentity** takes a {**lat**, **lng**} or full options object and returns {**id**, **pubkey**, **pubkeyHex**, **privateKey**, **region**, ...}.

The whole kernel has zero runtime dependencies beyond the host JavaScript environment (browser or Node 20+). No bundler, no transpilation, no build step. The same source files run in the browser via ES modules, in Node via direct **import**, and in the simulator via the in-process transport.

IV The Local Node

EACH PEER’S VIEW OF THE NETWORK lives in two data structures on the `NeuronNode`: the *synaptome* (its outgoing connections) and the *incoming synapse index* (peers known to point at it). Both are keyed by 264-bit `BigInt` `nodeId`.

A **Synapse** is a routing edge — it carries a peer ID, a measured latency, a learned weight, a stratum (the XOR-distance bucket), an inertia lock, and provenance (was this synapse added via handshake, hop-cache, triadic introduction, or sponsor bootstrap?). The synaptome is unbounded by design; the routing layer evicts synapses through the *Structural Survival Rule*, which guarantees navigability even under aggressive weight decay.

Routing decisions don’t pick the XOR-closest synapse blindly. Each candidate is scored by Activation Potential:

$$AP_c = \frac{\Delta \text{Distance}_c}{L_c} \times (1 + w \cdot W_c)$$

The first factor is progress velocity — how much XOR distance the hop closes per millisecond of measured latency. The second is a mild boost for synapses that have historically led to fast lookups. Weight only accrues on at-or-below-average-latency paths, so a high- W synapse genuinely represents a fast route, not just a frequently-used one. Two random hops with similar latency tie; the one that’s served us well before wins the tiebreak.

The rule, in one sentence: a synapse may only be pruned if at least one other synapse with the same stratum exists. The effect: every populated stratum keeps at least one live route, so the address space stays globally navigable even after heavy decay.

V The Protocol Layer

AXONAPEER IS THE ACTIVE OBJECT that turns a `NeuronNode` into a network participant. It owns:

- The handler installation. On `peer.start()` it registers the Axona routing handlers (`lookup_step`, `lookahead_probe`, `find_closest_set`, `route_msg`, `local_probe`, `triadic_introduce`, `hop_cache`, `lateral_spread`, `reinforce`) on the transport.
- The auto-admission loop. It asks `transport.bindPeers()` for currently-bound `nodeIds` and seeds them into the synaptome; it subscribes to `transport.onPeerBound` for future admissions.
- The pub/sub surface. `peer.pub`, `peer.sub`, `peer.pull`, `peer.metrics`: the developer-facing API. Each delegates to the `AxonaManager` after deriving the topic ID and (for `pub`) building a signed envelope.

- The direct-message surface. `peer.send`, `peer.notify`, `peer.onMessage`: per-peer request/response and fire-and-forget, outside the pub/sub overlay. For protocols that aren't topic-shaped.
- The lookup primitive. `peer.lookup(targetKey)` runs an iterative K-closest walk and returns `{found, path, hops, peerId, latency}`. This is what `AxonaManager` uses internally on every pub/sub to find the topic's axons.

`AxonaManager` is a separate object owned by the peer. Its job is the pub/sub overlay: when a subscriber sends `pubsub:subscribe-k` to the K-closest axons for a topic, the receiving `AxonaManager` stores the subscriber in its `role.children` for that topic. When a publisher fans `pubsub:publish-k` to its own K-closest, each receiving `AxonaManager` iterates its `role.children` and forwards `pubsub:deliver` to each subscriber. The publisher's signed envelope flows through unchanged; per-hop dedup prevents the fan-out tree from delivering duplicates.

The K-closest replication is what makes the system tolerant. A publisher sends to *five* root axons in parallel; a subscriber subscribes at *five* root axons in parallel. As long as one axon in each set overlaps, delivery succeeds. With $K = 5$ and synaptomes that converge under the bridge's peer-list dissemination, overlap is typically 4 or 5 of 5 — so a publish that loses two axons to transient WebRTC failures still delivers cleanly.

Three convergence mechanisms close the residual gaps where “one axon overlaps” is not guaranteed — a publish lands on the *publisher's* K-closest set, which need not be a given subscriber's. **Gap-safe replay** (kernel v2.37.0): a (re)subscriber reports the content-IDs it actually holds (a bounded *have* digest) and a root replays exactly the complement, so a hole in the middle is backfilled rather than masked by a single high-water timestamp. **Root-to-root anti-entropy** (v2.39.0): each root periodically exchanges digests of held content-IDs with its K-closest siblings and pulls what it is missing — and *re-verifies* every pulled message (publisher signature B-4 + content-address), so a sibling root cannot inject a forged or content-poisoned message and a killed message is never resurrected. **Kill convergence** (v2.36.0): recently-applied retractions are re-gossiped to the current root set, so a replica that missed the original kill still drops and tombstones the message rather than re-serving it to a reloading subscriber. All three are wire-additive — old and new nodes interoperate.

This section is the architectural placement of `AxonaPeer`; the field-by-field surface — every method, its parameters, and its return shape — is catalogued in the Application API section later in this document.

VI The Transport Layer

TRANSPORT IS WHERE ADDRESSES MEET SOCKETS. The kernel's **Transport** contract is a small 12-method surface: **start**, **stop**, **getLocalNodeId**, **openConnection**, **closeConnection**, **isConnected**, **send** (request/response), **notify** (fire-and-forget), **onRequest**, **onNotification**, **onPeerDied**, **getLatency**. Every concrete transport implements this surface and the upper layers don't care which one is wired in. (Peer *binding* — the authenticated **bindPeer**/**onPeerBound** admission flow — lives in the concrete web/node transports on top of this contract, not in the contract itself.)

Three concrete implementations ship:

transport/web/. The browser stack. A **CompositeTransport** fans out to two sub-transports based on which one owns each **nodeId**: a **BridgeTransport** that carries Axona wire frames over the same **WebSocket** the bridge uses for signaling, and a **WebRTCTransport** over a **MeshManager** that owns the data-channel pool. The factory **webTransport**({**bridgeUrl**, **identity**}) composes the two, drives the bridge's version gate, runs the hello/hello-ack admission, threads TURN credentials from the welcome frame into the mesh's ICE config, and resolves once everything is up.

The factory also owns the *connection lifecycle*, so an application doesn't hand-roll it: a 1 Hz ping/pong heartbeat with stale detection, auto-reconnect with exponential backoff on an unexpected socket close (re-running the version gate and hello/hello-ack on each re-open), and a small observability surface — a **bridgeState** machine (**connecting** → **open** → **stale** → **disconnected**, plus **upgrade-required** on a 4426 close), the captured **welcome** (**bridgeInfo**), and the live ping RTT (**bridgeRtt**). A 4426 version-gate rejection suppresses reconnect; **reconnectNow**() forces an immediate re-attempt (used on tab-resume and network-online).

The WebRTC mesh under **webTransport** self-heals on the same principle. **MeshManager** runs a 1 Hz ping/pong on every data channel and (kernel v2.1.1) evicts a peer the moment its channel goes dead: no pong for the heartbeat-timeout window, or a few consecutive throwing **send**()s — the signature of a channel a browser keeps reporting **open** after a laptop sleep, where the socket is gone but **readyState** lies. Eviction fires **onPeerLost** → **onPeerDied**, so the protocol layer drops the synapse and routes around the dead peer while the bridge's next peer-list rebuilds a fresh channel. This is the Transport contract's heartbeat-timeout obligation, and it means recovery from a connection freeze needs *zero*

application code. An application *may* add an accelerator: on a resume event (`visibilitychange`, `online`, `bfcache pageshow`) call `transport.mesh.reset() + reconnectNow()` to clear stale peers at once rather than waiting out the per-peer timeouts — both the demo and the axona.net peer do this. A macOS screensaver notably fires none of those resume events (the tab never goes hidden), which is why the transport-level eviction is the load-bearing mechanism and the app-level reset is only a speed-up.

transport/node/. A Node-side WebSocket transport. Used by the bridge's embedded peer to talk to connected browsers without simulating WebRTC — the bridge *is* the signaling server, so it doesn't need WebRTC to reach any browser; the WebSocket is direct.

transport/sim/. The simulator transport. Used by dht-sim's `TransportAxonaEngine` and by the kernel's smoke tests. Identical Transport contract, but messages are scheduled in the event loop with simulated latency derived from a real internet topology dataset.

The wire format is JSON, with a single quirk: `BigInt` is serialised as "`<digits>n`" (decimal digits with an `n` suffix) and revived back to `BigInt` on decode. `NodeId` fields are always hex strings on the wire — they never appear as `bigint`-suffixed literals because the kernel converts them at frame construction.

The same kernel test suite runs against all three transports. Cross-transport bugs — which is to say, the kind of bug where a fix lands in one transport's call path but not the others — can't ship.

VII The Bridge

THE BRIDGE IS AN INTRODUCER, not a server. Its job is to let browsers find each other; once they've found each other, WebRTC data channels carry every byte of application traffic peer-to-peer. The bridge stays in the picture only for two reasons: (1) it's the rendezvous for the version handshake before peers can talk at all; (2) it's a guaranteed-reachable mesh peer, which makes it a useful K-closest axon in small networks where two browsers might otherwise fail to find common axons.

The bridge handles three kinds of frames over the WebSocket:

Signaling. `welcome`, `peer-list`, `peer-joined`, `peer-left`, `signal`, `ping/pong`. Routes WebRTC SDP offers/answers and ICE candidates between browsers, broadcasts join/leave events, supplies short-lived TURN credentials in the welcome frame.

Version handshake. `client-hello` / `server-hello`. Enforces a minimum peer version. Rejects sub-minimum peers with WebSocket

close code 4426 and a human-readable reason; the peer’s UI shows a “please reload” banner.

Axona wire. {type: 'axona', payload: ...} frames carrying req/res/ntf messages between any connected browser and the bridge’s embedded peer. These are the same frames any two browsers exchange over WebRTC — the bridge embeds a full **AxonaPeer** of its own (**BridgeAxonaNode**) and participates in routing just like any other node, which is why it can be in K-closest sets and why it can hold subscriber children for pub/sub topics.

The bridge’s embedded peer has a fixed, disk-persisted identity (895bf0... in production, anchored in us-east — the 0x89 S2 cell, which is why it sits XOR-close to **us-east/*** topics and is so often a K-closest root). Its **health** surface exposes the kernel version it’s running and the synaptome it has built up, available at **/healthz** JSON. Operators check that endpoint to confirm a deploy.

How a New Node Joins

JOINING IS A MUTUAL, AUTHENTICATED BIND. A fresh node has no peers; it reaches the mesh through a bridge and then fans out to direct WebRTC links. Crucially, synaptome membership is established the *same* way on every link — bridge or mesh — so “joining” is just the first instance of the generic admit path, not a special case.

1. **Connect + version gate.** The node opens a WebSocket to a bridge and exchanges **client-hello/server-hello**. The bridge checks the wire major (**REQUIRED_WIRE_MAJOR**, currently 2 — the **axona/5** epoch) and a per-namespace version floor (**flagDayFloor**: a kernel-2.x floor or a peer-app-3.x floor); a peer on the wrong epoch or below the floor is closed with 4426 before it can do anything else. Because the epoch is folded into the signed transcript of the next step, this gate is not merely advisory: an **axona/4** node and an **axona/5** node *cannot* complete a handshake. The 2026-06-08 flag-day migrated production to **axona/5** and retired the **axona/4** network; any peer on a superseded epoch is partitioned out by construction.
2. **Welcome.** The bridge replies with **welcome**, carrying the connection id and a fresh **serverNonce**; both sides now hold the per-connection channel-binding value (CBV) the next step signs over.
3. **Authenticated handshake (axona/5).** Each side sends a **hello** (the responder a **hello-ack**) carrying its Ed25519 public key and a signature over the link’s CBV, and verifies the other: the presented

The kernel builds a peer’s **AxonaManager** lazily, on its first **pub/sub**. The bridge never publishes or subscribes — it only relays — so it must build that engine *eagerly* at startup, otherwise the **pubsub:subscribe-k/publish-k** handlers are never installed and every pub/sub frame addressed to the bridge is silently dropped. Because the bridge sits in nearly every peer’s K-closest set, that failure mode is a near-invisible delivery hole; the eager build is what keeps it an honest relay axon.

pubkey must hash to the 256-bit suffix of the nodeId it claims, and the signature must check out. A peer that cannot prove its id — or one being MITM'd, since the CBV binds the proof to *this* channel — fails here and is closed.

4. **Mutual bind + admit.** On success both ends record the binding and admit the other into their synaptome *as an ordinary Synapse*: the joining node calls `bindPeer(bridgeId, 'bridge')`, whose `onPeerBound` hook seeds the bridge into its synaptome (`addedBy: bootstrap`); the bridge's embedded peer runs `_completeHandshake`, binds the new node to that connection, and admits it via `_addByVitality` (`addedBy: handshake`). No bridge-specific code runs in either step — the bridge is simply the first guaranteed-reachable peer the node binds.
5. **Fan out to the mesh.** The bridge broadcasts `peer-joined` and hands the new node a `peer-list`. The node negotiates WebRTC data channels to those peers (SDP relayed via the bridge), and each channel runs the *same axona/5* hello — now additionally bound to the peer's DTLS-certificate fingerprint — firing the same `onPeerBound` admit. The synaptome fills with direct mesh peers; the bridge stops being the only neighbour.
6. **Converge.** Optionally `peer.join(sponsor)` seeds a specific first neighbour; thereafter vitality (`weight × recency`) and the 10s `refreshTick` grow and prune the synaptome and re-anchor subscriptions onto the converged K-closest set. The node is now a full participant.

VIII The Wire Protocol

EVERY FRAME ON THE WIRE is JSON. Every nodeId or topic ID field is a 66-character lowercase hex string. There's no binary framing, no protobuf, no length-prefixed records — the discipline of having one wire format that any of the three transports can carry is worth the few percent of size you give up.

Three kinds of message-typed envelopes flow over the wire:

Request (*k*: 'req'). A correlated request expecting a response.

Carries an integer `id` for correlation, a `type` string naming the handler (`lookup_step`, `route_msg`, `find_closest_set`, ...), and an opaque `body`.

Response (*k*: 'res'). The reply to a req. Echoes the `id` and adds an `ok` boolean and a `body`.

The admit is symmetric and identity-bound in both directions — each side admits the other's *authenticated* nodeId, so a peer can only ever be admitted under an id it proved it owns. Because every node runs step 4 against the bridge first, the bridge lands in every synaptome: its centrality is *emergent* from this bootstrap, not granted by any special rule (see § Future Directions, “demote the bridge to an ordinary peer”).

Notification (*k*: 'ntf'). Fire-and-forget. No id, no response. Used for hello/hello-ack admission, pubsub:subscribe-k, pubsub:publish-k, pubsub:deliver, triadic_introduce, hop_cache, lateral_spread, reinforce.

The handler types are versioned implicitly through the kernel — the bridge’s `minPeerVersion` gate excludes peers running an older protocol that wouldn’t understand newer messages. Major protocol changes bump `KERNEL_VERSION`’s minor; minor changes are additive (a peer that doesn’t know a handler type silently drops the message).

A pub/sub message envelope is signed:

```
{
  msgId:      "<sha256 of canonical(envelope) hex>",
  seq:        1716638400123,    // per-publisher monotonic (v2.9.0)
  ts:         1716638400000,
  topic:      "us-east/hello-world",
  message:    <application payload>,
  signerPubkey: "<Ed25519 pubkey hex>",
  signature:  "ed25519:<128 hex chars>"
}
```

Receivers verify the signature against the signer’s pubkey using `crypto.subtle.verify`; the kernel surfaces verification result on the delivered envelope so applications can decide whether to trust unsigned or unverifiable publishes.

As of kernel v2.9.0 (envelope format **v2**, finding C-2) the signature covers a domain-separated core (`axona:pubsub-envelope:v2`) and includes a per-publisher monotonic `seq`. Two properties follow. *Freshness*: a topic’s root axons reject a *live* publish whose signed `ts` is outside a ± 5 -minute window, so a captured envelope can’t be re-injected as new once it ages out (the legitimate replay-to-late-subscribers path is exempt — it serves cached history). *Order*: `seq` gives each publisher’s stream a total order, which is what makes “latest” and bounded-queue eviction (below) deterministic across replicas.

IX Routing

AXONA’S ROUTING IS LEARNING-ADAPTIVE. It replaces Kademlia’s k -bucket array with a dynamic synaptome that adjusts edge weights based on observed latency. The basic shape — iterative lookup, $\log_2(N)$ hop count — is unchanged from Kademlia; what changes is which peer you ask at each hop.

The five wire operations that drive Axona routing are:

lookup_step (*req*). “Here’s a target. What’s your best next hop?”

The receiver evaluates AP across its synaptome’s progress candidates, returns the winning synapse + its own routing state.

lookahead_probe (*req*). “If I forward to you for target *T*, who would you forward to?” A one-deep peek that lets the caller compare expected hop costs across candidates before committing.

find_closest_set (*req*). “Give me your *K* peers closest to this target.” The K-closest harvester used by AxonaManager’s pub/sub axon discovery.

route_msg (*req*). Generic routed-message forwarder. Used by AxonaManager for `__tunneled_direct__` delivery when the target isn’t directly bound.

reinforce (*ntf*). “This lookup succeeded; bump the weight on the synapse I used.” Fire-and-forget LTP signal back along the path. Only flows when the lookup hit at-or-below average latency for the path’s region. On identity-binding transports the handler is a no-op unless the *target* synapse’s peer is currently bound (the B-3 gate: an unauthenticated **reinforce** must not be able to refresh the eviction-inertia of a stale or unverified entry and pin it in the table). Weight is independently clamped to ≤ 1.0 , so inertia is the only lever the gate needs to close.

Two structural-plasticity messages let the network rewire itself without explicit instruction: **triadic_introduce** (“introduce me to your friend at *X*”) and **hop_cache** (“I’ve forwarded to *X* several times; install a direct synapse to *X*”). Both are fire-and-forget notifications that obey the synaptome’s admission rule (`_addByVitality`).

X Pub/Sub: K-Closest

THE PUB/SUB OVERLAY IS ITS OWN STRUCTURE. It sits on top of Axona routing but doesn’t reuse the synaptome for state — each peer maintains a separate set of *axon roles*, one per topic it’s chosen as an axon for. A role holds a topic ID, a set of children (subscribers), a replay cache, and a *K*-replicated peer set (the other root axons for the same topic).³

Subscriber-side flow:

1. Compute `topicId = sha256(publisher || topic-name)`.
2. Find the *K* peers closest to `topicId` in the local synaptome. (No network walk — the local reachable peer set is what matters; remote-walk K-closest returns ghost IDs that route fan-out into dead ends.)

Residual, by design: the gate keys on the *target* synapse being bound, not on the *sender*’s identity — so any bound peer can refresh the inertia of any other bound peer’s synapse. The blast radius is bounded to “keep an already-verified, live peer pinned slightly longer,” which is benign; binding the signal to the sender would add cost for no attacker-relevant gain. Revisit only if reinforce ever carries weight a sender could bias. The performance claim of Axona over Kademlia is roughly 40% fewer wall-clock milliseconds per lookup at 25,000 nodes, mostly from latency-weighted hop selection. The figure is documented in the companion paper *Axona: A Learning-Adaptive Distributed Hash Table and Axonal Publish/Subscribe* (v0.4.01, 2026-05-23).

³ Notation: this replication factor is the *root-set size* $R = 5$ (`AxonaManager.rootSetSize`)—distinct from the Kademlia *routing* parameter $k = 20$ (`AxonaDomain.k`), which sizes the closest-set used for *lookups*. “K-closest” throughout names the generic find-closest-*N* operation both use: $N = k$ for routing, $N = R$ for replication.

3. Send `pubsub:subscribe-k` to each, carrying the subscriber's `nodeId` and an optional `lastSeenTs` for replay.
4. Each receiver runs `_addOrRecruitChild`, stores the subscriber under `role.children`, and optionally serves a `pubsub:replay-batch` of cached messages newer than `lastSeenTs`.

Publisher-side flow:

1. Compute the same `topicId`.
2. Build a signed envelope.
3. Find the same K peers closest to `topicId` locally.
4. Send `pubsub:publish-k` to each.
5. Each receiver runs `_onPublishDirect`: dedup against `_seenPublishes`, add to replay cache, fan `pubsub:deliver` to every child in `role.children`.
6. Each child receives `pubsub:deliver`, dedups, fires the local `_deliveryCallback` which routes to the application's subscription handler.

When a publisher's K -closest peers also happen to be in the subscribers' K -closest sets, the tree converges and delivery is single-hop from publisher to axon to subscriber. When they don't, the system relies on the redundancy: a missing axon's subtree is covered by the four other axons. The Activation Potential weighting makes high-quality axons dominate the K -closest set over time.

Per-hop `_alreadySeenPublish` dedup is keyed on `publishId` (a string like `nodeId:counter`), separate from the content-hash `msgId` the application sees. Multiple arrivals of the same `publishId` at one peer collapse to a single `_deliveryCallback` firing.

`K-CLOSEST` IS COMPUTED AGAINST THE LOCAL SYNAPTOME, so in a browser mesh that fills in gradually the answer is a moving target: a peer that subscribed while it only knew the bridge picked a narrow axon set, and a publisher that joined later computes a wider one — they can fail to overlap, and the publish misses that subscriber. The result is cached per topic under an epoch counter to avoid recomputing on every `pub/sub`; the fix is to *invalidate* that cache as the mesh changes. `AxonaPeer` bumps the epoch on every `onPeerBound`, and `AxonaManager.refreshTick` (a 10s timer the application arms after subscribing) flushes it and re-issues every live subscription against the now-converged set. So a subscription anchored early re-anchors as the mesh widens, rather than staying pinned to a stale, narrow axon set.

Topic Metrics: Scatter-Down, Gather-Direct

`PEER.METRICS(TOPIC)` reports a topic's live state — `publishes`, `current_count` (events retained in the tree right now), `subscribers`, `deliveries`, `pulls`, `reshares` — by querying the same axon tree

that carries delivery. It is a *scatter-down, gather-direct* traversal, not a recursive sum back up the tree.

1. **Route to the tree.** `requestMetrics` issues a routed `pubsub:metricsReq` toward the `topicId`; a node that doesn't host the topic returns `'forward'`, so the request DHT-routes until it reaches a root axon.
2. **Scatter down.** On receipt a role replies with its *own* local view, then forwards the request as `pubsub:metricsBroadcast` to each of its children, which recurse — the probe walks the tree along the same parent→child edges that fan messages out. A bounded seen-request set (keyed on `requestId`) makes each relay process a given probe once, so redundant edges neither loop nor double-count.
3. **Gather direct.** Every responding relay sends its partial result *straight back to the original requester* (`pubsub:metricsResp`), not up through its parent. Scattering down but gathering directly keeps the root from becoming an aggregation hotspot.
4. **Bounded collection.** The requester accumulates responses for a fixed window (`timeoutMs`, default 500 ms) — there is no global barrier — then folds them. `relayCount` (distinct responders) is the honest coverage signal; the result is a best-effort snapshot, not a ledger.

WHY THE FOLD DIFFERS PER FIELD. Each relay holds a local slice: its replay cache, its `children` map, and per-post counters. The tree *replicates the message queue* across the K root axons but *partitions subscribers* across the tree, so a uniform sum would be wrong. The client combines each quantity according to its nature.

Field	Fold	Why
<code>publishes</code>	max	distinct posts are replicated across roots; summing would multiply by K
<code>current_count</code>	max	replicas hold the same retained set; max also tolerates a lagging replica
<code>deliveries,</code> <code>pulls,</code> <code>reshares</code>	sum	partitioned work counters — each relay tallies only the actions it performed
<code>subscribers</code>	max	membership is partitioned; exact for an un-split (single-root) topic, a lower bound once the tree splits
<code>relayCount</code>	count	distinct responders — a coverage indicator

Replicated set-cardinality folds with max; partitioned per-action counters fold with sum; partitioned membership folds with max (approximate).

A self-authenticating ownership gate means a relay answers only when the topic's publisher matches the requester, so today the owner alone gets non-empty results. `current_count` and live `subscribers` were added in kernel v2.11.0.

XI The Application API

EVERY AXONA APPLICATION TALKS TO THE PROTOCOL THROUGH ONE IMPORT. The barrel export at `@axona/protocol` is the entire contract surface — a small set of factory functions, one peer class, and a handful of helpers. This section is the field-by-field reference.

Identity

`deriveIdentity({ lat, lng })` *generate a fresh peer identity*

lat number. Latitude in decimal degrees.

lng number. Longitude in decimal degrees.

Returns a Promise resolving to `{ id, pubkey, pubkeyHex, privkey, privateKey, region }`. `id` is the 66-char hex `nodeId` (8-bit S2 prefix from `(lat,lng)` + 256-bit SHA-256 of the public key). `privateKey` is the in-memory `CryptoKey` ready for signing; `privkey` is the base64-encoded export for persistence. The whole envelope is what `AxonaPeer` expects as `identity`.

`dumpIdentity(identity)` *serialize for storage*

Returns a Promise resolving to a plain JSON object you can write to disk or IndexedDB. The reciprocal `loadIdentity(envelope)` reconstructs the `CryptoKey`-bearing form on the next launch.

Transports

`webTransport({ ... })` *browser transport: WebSocket to the bridge + WebRTC mesh*

bridgeUrl string. e.g. `wss://bridge.axona.net`.

identity identity envelope from `deriveIdentity`.

log (optional, default `() => {}`) function (`event`, `data`) that receives transport-internal diagnostic events.

WebSocketImpl (optional) constructor for the `WebSocket` class. Defaults to `globalThis.WebSocket`.

autoHandshake (optional, default `true`) when true, `transport.start()` drives the full bridge admission sequence end-to-end — the WebSocket version gate, the application-level `hello/hello-ack` exchange, and binding the bridge as a peer. Set to false only if you're driving the handshake yourself.

peerVersion (optional, default `KERNEL_VERSION`) semver string sent in `client-hello`.

handshakeTimeoutMs (optional, default 15000) how long to wait for the bridge's `hello` before rejecting the *first* `start()`.

reconnect (optional, default `true`) auto-reconnect with exponential backoff on an unexpected socket close (bounds tunable via `reconnectInitialMs` / `reconnectMaxMs`). A 4426 version-gate rejection is the one close that does *not* reconnect.

Returns a `CompositeTransport`. After `await transport.start(identity.id)` the bridge is reachable as a peer (`transport.bridgeNodeId`) and the WebRTC mesh is ready to accept channels. Beyond the `Transport` contract the returned object carries the connection-observability surface the factory now owns:

onBridgeState(cb) subscribe to state transitions `connecting` | `open` | `stale` | `disconnected` | `upgrade-required`; `bridgeState` is the current value, `upgradeReason` the 4426 detail.

onWelcome(cb) the bridge's `welcome` (`{connId, version, kernelVersion, turn}`), also cached on `bridgeInfo` and replayed to late subscribers.

bridgeRtt / *bridgeRttAvg* last and windowed ping→pong round-trip in ms.

onPingTraffic(cb) per-frame (`nodeId`, `'sent' | 'recv'`) pulses for live link indicators.

reconnectNow() force an immediate reconnect with the backoff reset (`tab-resume` / `network-online`).

`nodeTransport.server({ ... })` *server-side WebSocket transport (Node 20+)*

`nodeTransport.client({ ... })` *Node WebSocket client*

Used by `axona-bridge`'s embedded peer and by Node applications. Same `Transport` contract as the browser transport; no WebRTC.

`simTransport({ network, identity })` *in-process simulation transport*

network `SimNetwork` shared across peers.

identity identity envelope.

For tests and `dht-sim`. Messages are scheduled in the event loop with synthetic latency from a real internet topology dataset.

AxonaPeer

`new AxonaPeer({ ... })` *the per-node DHT instance*

node `NeuronNode` instance for this peer (carries id, lat, lng, the synaptome).

domain `AxonaDomain` the peer joins. Construct one per network with `new AxonaDomain({ k: 20 })`, where `k` (*default 20*) is the *K*-replication factor.

transport the transport instance from `webTransport`, `nodeTransport`, or `simTransport`.

identity (*optional*) identity envelope. Required for signed publishes (the default). Apps that only call `peer.pub(t, m, { sign: false })` can omit it.

axonaManager (*optional*) explicit `AxonaManager`. When omitted, one is built automatically on `peer.start()`.

persist (*optional*) `PersistenceAdapter` for synaptome / subscription / replay-cache checkpoints. Default: no persistence.

Lifecycle

`await peer.start()` *wire handlers, install delivery hook*

`await peer.stop()` *detach handlers and event listeners*

`await peer.join(sponsor?)` *production bootstrap path*

sponsor (*optional*) 66-char hex node ID. If given, opens a transport channel to that peer and seeds the synaptome from it. If omitted, returns immediately — the peer is “joined” but isolated; inbound connections from other peers populate the synaptome.

`await peer.leave(opts?)` *leave the network gracefully*

opts.drain (*default true*) wait for in-flight publishes to settle.

opts.notify (*default true*) send a `peer-leaving` notification to every peer in the synaptome.

opts.timeoutMs (*default 5000*) maximum drain window in ms.

Pub / Sub

`await peer.pub(topic, message, opts?)` *publish to a topic*

topic string. Application-level topic name (any non-empty string).

message JSON-serializable payload.

opts.sign (default **true**) sign the envelope with the identity's Ed25519 key.

opts.publisher (default: *this peer's nodeId*) controls where the topic is anchored in the DHT. **undefined** means this peer publishes its own topic; **null** means the public well-known mode (**topicId**=SHA256(**topic**), anyone can publish); a 66-char hex string lets a peer relay under that publisher's anchor (rare).

Returns a Promise resolving to the **msgId** — the 64-char hex content hash of the envelope. Throws **PublishError** on invalid inputs or signing failure.

`await peer.sub(topic, handler, opts?)` *subscribe to a topic*
topic string.

handler function (envelope) => void invoked for each delivery.

The envelope is { **msgId**, **seq**, **ts**, **topic**, **message**, **publisher**, **signerPubkey** }. A retraction (see **peer.kill**) arrives on the same handler as a marker, { **deleted**: **true**, **msgId**, **topic** } — branch on **env.deleted** to drop your local copy.

opts.publisher (default: *this peer's nodeId*) same semantics as **pub** — must match the publisher's **publisher** value for the subscription to land at the right topic ID.

opts.since (default: *live tail*) replay control. Pass **'latest'** for the most-recent cached message plus future deliveries, **'all'** for the full replay cache plus future, or a Unix-ms **number** for everything newer than that timestamp.

Returns a Promise resolving to a **Subscription** handle. Call **sub.stop()** to cancel that handle, or **peer.unsub(topic, opts)** (kernel v2.10.0) to leave the topic by name — stops every local subscription for it and, when the last goes, sends the network unsubscribe (self-only by construction).

`await peer.pull(msgId, opts)` *fetch a message from the relay's replay cache — by ID, or the latest*

msgId 64-char hex content hash, or **null**/omitted to fetch the topic's most-recent message (highest **seq**; kernel v2.10.0).

opts.topic string. The topic name (required, used to locate the relay).

opts.publisher 66-char hex publisher ID, or **null** for public topics.

opts.timeoutMs (default 1000).

Returns the envelope or **null** if not found (which includes a message that has *expired* past its hold time). A successful pull slides the message's hold forward, bounded by its 48 h ceiling.

`await peer.metrics(topic, opts)` *aggregate counters across the relay tree*

topic string.

opts.publisher 66-char hex publisher ID, or `null`.

opts.timeoutMs (default 500).

Returns { `publishes`, `subscribers`, `deliveries`, `pulls`, `reshares`, `relayCount` }.

`await peer.host(topic?, opts?)` *store + serve a topic for others, without subscribing (kernel v2.40.0)*

topic string, or omitted. With a topic, the node hosts that one topic.

Omitted, the node hosts its **own keyspace neighborhood** — it is recruited as a root for whatever topics fall near its node ID (“host whatever lands near me”), the natural mode for a relay.

opts.publisher same publisher-anchor semantics as `sub`, selecting which feed’s topic ID to host.

Hosting is **decoupled from subscribing**: a host stores and serves a topic for other peers but is *not* a consumer of it — it registers no delivery handler and is never added to the node’s own subscription set. Mechanically it announces itself with the *same* `pubsub:subscribe-k` a subscriber already sends, so existing roots recruit it with no wire change (no flag day), then caches and serves replays exactly as a root does. It respects the proximity gate (B-2): a host only ends up root-ing topics it is genuinely among the K-closest to, so the primitive cannot force a node into the serving set for arbitrary or distant topics. `peer.unhost(topic?)` reverses it; `health().hosting` reports { `keyspace`, `topics` }. This is the participation primitive for relays and other infrastructure that should *carry* a topic without consuming it — before it, a node entered a topic’s serving fabric only by subscribing, so a relay that meshed but never subscribed carried nothing.

Message lifecycle (kernel v2.10.0)

Three lifecycle controls, each **self-authenticating** — the authority to act is proven by the same keypair that proved authorship/ownership, with no registry or central gatekeeper.

`await peer.kill(topic, msgId, opts?)` *retract a message you published (creator-only)*

A topic’s root axons accept the kill only if it is signed by the *same key that signed the original message*; they drop it from their replay caches, record a short-lived tombstone (so a lagging replica can’t resurrect it), and forward a delete marker to subscribers. It is best-effort redaction,

not a cryptographic un-send — a subscriber that already received the plaintext can keep it — and an anonymous (**sign: false**) message has no provable creator and so cannot be killed. Throws **KillError**.

await peer.unpub(topic, opts?) *remove a topic's whole queue (owner-only)*

The *owner* is the identity whose **nodeId** seeds the topic ID; roots verify this self-authenticatingly (the signer's pubkey must bind to the owner **nodeId** and that **nodeId** must derive the topic ID). **opts.destroy** (default **false**) escalates from “clear the messages” to total removal (messages and config/ownership state). Public (ownerless) topics cannot be unpublished. Throws **UnpubError**.

Topic limits

Each topic's replay queue is bounded (kernel v2.10.0): **max messages** (default 100, ceiling 256) with deterministic eviction of the lowest-ordered message by signed (**seq**, **ts**, **msgId**) — a max of 1 is a retained / latest-value slot; a **hold time** (default 24 h, hard ceiling 48 h) after which a message expires and is swept; and, on *open* (anyone-may-publish) topics, a **per-publisher quota** (one-quarter of the queue) so a single publisher can't flood out the rest. Owner-gated topics are governed by their publish rules instead.

Direct messaging

await peer.send(targetId, message) *RPC-style; awaits the remote handler's return value*

peer.notify(targetId, message) *fire-and-forget*

peer.onMessage(handler) *register a handler for incoming direct messages*

targetId 66-char hex node ID. The peer must already be in the synaptome.

message JSON-serializable payload.

handler function (**fromId**, **message**) => **result**. The return value is sent back to the caller on **send**; ignored on **notify**. Async handlers are awaited.

Mesh introspection

peer.peers() returns an array of 66-char hex node IDs currently in the synaptome.

peer.onPeerJoin(handler) registers (**nodeId**) => **void** for new arrivals.

`peer.onPeerLeave(handler)` registers `(nodeId) => void` for departures.
`peer.health()` returns `{ started, transportConnected, synapseCount, subscriptionCount, kernelVersion }`.
`peer.onLog(level, handler)` `level` is 'info', 'warn', or 'error'.
`peer.onError(handler)` registers `(err) => void` for unhandled errors.

Lookup

`await peer.lookup(targetKey)` *iterative K-closest walk*

targetKey 66-char hex 264-bit key. Usually a topic ID or a node ID.

Returns `{ found, path, hops, peerId, latency }`. `found` is the closest reachable peer; `path` is the list of nodes traversed; `hops` is the count; `latency` is wall-clock ms.

Snapshots

`await peer.snapshot()` returns a JSON object capturing identity, synaptome, axon roles, subscriptions, replay caches. Symmetric with `AxonaPeer.fromSnapshot(snap, deps)`, which reconstructs a peer from one of these dumps without touching the network.

Helpers

`await deriveTopicId(publisher, topic)` returns the 66-char hex topic ID. `publisher` may be null (public mode) or a 66-char hex node ID. Useful for displaying the publish location in a UI.

`geoCellId(lat, lng, bits)` returns the integer S2 cell at the given precision. `bits = 8` matches the top 8 bits of every Axona `nodeId`.

XII Identity and Security

IDENTITY IS AN ED25519 KEYPAIR. The kernel uses `crypto.subtle.generateKey` to produce a curve-Ed25519 pair; the public key is hashed with SHA-256 to produce the bottom 256 bits of the `nodeId`; the top 8 bits are the S2 cell of the user's (`lat`, `lng`). The private key never leaves the host environment, but the public key is broadcast on every connection (in `hello/hello-ack`) so peers can verify each other.

Every signed publish carries the publisher's public key (`signerPubkey`) inline. The kernel verifies the `ed25519:<hex>` signature against `signerPubkey` using `crypto.subtle.verify`. Verification is per-envelope, not per-peer — a peer that's untrusted can still relay signed

publishes from a trusted publisher unchanged, and any receiver along the way can verify the chain of custody.

Persistence is opt-in. The kernel’s `PersistenceAdapter` contract exposes `write(namespace, value)` and `read(namespace)`; concrete adapters target IndexedDB (browser) and the filesystem (Node). When wired, AxonaPeer checkpoints identity, synaptome seeds, and pending subscriptions on a 5-second debounce. On restart, the identity envelope is reloaded and validated (the stored `nodeId` must match the freshly-derived one); subscriptions are reinstalled as soon as the mesh has stabilised.

The `crypto.subtle` dependency is a hard requirement: the demo and the peer fail to boot in insecure contexts (HTTP pages outside localhost) because the Web Crypto API simply isn’t available. The bridge serves over HTTPS; the peer’s GitHub Pages deployment enforces HTTPS at the Pages level; the demo at `demo.axona.net` required Let’s Encrypt cert provisioning before identity derivation would run at all.

Proof of work, decoupled from the address

IDENTITY MINTING IS GATED BY A PROOF OF WORK (the E-1 placement defense), shipped *inert* at difficulty 0 in kernel v2.34 so that difficulty is a tunable protocol parameter rather than a format change.⁴ The work is a search for a nonce such that $H(\text{domain} \parallel \text{role} \parallel \text{pubkey} \parallel \text{nonce})$ carries at least N leading zero bits; the winning nonce rides as a sibling wire field (`pow` on the hello, `signerPow` on the envelope). Per-role difficulty is intentional — a transport identity can be ground into a root position (eclipse) and warrants a higher cost; a publish key can only flood, so its cost is calibrated to anti-flood, not anti-eclipse.

CRUCIALLY, THE WORK IS NEVER ON THE ADDRESS. The `nodeId` remains `[prefix] || SHA-256(pubkey)` and is unconstrained by the PoW at any difficulty. Encoding the work *into* the address — the “vanity address” form, S/Kademlia’s static puzzle — would be catastrophic for a heterogeneous keyspace: forcing leading-zero address bits clusters every PoW’d publisher (and the owned topics it anchors) onto `0x000...`, the keyspace edge, collapsing the DHT’s uniform load distribution — and only a *subset* of identities (publish keys) is PoW’d, while transport IDs and topics stay uniform. A trailing-zero variant would preserve placement (XOR closeness is dominated by the high bits) but re-couple difficulty to the address format, reviving the flag-day. The separate nonce avoids both: zero address bits constrained, full address entropy, difficulty tunable in place.

Because identity is bound to a region prefix, moving geographic location changes the `nodeId`. axona-peer’s `identity.js` preserves the existing identity across reloads unless the user explicitly re-selects a different region. A roaming user keeps the same DHT identity for the life of the browser profile.

⁴ ~~new crypto.~~ A. Back, *Hashcash — A Denial of Service Counter-Measure* (2002; orig. proposed 1997) — the $H(\text{input} \parallel \text{nonce})$ leading-zero-bits primitive; I. Baumgart & S. Mies, *S/Kademlia: A Practicable Approach Towards Secure Key-Based Routing* (ICPADS 2007) — crypto-puzzle Sybil/eclipse resistance for DHTs, and the *static* (ID-constraining) vs *dynamic* (separate-nonce) puzzle distinction this design turns on; J. R. Douceur, *The Sybil Attack* (IPTPS 2002) — the upstream result that a self-authenticating network needs a resource proof at all.

Before any difficulty > 0 ships, the scaffolding work hash (currently SHA-256) must be replaced with a memory-hard function so phone/browser minters aren’t out-classed by ASIC grinders.

XIII Production Deployment

THE PUBLIC STACK RUNS AT THREE URLS.

bridge.axona.net. The deployed bridge. Runs on a DigitalOcean droplet, started by systemd, kept current via `git pull && npm ci --omit=dev && systemctl restart`. Exposes `/healthz` for monitoring. Listens on `wss://` only; HTTP requests are redirected.

axona.net. The reference peer application, served via GitHub Pages from the `axona-peer` repo's main branch. Every push redeploys. The vendored `axona-protocol` kernel ships with the app; the version row in the UI confirms which kernel an open tab is running.

demo.axona.net. The minimal pub/sub demo, served via GitHub Pages from the `axona-protocol` repo's main branch (where `examples/minimal-pubsub-browser/` lives). Same Pages deployment pattern; HTTPS-enforced because `crypto.subtle` requires a secure context.

A coordinated release works as follows: a fix lands in `axona-protocol`, bumps `KERNEL_VERSION` and the package version, tags the commit. `axona-peer` resyncs its `vendor/axona-protocol/` from the new kernel, bumps `PEER_VERSION`, and pushes (Pages picks up). The bridge updates its `@axona/protocol` dependency in `package.json`, bumps `VERSION` in `server.js`, pushes, and the operator runs the deploy script. Every component ends up reporting the same `kernelVersion` in its visible metadata.

XIV Future Directions

THE ARCHITECTURE IS STABLE. The shape of the layering, the wire format, the BigInt-canonical address discipline, the K-closest pub/sub semantics — these are not expected to change. What's coming sits inside the existing seams:

Federated bridges. Multiple bridges in different geographic regions, gossiping peer-lists so a peer connected to one bridge can be reached from another. Single-bridge today is a deployment simplification, not a protocol constraint.

Demote the bridge to an ordinary peer. Federated bridges remove the single *rendezvous*; this companion change removes the single *super-peer*. The routing and pub/sub layers already treat a bridge as just another `nodeId` — there is no `isBridge` flag, no eviction exemption, no routing privilege.⁵ The residual special status is purely at the *bootstrap* layer: a peer holds its bridge WebSocket

Every Axona component is required to surface its `KERNEL_VERSION` at a runtime-inspectable point: the bridge in `/healthz` and the welcome frame, the peer in its UI version row, the demo in its footer. The rule exists because before it did, “which kernel is everyone on?” required SSH-ing into the bridge and grepping `node_modules`.

⁵ The bridge's present dominance is *emergent*, not granted: every peer dials it first, so it lands in every synaptome and carries the highest degree; and its `0x89` geo-identity sits XOR-close to every `us-east/*` topic. The ordinary K-closest and vitality rules then keep selecting it as a root. The topology makes it central; the algorithm gives it nothing.

open for signaling and *auto-readmits* the bridge to its synaptome on every (re)connect, so the bridge never actually leaves the routing table. Once a peer holds enough healthy mesh peers it should *drop the bridge from routing and root candidacy* — keeping the socket open for signaling only — and stop auto-readmitting its `nodeId`. With several bridges sharing the load, this also dissolves the *correlated* failure mode in which one bridge restart removes a root from nearly every topic at once. The goal state: a bridge is a signaling endpoint that any operator can stand up, indistinguishable from any other peer to the routing fabric. Validate convergence and churn behaviour in `dht-sim` before shipping.

Bridge directory. A well-known public topic (`deriveTopicId(null, "axona.bridge-directory")`)

where every bridge republishes itself once per hour: URL, region, kernel version, advertised TURN credentials, last-seen timestamp.⁶ Peers subscribe at startup and cache the directory; a peer whose primary bridge disappears has a list of alternates to try without out-of-band coordination. To prevent spam and DDoS of the channel, the topic uses an *allowed-signer set*: only Ed25519 keys with a valid attestation chain back to one of the genesis bridge keys (baked into the kernel) can publish. Subscribers verify the chain on every publish and drop unattested envelopes at the AxonaManager layer, before they reach the application. Genesis keys can sign attestations admitting new bridges without a kernel release. This turns “federated bridges” from a deployment story into something clients can actually discover at runtime.

⁶ Subscribers keep the newest entry per `signerPubkey` and prune entries older than 24 hours. Being a public topic, its ID carries the `0x00` prefix, so every peer derives the same directory ID and sees the same listing regardless of geography.

Capability hardening (endo / SES). Adopt the Agoric and MetaMask object-capability stack to harden the host-side surface, in three landings: `@endo/marshal` for envelope payloads (so payloads can carry `BigInt`, remotable references, and the slot-and-body wire shape `captp` expects);⁷ an opt-in `AxonaPeer.create({secure: true})` that calls `lockdown()` to freeze JavaScript primordials before kernel init, defending against prototype pollution and supply-chain mutation; and `@endo/captp` layered over `peer.send/peer.onMessage` to upgrade direct messaging from string-handler RPC to capability-style RPC with promise pipelining and revocable handles. A later phase ships `@endo/import-bundle` for content-addressed, signature-verified code distribution — the substrate for running third-party agents and creator-published filters inside peer compartments with only the authority each receiver chose to grant.

⁷ Marshal output is JSON-shaped with `@qclass` annotations on typed values, so legacy peers see plain objects and capability-aware peers see typed data. Negotiated per-envelope via a `payloadFormat` tag.

Persistent topics. Today’s replay cache survives in-memory only.

A pluggable persistence layer for axon `role.replayCache` (with a TTL) lets late subscribers catch up across bridge restarts.

Topic access control. Public-mode topics (`publisher: null`) let

anyone publish. Private topics (`publisher: <fixed-pubkey>`) accept publishes only from that key. A capability layer where subscribers verify per-publish signatures against an allowed-signer set lives at the application layer today; lifting it into the kernel is straightforward.

Cross-tab state sharing. Multiple tabs on the same browser today derive the same identity and form independent peers; a BroadcastChannel-based shared peer would reduce mesh load proportional to tab count.

Native transport. A C-based or Rust-based embedded kernel for IoT-class devices, sharing the wire format with the JS kernel so existing peers can talk to it without adaptation.

Connection-time optimization. The authenticated-identity handshake costs only sub-millisecond crypto (one Ed25519 sign plus one verify per link); the felt cost of a cold start is dominated instead by WebRTC ICE/DTLS bring-up, un-bundled module loading, and the few added round-trips of the mesh handshake — none of which touches steady-state routing or delivery latency. The recovery levers all sit at the client and transport layer, independent of the wire protocol: bundle the kernel modules, render the UI and authenticate in the background, and fold the mesh handshake into the DTLS/SDP setup so the signed hello rides the first data-channel frame.⁸ Not yet mission-critical; deferred until adoption makes first-connect time the binding constraint.

Black-hole detection (forwarding vitality). Every hardening to date answers “*is this peer who it claims to be, and may it do this?*” — identity, channel binding, ingress authorization, verified routing admission. A *black-hole node* passes all of them and then silently drops (or, equivalently, fails to forward) the traffic it should relay. This is a Byzantine behavioural fault, not an identity one. Tampering is already defeated (end-to-end signatures + A-1) and replay/re-order by C-2’s `seq`+freshness; *omission*, though, is only *inferable*, never provable — you cannot prove a negative, and a malicious drop is indistinguishable from a crash or a bad link. The direction is a *local* forwarding-vitality signal — each node’s own first-party measurement of whether traffic routed through a peer reached an acked destination, over the existing K-closest path redundancy — driving route-around, never a gossiped reputation verdict (which the no-central-authority constraint forbids). Equivocation is the one bad-faith behaviour that *is* provable (signed conflicting artifacts) and can land earlier. The gating question — the false-positive rate against honest-but-flaky nodes — must be answered in simulation first; see the black-hole threat note and the heterogeneous-protocol

⁸ The last lever pairs naturally with binding the channel-binding value to the DTLS certificate fingerprint: once the CBV derives from the fingerprint already exchanged in the SDP, no nonce round-trip is needed and the handshake adds no extra latency — security and speed land together.

simulation plan in `axona-docs`. This is the behavioural complement to verified admission: that governs *who is in your routing table*; this governs *whether they actually do their job*.

Decoupled publish identity. Today one Ed25519 key serves four roles — transport auth, routing identity, subscribe origin, and *authorship* (the signature on every published envelope). A publication therefore carries the publisher’s transport identity, with its persistent `nodeId`, region prefix, and channel-bound mesh presence. The direction is to let a peer sign publishes with a *separate* publish key, so authorship becomes a free-standing pseudonym independent of the node it was published from.⁹ The envelope format already supports it — the signature verifies against the envelope’s `signerPubkey` field and the `msgId` folds in that same key as the author, so neither verification (B-4) nor the wire changes; the only coupling is that `peer.pub` defaults the signer to the node key. *This is the rare change that is additive but effectively permanent*: it needs no flag-day to introduce (old and new envelopes verify identically), yet once content carries publish-key authorship it cannot be cleanly retracted — so it can be added at any time, on its own schedule. Three things must be designed alongside it, none in the wire format: **(1)** it anonymizes the *author field*, not the *act* — a first-hop axon can still time-correlate origination, and a node-anchored `topic_id` re-leaks the publisher, so anonymized publishes must target region-synthetic or publish-key-anchored topics; **(2)** anti-abuse must re-bind — a free, infinite supply of publish keys defeats per-publisher quota, so rate-limiting attaches to the transport node or the publish key carries proof-of-work (it meets the open E-1 / quota threads here); and **(3)** the per-publisher `seq` (C-2) is per-*signer*, so the same publish key used from two devices needs coordinated sequencing. Authorship trust then becomes a key-distribution question (which `pub_pk` is whom), exactly as for any pseudonymous-identity system.

⁹ This is the Nostr `npub` model: a portable author identity that can publish from several devices (different transport keys, same publish key), and that survives node churn.

The job of the architecture is to make these additions land cleanly. The job of the existing layers is to keep working while they do.

Companion documents: the Axona Programmer Guide and Axona API Reference, the Axona Explainer, and the source repositories at github.com/axona-net.

A Appendix: The Demo Application

THIS IS THE COMPLETE SOURCE of the minimal-pubsub demo deployed at `demo.axona.net`. The file lives in the kernel repository under `examples/`

`minimal-pubsub-browser/index.html`. Just over six hundred lines of HTML, CSS, and JavaScript, no build step, no dependencies outside the kernel barrel export. Open it in two tabs; one publishes, the other receives. The structure maps to the API reference in §XI section by section: identity, transport, peer construction, lifecycle, subscribe, publish — and, since kernel v2.1, the connection-state and ping-traffic hooks that drive the dot strip and survive a bridge restart.

examples/minimal-pubsub-browser/index.html

```
<!doctype html>
<html lang="en">
<head>
  <meta charset="utf-8">
  <meta name="viewport" content="width=device-width, initial-scale=1, viewport-
    fit=cover">
  <meta name="theme-color" content="#2d7373">
  <title>Axona Demo</title>

  <style>
    :root {
      --bg:          #ffffff8;
      --ink:         #1a1a1a;
      --muted:       #6b6b6b;
      --card:        #ffffff;
      --border:      #e2dfd6;
      --accent:      #2d7373;
      --accent-soft: #d4e4e1;
      --dot-bridge:  #2d7373;
      --dot-open:    #22c55e;
      --dot-pending: #f59e0b;
      --dot-dead:    #cbbc4;
      --tile-bg:     #f4f1e8;
    }
    * { box-sizing: border-box; }
    html, body { margin: 0; padding: 0; background: var(--bg); color: var(--
      ink); }
    body {
      font-family: system-ui, -apple-system, 'Segoe UI', Roboto, sans-serif;
      line-height: 1.5;
      min-height: 100vh;
      padding: env(safe-area-inset-top) env(safe-area-inset-right) env(safe-
        area-inset-bottom) env(safe-area-inset-left);
    }
    .wrap {
      max-width: 880px;
      margin: 0 auto;
      padding: 1.4rem 1.1rem 3rem;
    }

    /* Header */
    header {
      display: flex;
      align-items: baseline;
      justify-content: space-between;
      gap: 1rem;
```

```

    margin-bottom: 0.6rem;
}
h1 {
    font-size: 1.5rem;
    font-weight: 600;
    letter-spacing: -0.01em;
    margin: 0;
    color: var(--accent);
}
.status {
    font-size: 0.8rem;
    color: var(--muted);
    font-variant-numeric: tabular-nums;
}

/* Dot strip (mesh) */
.dots {
    display: flex;
    flex-wrap: wrap;
    gap: 4px;
    padding: 0.5rem 0 1rem;
    min-height: 18px;
}
.dot {
    width: 7px;
    height: 7px;
    border-radius: 50%;
    background: var(--dot-dead);
    transition: background 0.15s, transform 0.15s, opacity 0.15s;
    opacity: 0.55; /* resting (no recent traffic) */
}
.dot.bridge { background: var(--dot-bridge); width: 9px; height: 9px; }
.dot.open { background: var(--dot-open); }
.dot.pending { background: var(--dot-pending); }
.dot.dead { background: var(--dot-dead); }

/* v2.0.2: pulse-on-ping. The dot brightens when a ping/pong
   frame crosses its channel. Class added on event, removed
   400 ms later → dot fades back to its resting opacity. Steady
   channels at 1 Hz visibly blink twice/second; stalled channels
   go dim within ~1 s. */
.dot.pulse {
    opacity: 1;
    transform: scale(1.35);
}

/* Topic info row */
.topic-row {
    display: flex;
    flex-wrap: wrap;
    gap: 0.5rem;
    margin-bottom: 1.2rem;
}
.tile {
    background: var(--tile-bg);
    border: 1px solid var(--border);
    border-radius: 6px;
    padding: 0.4rem 0.7rem;
    font-size: 0.82rem;
    color: var(--ink);
    display: flex;
    flex-direction: column;
    gap: 1px;
}
.tile.k { color: var(--muted); font-size: 0.7rem; letter-spacing: 0.04em;
    text-transform: uppercase; }
.tile.v {
    font-family: ui-monospace, 'SF Mono', Menlo, monospace;

```

```

    font-size: 0.86rem;
}
.tile.topic .v { color: var(--accent); font-weight: 600; }
.tile.prefix .v { color: var(--accent); font-weight: 600; }
.tile.topcid .v { color: var(--muted); font-size: 0.74rem; word-break:
break-all; }

/* Two-pane grid */
.panes {
    display: grid;
    grid-template-columns: 1fr 1fr;
    gap: 1rem;
}
@media (max-width: 700px) {
    .panes { grid-template-columns: 1fr; }
}
.pane {
    background: var(--card);
    border: 1px solid var(--border);
    border-radius: 10px;
    padding: 1rem 1.1rem;
}
.pane h2 {
    font-size: 0.78rem;
    font-weight: 600;
    letter-spacing: 0.06em;
    text-transform: uppercase;
    color: var(--muted);
    margin: 0 0 0.7rem 0;
}

/* Publish form */
label.field {
    display: block;
    margin-bottom: 0.7rem;
}
label.field span {
    display: block;
    font-size: 0.74rem;
    color: var(--muted);
    margin-bottom: 0.2rem;
}
input[type="text"] {
    width: 100%;
    padding: 0.6rem 0.7rem;
    border: 1px solid var(--border);
    border-radius: 6px;
    font-family: inherit;
    font-size: 1rem;
    color: var(--ink);
    background: white;
    transition: border-color 0.15s, box-shadow 0.15s;
}
input[type="text"]:focus {
    outline: none;
    border-color: var(--accent);
    box-shadow: 0 0 0 3px var(--accent-soft);
}
button.publish {
    width: 100%;
    padding: 0.75rem 1rem;
    border: none;
    border-radius: 6px;
    background: var(--accent);
    color: white;
    font-family: inherit;
    font-size: 1rem;
    font-weight: 600;
}

```

```

        cursor: pointer;
        transition: background 0.15s;
        margin-top: 0.3rem;
    }
    button.publish:hover { background: #1d5959; }
    button.publish:active { background: #144545; }
    button.publish:disabled { background: #9aafad; cursor: not-allowed; }

    /* Subscribe feed */
    .feed {
        min-height: 220px;
        max-height: 420px;
        overflow-y: auto;
        display: flex;
        flex-direction: column;
        gap: 0.5rem;
    }
    .feed .empty {
        color: var(--muted);
        font-style: italic;
        font-size: 0.92rem;
        padding: 0.4rem 0;
    }
    .msg {
        padding: 0.55rem 0.7rem;
        background: var(--tile-bg);
        border-radius: 6px;
        font-size: 0.95rem;
    }
    .msg .meta {
        font-size: 0.72rem;
        color: var(--muted);
        margin-bottom: 0.15rem;
        display: flex;
        gap: 0.6rem;
        align-items: baseline;
    }
    .msg .name { font-weight: 600; color: var(--accent); }
    .msg .time { font-variant-numeric: tabular-nums; }
    .msg .body { color: var(--ink); white-space: pre-wrap; word-wrap: break-word; }

    /* Footer */
    footer {
        margin-top: 1.6rem;
        padding-top: 0.9rem;
        border-top: 1px solid var(--border);
        font-size: 0.74rem;
        color: var(--muted);
        display: flex;
        justify-content: space-between;
        gap: 1rem;
        flex-wrap: wrap;
        font-family: ui-monospace, 'SF Mono', Menlo, monospace;
    }
    footer a { color: var(--accent); text-decoration: none; }
    footer a:hover { text-decoration: underline; }
</style>
</head>
<body>
<div class="wrap">

    <header>
        <h1>Axona Demo</h1>
        <div class="status" id="status">...connecting</div>
    </header>

    <div class="dots" id="dots" aria-label="connected peers"></div>

```

```

<div class="topic-row">
  <div class="tile topic">
    <span class="k">Topic</span>
    <span class="v">hello-world</span>
  </div>
  <div class="tile prefix">
    <span class="k">Anchor</span>
    <span class="v" id="prefix">-</span>
  </div>
  <div class="tile topicid">
    <span class="k">Topic ID</span>
    <span class="v" id="topicId">-</span>
  </div>
</div>

<div class="panes">
  <div class="pane">
    <h2>Publish</h2>
    <label class="field">
      <span>Name</span>
      <input type="text" id="nameInput" value="Alice" autocomplete="off"
autocapitalize="words">
    </label>
    <label class="field">
      <span>Message</span>
      <input type="text" id="msgInput" value="Hello, world." autocomplete="
off">
    </label>
    <button class="publish" id="pubBtn" disabled>Publish</button>
  </div>

  <div class="pane">
    <h2>Subscribe</h2>
    <div class="feed" id="feed">
      <span class="empty">Waiting for ...messages</span>
    </div>
  </div>
</div>

<footer>
  <span>demo v<span id="demoVer">...</span> · kernel v<span id="kernelVer">...</
span></span>
  <a href="https://github.com/axona-net/axona-protocol">source</a>
</footer>

</div>

<script type="module">
import {
  AxonaPeer, AxonaDomain, NeuronNode,
  deriveIdentity, deriveTopicId, geoCellId,
} from '/src/index.js';
import { webTransport } from '/src/transport/web/index.js';

// Constants
// Bump the middle digit on every code change to this file so the
// footer reflects what's actually running and stale browser tabs are
// distinguishable.
const DEMO_VERSION = '1.6.0';
const BRIDGE_URL = new URLSearchParams(window.location.search).get('bridge')
  || 'wss://bridge.axona.net';
const VIRGINIA = { lat: 38.0, lng: -77.0 };
// axona-peer composes topic names as `${regionId}/${eventName}` (see
// axona-peer/src/client.js line 1230). Use the same composition here
// so the demo's hello-world topic at us-east is the EXACT topic that
// axona-peer subscribers on us-east/hello-world see. Without this,
// the 0x89 anchor matches but the body hash differs and the two apps

```

```

// never converge. EVENT is what we show in the UI tile; TOPIC is
// what crosses the wire.
const REGION_ID = 'us-east';
const EVENT = 'hello-world';
const TOPIC = `${REGION_ID}/${EVENT}`;
const MAX_DOTS = 100;

// DOM
const $dots = document.getElementById('dots');
const $status = document.getElementById('status');
const $prefix = document.getElementById('prefix');
const $topicId = document.getElementById('topicId');
const $kernel = document.getElementById('kernelVer');
const $name = document.getElementById('nameInput');
const $msg = document.getElementById('msgInput');
const $pubBtn = document.getElementById('pubBtn');
const $feed = document.getElementById('feed');

// Footer versions
document.getElementById('demoVer').textContent = DEMO_VERSION;
try {
  const pkg = await fetch('/package.json').then(r => r.json());
  $kernel.textContent = pkg.version;
} catch { $kernel.textContent = '?'; }

// Topic info (Virginia-anchored)
const virginiaPrefix = geoCellId(VIRGINIA.lat, VIRGINIA.lng, 8);
$prefix.textContent = '0x' + virginiaPrefix.toString(16).padStart(2, '0').
  toUpperCase();

// Synthetic publisher: nodeId-shaped hex whose top 8 bits = Virginia
// S2 cell. Anchors the topic at Virginia regardless of where the
// subscribing peer lives - every demo user worldwide computes the
// same topicId.
const publisher = virginiaPrefix.toString(16).padStart(2, '0') + '0'.repeat
  (64);
const topicId = await deriveTopicId(publisher, TOPIC);
$topicId.textContent = topicId.slice(0, 12) + ... + topicId.slice(-6);

// Geolocation (for our own identity only)
async function getLocation() {
  // Even if denied, identity falls back to (0, 0) + cell 0x10 (Gulf
  // of Guinea). We never end up with a 0x00 identity prefix unless
  // the user is at the South Pole. The topic anchor above is
  // independent of this - it always stays at Virginia (0x89).
  const FALLBACK = { lat: 0, lng: 0 };
  if (!navigator.geolocation) return FALLBACK;
  return new Promise(resolve => {
    navigator.geolocation.getCurrentPosition(
      pos => resolve({ lat: pos.coords.latitude, lng: pos.coords.longitude })
    ,
      () => resolve(FALLBACK),
      { enableHighAccuracy: false, timeout: 6000, maximumAge: 600_000 },
    );
  });
}
const geo = await getLocation();

// Identity + transport + peer
const identity = await deriveIdentity({ lat: geo.lat, lng: geo.lng });
const transport = webTransport({ bridgeUrl: BRIDGE_URL, identity });
const node = new NeuronNode({ id: BigInt('0x' + identity.id), lat: geo.lat
  , lng: geo.lng });
node.transport = transport;
const domain = new AxonaDomain({ k: 20 });
const peer = new AxonaPeer({ domain, node, identity, transport });

// Peer-state tracking for the dot strip

```

```

//
// We render one dot per known peer. Bridge is always the first
// (larger, accent-colored) dot, drawn whenever the bridge connection
// is alive. Each subsequent dot is a WebRTC peer.
//
// open    - green    (bound peer, isConnected = true)
// pending  - amber    (in synaptome but not yet bound - fresh signal)
// dead     - grey     (was bound, transport reported death)
//
// Cap at MAX_DOTS; if more peers, we just truncate (rare in a single
// browser tab - typically ~520 connections).
// Render the dot strip directly from transport.boundPeers() - the
// transport's own ground truth. This matches axona-peer's pattern
// (client.js:455 - for each render tick it diffs mesh.peers()
// against the existing rows and DOM-removes anything no longer
// present). There's no separate death-tracking map: a peer that
// dies just stops appearing. No grey-dot accumulation.
let bridgeUp = false;

// Transport callbacks deliver peerIds as BigInts (kernel canonical
// form). Normalise to lowercase 66-char hex for the dot's tooltip.
function asHex(peerId) {
  if (typeof peerId === 'bigint') return peerId.toString(16).padStart(66, '0')
  ;
  if (typeof peerId === 'string') return peerId.toLowerCase();
  return String(peerId);
}

// Map peerHex → its current dot DOM node, so onPingTraffic can flash
// the right dot without re-walking the children every event. Cleared
// on each renderDots(); rebuilt as we create children.
const dotByPeer = new Map();
let bridgeDot = null;

function renderDots() {
  $dots.innerHTML = '';
  dotByPeer.clear();
  // Bridge dot first
  bridgeDot = document.createElement('div');
  bridgeDot.className = 'dot bridge' + (bridgeUp ? ' open' : ' pending');
  bridgeDot.title = 'bridge';
  $dots.appendChild(bridgeDot);
  // Currently-bound peers (kernel reports only live channels)
  const peers = [...(transport.boundPeers?.() || [])].slice(0, MAX_DOTS - 1);
  for (const peerId of peers) {
    const hex = asHex(peerId);
    const d = document.createElement('div');
    d.className = 'dot open';
    d.title = hex.slice(0, 12) + ...;
    $dots.appendChild(d);
    dotByPeer.set(hex, d);
  }
}
renderDots();

// v2.0.2: pulse a dot on every ping-sent / pong-received frame on
// that peer's channel. Class is removed 400 ms later so a steady
// 1 Hz traffic looks like a soft, regular blink and a stalled
// channel goes dim within a second.
function pulse(dot) {
  if (!dot) return;
  dot.classList.add('pulse');
  // clearTimeout existing handle if any, to avoid premature class removal
  // when two events arrive in quick succession.
  if (dot._pulseTimer) clearTimeout(dot._pulseTimer);
  dot._pulseTimer = setTimeout(() => dot.classList.remove('pulse'), 400);
}

```

```

transport.onPingTraffic?.((peerId /*, kind */) => {
  const hex = asHex(peerId);
  // Bridge ping/pong attributes to the bridge's nodeId - the same
  // hex appears in boundPeers() as the bridge connection. If the
  // hex matches the bridge's nodeId, pulse the bridge dot; otherwise
  // look up the per-peer dot.
  const bridgeHex = transport.bridgeNodeId ? asHex(transport.bridgeNodeId) :
    null;
  if (bridgeHex && hex === bridgeHex) pulse(bridgeDot);
  else pulse(dotByPeer.get(hex));
});

// onPeerBound / onPeerDied are just render triggers - boundPeers()
// is the source of truth, so we re-read it on every event.
transport.onPeerBound?.(() => renderDots());
transport.onPeerDied?.(() => renderDots());

// Connect
//
// Two-phase: transport+peer start first (handshake with the bridge),
// then wait for the mesh to fill in before subscribing. Subscribing
// while synaptome.size === 1 (just the bridge) registers our
// subscription at the bridge only. When other peers come online and
// other publishers route their `pubsub:publish-k` through a fuller
// K-closest set, those publishes can miss us because our subscription
// is anchored at the wrong axon. Gating the subscribe behind a
// stable synaptome matches axona-peer's manual "Add subscription"
// pattern and what dht-sim does after warm-up.
//
// READY_SYNAPSE_COUNT = bridge + ~3 WebRTC peers; READY_TIMEOUT_MS
// puts a ceiling on solo-tab runs (when there's no one else to find).
const READY_SYNAPSE_COUNT = 4;
const READY_TIMEOUT_MS = 10_000;

async function waitForMeshReady() {
  const t0 = Date.now();
  while (Date.now() - t0 < READY_TIMEOUT_MS) {
    if (node.synaptome.size >= READY_SYNAPSE_COUNT) return node.synaptome.size;
    ;
    await new Promise(r => setTimeout(r, 200));
  }
  return node.synaptome.size;
}

try {
  await transport.start(identity.id);
  await peer.start();
  bridgeUp = true;
  $status.textContent = 'stabilising ...mesh';
  renderDots(); // bridge dot + any peers already bound

  await waitForMeshReady();
  // Subscription happens just below (top-level await) - gate the
  // button on success after sub resolves.
} catch (err) {
  $status.textContent = 'connection failed';
  console.error('[demo] connect failed', err);
}

// Subscribe
let messageCount = 0;
function appendMessage(name, body, when) {
  if (messageCount === 0) $feed.innerHTML = '';
  messageCount++;
  const msg = document.createElement('div');
  msg.className = 'msg';
  const meta = document.createElement('div');
  meta.className = 'meta';

```

```

const nameEl = document.createElement('span');
nameEl.className = 'name';
nameEl.textContent = name;
const timeEl = document.createElement('span');
timeEl.className = 'time';
timeEl.textContent = when.toLocaleTimeString([], { hour: '2-digit', minute:
  '2-digit' });
meta.appendChild(nameEl);
meta.appendChild(timeEl);
const bodyEl = document.createElement('div');
bodyEl.className = 'body';
bodyEl.textContent = body;
msg.appendChild(meta);
msg.appendChild(bodyEl);
$feed.appendChild(msg);
// Auto-scroll to the newest message
$feed.scrollTop = $feed.scrollHeight;
}

const sub = await peer.sub(TOPIC, (envelope) => {
  // Envelopes from the wire are "Name: Body" strings.
  const wire = envelope.message || '';
  const colonAt = wire.indexOf(':');
  const name = colonAt > 0 ? wire.slice(0, colonAt).trim() : '(unknown)';
  const body = colonAt > 0 ? wire.slice(colonAt + 1).trim() : wire;
  appendMessage(name, body, new Date());
}, { publisher, since: 'all' });

// Arm the kernel's AxonaManager periodic refresh timer. By default
// AxonaPeer._buildDefaultAxonaManager() does NOT call am.start() (the
// comment in AxonaPeer.js:1820 says "applications call peer.sub after
// the mesh has stabilised, so the initial K-closest is already wide
// and refresh isn't needed"). That assumption holds for the dht-sim
// simulator's steady-state runs but NOT for a browser mesh, where
// WebRTC peers come and go on every page reload. Without this, a
// subscription registered at the initial K-closest set becomes stale
// as the mesh churns, and publishes routed through a new K-closest
// set miss the stale subscribers. Arming refreshTick at the 10 s
// default keeps the subscription pinned to whoever is K-closest
// right now.
peer._axonaManager?.start?.();

// Backstop re-render every 2 s. onPeerBound / onPeerDied catch
// most transitions, but some peers die silently (network drop, page
// suspended, browser throttled tab). A periodic re-read of
// transport.boundPeers() keeps the strip honest without a
// state-tracking map.
setInterval(renderDots, 2000);

// Resume recovery (laptop sleep / tab background)
// The kernel transport already self-heals after a connection freeze:
// dead WebRTC channels are evicted on heartbeat timeout (mesh v2.1.1)
// and the bridge socket auto-reconnects with backoff. This is the
// optional *application-layer* accelerator the peer at axona.net also
// uses - on a genuine resume event, tear the mesh down at once
// (mesh.reset() fires onPeerLost so the synaptome clears too) and force
// an immediate bridge reconnect, so recovery is snappy rather than
// waiting out the per-peer timeouts. A macOS screensaver may fire NONE
// of these events (the tab never goes hidden); that's exactly the case
// the kernel-side eviction backstops.
function onResume(reason) {
  console.log('[demo] resume:', reason);
  try { transport.mesh?.reset?.(); } catch { /* mesh may already be gone */ }
  transport.reconnectNow?.();
}
let hiddenAt = 0;
document.addEventListener('visibilitychange', () => {
  if (document.visibilityState === 'hidden') { hiddenAt = Date.now(); return;

```

```

    }
    if (hiddenAt && Date.now() - hiddenAt >= 5000) onResume('visible');
    hiddenAt = 0;
  });
window.addEventListener('online', () => onResume('online'));
window.addEventListener('pageshow', (e) => { if (e.persisted) onResume('bfcache'); });

// Subscription is registered - safe to publish now.
$status.textContent = 'connected';
$pubBtn.disabled = false;

// Publish
async function publish() {
  const name = ($name.value || 'Alice').trim();
  const msg = ($msg.value || 'Hello, world.').trim();
  if (!msg) return;
  $pubBtn.disabled = true;
  try {
    await peer.pub(TOPIC, name + ': ' + msg, { publisher });
  } catch (err) {
    console.error('[demo] publish failed', err);
  } finally {
    $pubBtn.disabled = false;
    $msg.focus();
    $msg.select();
  }
}
$pubBtn.addEventListener('click', publish);
$msg.addEventListener('keydown', (e) => {
  if (e.key === 'Enter' && !$pubBtn.disabled) publish();
});
</script>

</body>
</html>

```