

AXONA

Architecture

David A. Smith
Axona.net

An Axona Architecture Note v0.8.6 2026-06-22

A protocol is the contract between layers that don't trust each other.

What I cannot create, I do not understand.

RICHARD FEYNMAN · 1988

I The Stack

AXONA SHIPS AS FOUR REPOSITORIES that compose into one running system. Each layer has a single responsibility,¹ a single canonical version identifier (`KERNEL_VERSION`), and a contract surface narrow enough to fit on a postcard.

@axona/protocol (kernel). The pure-JS library. Owns the DHT routing, the pub/sub overlay, the cryptography, and the transport contracts. No browser DOM, no Node-specific I/O, no network code outside abstract transport classes. Runs identically in browsers, in Node, in dht-sim's simulated network.

axona-peer (browser application). The reference end-user app. Provides identity persistence, region selection, WebRTC mesh formation, a chat-style UI for pub/sub. Imports the kernel as a vendored copy under `vendor/axona-protocol/` so the version that ran when the page was built is the version that runs in the browser.

axona-bridge (deployed server). The signaling + introducer service at `bridge.axona.net`. Runs an embedded `AxonaPeer` of its own (so it's a full DHT participant, not just a relay), gates connecting peers behind a version handshake, mints short-lived TURN credentials, relays WebRTC signaling between browsers, and forwards Axona wire frames addressed to its own embedded peer.

axona-relay (headless supernode). A Node process that joins the mesh over *real* WebRTC (a spec `RTCPeerConnection` from `node-datachannel`) and runs the same kernel stack a browser peer runs — routing lookups, rooting pub/sub topics, relaying signaling for others — but with no DOM, no public server, and no TURN minting. It is a strict *subset* of the bridge: it dials a bridge once for bootstrap, then is a first-class peer. Running several relays strengthens routing and pub/sub durability without any of them being an inbound server.

Each component carries its own semver track because each ships to a different audience: kernel changes need API-level discipline; peer changes are user-visible UI; bridge and relay changes are operational deploys. But they all share the same kernel version, visible at runtime through `/healthz` (bridge), the `welcome` frame (any peer connecting), and the version row in the peer UI.

¹ The pattern is intentional. A peer's business logic should never need to know the wire format; a bridge should never need to know what topic a subscriber cares about; the kernel should never need to know which application is using it. Every interface in this document is a place where one layer was given the right to be wrong about the layer below.

A fifth repository, `dht-sim`, runs the kernel against a simulated network for protocol research and regression benchmarks. It's not part of the production stack but shares the same `@axona/protocol` library; every change that ships to production also has to pass the simulator's 25,000-node battery.




Figure 1: *

Five layers in one peer. Every arrow is a contract; every contract has a smoke-test suite.

II Identity and Addressing

AXONA KEEPS THREE CONCERNS SEPARATE, each with its own primitive. *Connection* is a node identity — an Ed25519 keypair whose public key yields a 264-bit *Node ID*; it is ephemeral and anchored to a region. *Authorship* is an author identity — a location-free Ed25519 keypair whose public key *is* the *Author ID* (the `signerPubkey` on the wire); it is durable and a peer may hold several. *Addressing* is the topic descriptor and message ID, derived independently of either identity. The kernel mints the two identities with separate factories (`createNodeIdentity`, `createAuthorIdentity`); a publish is signed per-call with an author, never with the connection key.

THE CONNECTION IDENTITY IS A 264-BIT IDENTIFIER. The top 8 bits are a Google S2 cell prefix² computed from the peer’s latitude and longitude; the bottom 256 bits are the SHA-256 of the peer’s Ed25519 public key. A peer in London has an ID starting with `0x44`; a peer in Virginia starts with `0x89`. Two peers in the same region share a leading byte; two random peers share only the bits their keys happen to share.

The address space is treated as `BigInt` in every line of running code. The wire format is 66-character lowercase hex — the *ONLY* representation that ever appears in JSON. The conversion happens at exactly two places: when the transport serialises an outgoing frame, and when the dispatcher parses an incoming one.

An author identity carries no location. Where the node ID binds a key to a place, the Author ID is the bare Ed25519 public key

² S2 projects the sphere onto the six faces of a cube and partitions each face with an equal-area quadratic transform. We bin each face into a 4×8 grid (32 cells per face) — “S2 level 2.5,” one bit short of level 3’s 8×8 — giving $6 \times 32 = 192$ cells, which fit in 8 bits with byte values 192–255 left reserved for future system use. The quadratic projection keeps cell areas within $\sim 1.7\times$ of each other globally — no polar shrinkage — and any standard S2 library can recompute our prefixes.

Each of the 192 cells carries exactly *one* human-readable name, so a region always presents the same label (no location-dependent flip-flop). Where a cell straddles an ocean and a landmass the land name wins; a multi-country cell takes its dominant city; homogeneous cells (one country, open ocean) keep their name. A name is usually unique to one code, though an area larger than one cell may span adjacent codes. The kernel ships the canonical table and the `regionName/regionCode/resolveRegion` helpers; the `s2-region-visualizer`

and nothing else — no region byte, no coordinates. It is what a reader recognises across sessions and devices; the same author key signs from anywhere. Addressing never depends on it being co-located with the publisher.

Topic identifiers share the address space. A topic is not a bare string but a structured descriptor — `{region, owner?, name, write}` — and its ID is `topicId = regionByte || sha256(canonical({owner, name, write}))`, so the top byte is always a *real*, routable S2 cell. The region comes from one of two rules: an explicit `region` places the topic where the application chooses (given as a region *name or its code* — `'useast' ≡ 0x89 ≡ 137`, all normalising to the same region byte); if `region` is omitted, it defaults to the *publisher's own node region* — the top byte of its node ID, which is a real cell the publisher already occupies. The region is *never* derived from the author key: an Author ID has no region, and hashing one into a region would dump every author's topics into a single arbitrary cell, manufacturing a hot spot in whatever populated region happened to fall closest. There is no synthetic publisher and no 00 “global” bucket: every topic lands in a populated region, and peers in that region are XOR-closest to it, which is what makes regional pub/sub route naturally.³

The Activation Potential (AP) routing formula, the K-closest selector, and the synaptome's stratum buckets all operate on this one space. There is no separate “topic ID space” or “peer ID space”; they share addresses because they share routing.

III The Kernel: @axona/protocol

THE KERNEL IS THE PROTOCOL. Everything else is its clients. The directory structure mirrors the contract surface:

src/contracts/. Abstract base classes: `Transport`, `DHT`, `DHTNode`, `BootstrapService`, `PersistenceAdapter`. These are the seams where alternate implementations plug in. A custom transport (libp2p, QUIC, embedded) extends `Transport`; a database-backed persistence (IndexedDB, SQLite) extends `PersistenceAdapter`.

src/dht/. The active routing layer: `NeuronNode`, `Synapse`, `AxonaPeer`, `AxonaDomain`, `Subscription`. This is where the Axona routing logic lives, and where the developer-facing API surface (`peer.pub`, `peer.sub`, `peer.lookup`) is declared.

src/pubsub/. The pub/sub overlay: `AxonaManager`, `post.js` (the topic descriptor + `resolveTopic` / `deriveTopicId` derivation), `envelope.js` (signed publish envelopes + write-policy checks),

³ Because the omitted-region default is the *publisher's* region, two co-located peers converge on the same topic id without naming a region; cross-region discovery names the region explicitly (typically alongside the owner and name). The `write` field also folds into the hash: `open` lets anyone publish (self-signed), `owner` restricts publishing to the named owner Author ID. Because `owner` and `write` live inside the signed descriptor, a host recomputes the topic ID and enforces write authorisation statelessly — see § XII.

`ed25519.js` (Ed25519 binding for `crypto.subtle`). `AxonaManager` owns the K-closest tree, the replay caches, the dedup LRU, the metrics counters.

src/transport/. Concrete transports: `transport/web/` (browser: BridgeTransport over WebSocket + WebRTCTransport over DataChannels), `transport/node/` (server-side WebSocket), `transport/sim/` (in-process simulation for dht-sim and tests). All implement the abstract `Transport` contract identically.

src/utils/. Pure helpers: `hexid.js` (`BigInt`↔hex conversions, XOR distance, leading-zero count), `s2.js` (S2 cell computation), JSON codec (`transport/wire.js`).

src/identity/. Ed25519 keypair derivation + persistence envelopes. Two factories, one per concern. `createNodeIdentity({lat, lng})` mints the *connection* identity: an ephemeral, location-anchored keypair returning `{id, pubkey, pubkeyHex, privateKey, region, ...}`, where `id` is the 264-bit Node ID. `createAuthorIdentity({persistAs?, store?})` mints the *authorship* identity: a location-free keypair returning `{authorId, pubkey, pubkeyHex, privateKey, ...}` with no `id` and no `region` — `authorId` is the public key, i.e. the `signerPubkey` on the wire. Persistence is an option on the author factory, not a separate call.

The whole kernel has zero runtime dependencies beyond the host JavaScript environment (browser or Node 20+). No bundler, no transpilation, no build step. The same source files run in the browser via ES modules, in Node via direct `import`, and in the simulator via the in-process transport.

IV The Local Node

EACH PEER’S VIEW OF THE NETWORK lives in two data structures on the `NeuronNode`: the *synaptome* (its outgoing connections) and the *incoming synapse index* (peers known to point at it). Both are keyed by 264-bit `BigInt` `nodeId`.

A `Synapse` is a routing edge — it carries a peer ID, a measured latency, a learned weight, a stratum (the XOR-distance bucket), an inertia lock, and provenance (was this synapse added via handshake, hop-cache, triadic introduction, or sponsor bootstrap?). The synaptome is unbounded by design; the routing layer evicts synapses through the *Structural Survival Rule*, which guarantees navigability even under aggressive weight decay.

The rule, in one sentence: a synapse may only be pruned if at least one other synapse with the same stratum exists. The effect: every populated stratum keeps at least one live route, so the address space stays globally navigable even after heavy decay.

Routing decisions don't pick the XOR-closest synapse blindly. Each candidate is scored by Activation Potential:

$$AP_c = \frac{\Delta \text{Distance}_c}{L_c} \times (1 + w \cdot W_c)$$

The first factor is progress velocity — how much XOR distance the hop closes per millisecond of measured latency. The second is a mild boost for synapses that have historically led to fast lookups. Weight only accrues on at-or-below-average-latency paths, so a high- W synapse genuinely represents a fast route, not just a frequently-used one. Two random hops with similar latency tie; the one that's served us well before wins the tiebreak.

V The Protocol Layer

AXONAPEER IS THE ACTIVE OBJECT that turns a **NeuronNode** into a network participant. It owns:

- The handler installation. On `peer.start()` it registers the Axona routing handlers (`lookup_step`, `lookahead_probe`, `find_closest_set`, `route_msg`, `local_probe`, `triadic_introduce`, `hop_cache`, `lateral_spread`, `reinforce`) on the transport.
- The auto-admission loop. It asks `transport.boundPeers()` for currently-bound nodeIds and seeds them into the synaptome; it subscribes to `transport.onPeerBound` for future admissions.
- The pub/sub surface. `peer.pub`, `peer.sub`, `peer.pull`, `peer.metrics`: the developer-facing API. Each delegates to the **AxonaManager** after deriving the topic ID and (for pub) building a signed envelope.
- The direct-message surface. `peer.send`, `peer.notify`, `peer.onMessage`: per-peer request/response and fire-and-forget, outside the pub/sub overlay. For protocols that aren't topic-shaped.
- The lookup primitive. `peer.lookup(targetKey)` runs an iterative K-closest walk and returns `{found, path, hops, peerId, latency}`. This is what **AxonaManager** uses internally on every `pub/sub` to find the topic's axons.

AxonaManager is a separate object owned by the peer. Its job is the pub/sub overlay: when a subscriber sends `pubsub:subscribe-k` to the K-closest axons for a topic, the receiving **AxonaManager** stores the subscriber in its `role.children` for that topic. When a publisher fans `pubsub:publish-k` to its own K-closest, each receiving **AxonaManager** iterates its `role.children` and forwards `pubsub:deliver` to each subscriber. The publisher's signed envelope flows through unchanged; per-hop dedup prevents the fan-out tree from delivering duplicates.

This section is the architectural placement of **AxonaPeer**; the field-by-field surface — every method, its parameters, and its return shape — is catalogued in the Application API section later in this document.

The K-closest replication is what makes the system tolerant. A publisher sends to *five* root axons in parallel; a subscriber subscribes at *five* root axons in parallel. As long as one axon in each set overlaps, delivery succeeds. With $K = 5$ and synaptomes that converge under the bridge’s peer-list dissemination, overlap is typically 4 or 5 of 5 — so a publish that loses two axons to transient WebRTC failures still delivers cleanly.

Three convergence mechanisms close the residual gaps where “one axon overlaps” is not guaranteed — a publish lands on the *publisher’s* K-closest set, which need not be a given subscriber’s. **Gap-safe replay** (kernel v2.37.0): a (re)subscriber reports the content-IDs it actually holds (a bounded **have** digest) and a root replays exactly the complement, so a hole in the middle is backfilled rather than masked by a single high-water timestamp. **Root-to-root anti-entropy** (v2.39.0): each root periodically exchanges digests of held content-IDs with its K-closest siblings and pulls what it is missing — and *re-verifies* every pulled message (publisher signature B-4 + content-address), so a sibling root cannot inject a forged or content-poisoned message and a killed message is never resurrected. **Kill convergence** (v2.36.0): recently-applied retractions are re-gossiped to the current root set, so a replica that missed the original kill still drops and tombstones the message rather than re-serving it to a reloading subscriber. All three are wire-additive — old and new nodes interoperate.

VI The Transport Layer

TRANSPORT IS WHERE ADDRESSES MEET SOCKETS. The kernel’s **Transport** contract is a small 12-method surface: **start**, **stop**, **getLocalNodeId**, **openConnection**, **closeConnection**, **isConnected**, **send** (request/response), **notify** (fire-and-forget), **onRequest**, **onNotification**, **onPeerDied**, **getLatency**. Every concrete transport implements this surface and the upper layers don’t care which one is wired in. (Peer *binding* — the authenticated **bindPeer**/**onPeerBound** admission flow — lives in the concrete web/node transports on top of this contract, not in the contract itself.)

Three concrete implementations ship:

transport/web/. The browser stack. A **CompositeTransport** fans out to two sub-transports based on which one owns each `nodeId`: a **BridgeTransport** that carries Axona wire frames over the same `WebSocket` the bridge uses for signaling, and a **WebRTCTransport** over a **MeshManager** that owns the data-channel pool. The factory `webTransport({bridgeUrl, identity})` composes the two, drives the bridge’s version gate, runs the hello/hello-ack admission,

threads TURN credentials from the welcome frame into the mesh's ICE config, and resolves once everything is up.

The factory also owns the *connection lifecycle*, so an application doesn't hand-roll it: a 1 Hz ping/pong heartbeat with stale detection, auto-reconnect with exponential backoff on an unexpected socket close (re-running the version gate and hello/hello-ack on each re-open), and a small observability surface — a **bridgeState** machine (**connecting** → **open** → **stale** → **disconnected**, plus **upgrade-required** on a 4426 close), the captured **welcome** (**bridgeInfo**), and the live ping RTT (**bridgeRtt**). A 4426 version-gate rejection suppresses reconnect; **reconnectNow()** forces an immediate re-attempt (used on tab-resume and network-online).

The WebRTC mesh under **webTransport** self-heals on the same principle. **MeshManager** runs a 1 Hz ping/pong on every data channel and (kernel v2.1.1) evicts a peer the moment its channel goes dead: no pong for the heartbeat-timeout window, or a few consecutive throwing **send()**s — the signature of a channel a browser keeps reporting **open** after a laptop sleep, where the socket is gone but **readyState** lies. Eviction fires **onPeerLost** → **onPeerDied**, so the protocol layer drops the synapse and routes around the dead peer while the bridge's next peer-list rebuilds a fresh channel. This is the Transport contract's heartbeat-timeout obligation, and it means recovery from a connection freeze needs *zero* application code. An application *may* add an accelerator: on a resume event (**visibilitychange**, **online**, **bfcache** **pageshow**) call **transport.mesh.reset()** + **reconnectNow()** to clear stale peers at once rather than waiting out the per-peer timeouts — both the demo and the axona.net peer do this. A macOS screensaver notably fires none of those resume events (the tab never goes hidden), which is why the transport-level eviction is the load-bearing mechanism and the app-level reset is only a speed-up.

transport/node/. A Node-side WebSocket transport. Used by the bridge's embedded peer to talk to connected browsers without simulating WebRTC — the bridge *is* the signaling server, so it doesn't need WebRTC to reach any browser; the WebSocket is direct.

transport/sim/. The simulator transport. Used by dht-sim's **TransportAxonaEngine** and by the kernel's smoke tests. Identical Transport contract, but messages are scheduled in the event loop with simulated latency derived from a real internet topology dataset.

The wire format is JSON, with a single quirk: **BigInt** is serialised as "**<digits>n**" (decimal digits with an **n** suffix) and revived back to

The same kernel test suite runs against all three transports. Cross-transport bugs — which is to say, the kind of bug where a fix lands in one transport's call path but not the others — can't ship.

BigInt on decode. NodeId fields are always hex strings on the wire — they never appear as bigint-suffixed literals because the kernel converts them at frame construction.

VII The Bridge

THE BRIDGE IS AN INTRODUCER, not a server. Its job is to let browsers find each other; once they’ve found each other, WebRTC data channels carry every byte of application traffic peer-to-peer. The bridge stays in the picture only for two reasons: (1) it’s the rendezvous for the version handshake before peers can talk at all; (2) it’s a guaranteed-reachable mesh peer, which makes it a useful K-closest axon in small networks where two browsers might otherwise fail to find common axons.

The bridge handles three kinds of frames over the WebSocket:

Signaling. `welcome`, `peer-list`, `peer-joined`, `peer-left`, `signal`, `ping/pong`. Routes WebRTC SDP offers/answers and ICE candidates between browsers, broadcasts join/leave events, supplies short-lived TURN credentials in the welcome frame.

Version handshake. `client-hello`/`server-hello`. Enforces a minimum peer version. Rejects sub-minimum peers with WebSocket close code 4426 and a human-readable reason; the peer’s UI shows a “please reload” banner.

Axona wire. `{type: 'axona', payload: ...}` frames carrying `req/res/ntf` messages between any connected browser and the bridge’s embedded peer. These are the same frames any two browsers exchange over WebRTC — the bridge embeds a full **AxonaPeer** of its own (**BridgeAxonaNode**) and participates in routing just like any other node, which is why it can be in K-closest sets and why it can hold subscriber children for pub/sub topics.

The bridge’s embedded peer runs an *ephemeral* identity — a fresh Ed25519 keypair minted on every start, never written to disk, so there is no keypair file to migrate across deploys.⁴ It still lands in us-east — the 0x89 S2 cell, from its default us-east coordinates — which is why it sits XOR-close to `us-east/*` topics and is so often a K-closest root; only the low 256 bits rotate per restart. Its `health` surface exposes the kernel version it’s running and the synaptome it has built up, available at `/healthz` JSON. Operators check that endpoint to confirm a deploy.

The kernel builds a peer’s **AxonaManager** lazily, on its first `pub/sub`. The bridge never publishes or subscribes — it only relays — so it must build that engine *eagerly* at startup, otherwise the `pubsub:subscribe-k/publish-k` handlers are never installed and every `pub/sub` frame addressed to the bridge is silently dropped. Because the bridge sits in nearly every peer’s K-closest set, that failure mode is a near-invisible delivery hole; the eager build is what keeps it an honest relay axon.

⁴Since the v3.0 identity model the bridge `nodeId` is disposable: it authenticates the connection, not authorship, and nothing depends on it persisting. A restart rotates the `nodeId` and the mesh simply re-binds.

How a New Node Joins

JOINING IS A MUTUAL, AUTHENTICATED BIND. A fresh node has no peers; it reaches the mesh through a bridge and then fans out to direct WebRTC links. Crucially, synaptome membership is established the *same* way on every link — bridge or mesh — so “joining” is just the first instance of the generic admit path, not a special case.

1. **Connect + version gate.** The node opens a WebSocket to a bridge and exchanges `client-hello/server-hello`. The bridge checks the wire major (`REQUIRED_WIRE_MAJOR`, currently 3 — the v0.3 descriptor-envelope wire). A peer that sends a different wire major — or, being pre-flag-day, omits `wireVersion` entirely — belongs to the old network and is closed with 4426 before it can do anything else. This wire-major check is the load-bearing, hermetic partition gate: a wire-2 (descriptor-less) envelope and a wire-3 envelope are mutually unverifiable, so the bridge refuses the mismatch up front rather than admitting a peer that then fails silently at the envelope layer. The older pernamespace version floor (`flagDayFloor`, with `MIN_KERNEL_VERSION` and `MIN_PEER_APP_VERSION` both 3.0.0) still runs but no longer discriminates — now that the kernel itself is on the 3.x line both namespaces collapse into one bucket, so the floor only surfaces the reported semver in the close reason. Beneath the wire gate the `axona/5` auth epoch is the hermetic backstop from the earlier 2026-06 partition (it did *not* change in v0.3): the epoch is folded into the signed transcript of the next step, so an `axona/4` node and an `axona/5` node still *cannot* complete a handshake.
2. **Welcome.** The bridge replies with `welcome`, carrying the connection id and a fresh `serverNonce`; both sides now hold the perconnection channel-binding value (CBV) the next step signs over.
3. **Authenticated handshake (axona/5).** Each side sends a `hello` (the responder a `hello-ack`) carrying its Ed25519 public key and a signature over the link’s CBV, and verifies the other: the presented pubkey must hash to the 256-bit suffix of the `nodeId` it claims, and the signature must check out. A peer that cannot prove its id — or one being MITM’d, since the CBV binds the proof to *this* channel — fails here and is closed.
4. **Mutual bind + admit.** On success both ends record the binding and admit the other into their synaptome *as an ordinary Synapse*: the joining node calls `bindPeer(bridgeId, 'bridge')`, whose `onPeerBound` hook seeds the bridge into its synaptome (`addedBy: bootstrap`); the bridge’s embedded peer runs `_completeHandshake`,

binds the new node to that connection, and admits it via `_addByVitality` (`addedBy: handshake`). No bridge-specific code runs in either step — the bridge is simply the first guaranteed-reachable peer the node binds.

5. **Fan out to the mesh.** The bridge broadcasts `peer-joined` and hands the new node a `peer-list`. The node negotiates WebRTC data channels to those peers (SDP relayed via the bridge), and each channel runs the *same* `axona/5` hello — now additionally bound to the peer’s DTLS-certificate fingerprint — firing the same `onPeerBound` admit. The synaptome fills with direct mesh peers; the bridge stops being the only neighbour.
6. **Converge.** Optionally `peer.join(sponsor)` seeds a specific first neighbour; thereafter vitality (`weight × recency`) and the 10s `refreshTick` grow and prune the synaptome and re-anchor subscriptions onto the converged K-closest set. The node is now a full participant.

VIII The Wire Protocol

EVERY FRAME ON THE WIRE is JSON. Every `nodeId` or topic ID field is a 66-character lowercase hex string. There’s no binary framing, no protobuf, no length-prefixed records — the discipline of having one wire format that any of the three transports can carry is worth the few percent of size you give up.

Three kinds of message-typed envelopes flow over the wire:

Request (*k*: `'req'`). A correlated request expecting a response.

Carries an integer `id` for correlation, a `type` string naming the handler (`lookup_step`, `route_msg`, `find_closest_set`, ...), and an opaque `body`.

Response (*k*: `'res'`). The reply to a `req`. Echoes the `id` and adds an `ok` boolean and a `body`.

Notification (*k*: `'ntf'`). Fire-and-forget. No `id`, no response. Used for hello/hello-ack admission, `pubsub:subscribe-k`, `pubsub:publish-k`, `pubsub:deliver`, `triadic_introduce`, `hop_cache`, `lateral_spread`, `reinforce`.

The handler types are versioned implicitly through the kernel — the bridge’s `minPeerVersion` gate excludes peers running an older protocol that wouldn’t understand newer messages. Major protocol changes bump `KERNEL_VERSION`’s minor; minor changes are additive (a peer that doesn’t know a handler type silently drops the message).

A pub/sub message envelope is signed:

The admit is symmetric and identity-bound in both directions — each side admits the other’s *authenticated* `nodeId`, so a peer can only ever be admitted under an `id` it proved it owns. Because every node runs step 4 against the bridge first, the bridge lands in every synaptome: its centrality is *emergent* from this bootstrap, not granted by any special rule (see §Future Directions, “demote the bridge to an ordinary peer”).

```

{
  msgId:      "<sha256 of canonical({publisher, message}) hex>",
  seq:        1716638400123,    // per-publisher monotonic (v2.9.0)
  ts:         1716638400000,
  topic:      {                  // v0.3: signed topic DESCRIPTOR,
    region: "useast",            // not a bare name. The signature
    owner:  "<owner pubkey|null>", // binds the whole descriptor,
    name:   "hello-world",       // incl. the write policy.
    write:  "open"                // "open" | "owner"
  },
  message:    <application payload>,
  signerPubkey: "<Ed25519 pubkey hex>",
  signature:   "ed25519:<128 hex chars>"
}

```

The `msgId` is the SHA-256 of the canonical `{publisher, message}` core (`publisher` being the `signerPubkey`, or `null` when anonymous) — a stable content address of author-plus-payload that deliberately excludes `ts`, `seq`, `topic`, and `signature`, so a re-publish of identical content by the same author dedups to one ID.

Receivers verify the signature against the signer’s pubkey using `crypto.subtle.verify`; the kernel surfaces verification result on the delivered envelope so applications can decide whether to trust unsigned or unverifiable publishes.

As of kernel v2.9.0 (envelope format **v2**) the signature covers a domain-separated core (`axona:pubsub-envelope:v2`, finding E-4) and includes a per-publisher monotonic `seq` with TTL freshness (finding C-2). Two properties follow. *Freshness*: a topic’s root axons reject a *live* publish whose signed `ts` is outside a ± 5 -minute window, so a captured envelope can’t be re-injected as new once it ages out (the legitimate replay-to-late-subscribers path is exempt — it serves cached history). *Order*: `seq` gives each publisher’s stream a total order, which is what makes “latest” and bounded-queue eviction (below) deterministic across replicas.

As of kernel v3.0.0 (envelope wire **v3**, the v0.3 identity/authorship flag-day) the signed `topic` is a *descriptor* `{region, owner, name, write}` rather than a bare name, and the signature covers the whole descriptor. This buys two ingress-time guarantees a root enforces against the topic it actually routed the message to. *Binding*: a root recomputes the topic ID from the signed descriptor and drops the message unless it equals the routed topic — a message signed for topic A cannot be injected onto topic B. *Write policy*: for a `write: 'owner'` topic the `signerPubkey` must equal the descriptor’s `owner`, or the publish is rejected (`WRITE_POLICY_VIOLATION`); only the owner key

may publish to an owned feed. Both checks run at *every* ingress path — direct publish, K-closest relay-accept, and anti-entropy sync — so a forged or mis-targeted envelope is dropped wherever it enters, and a pre-v3 bare-string topic is dropped outright: that is the flag-day boundary the wire-major-3 gate enforces up front.

IX Routing

AXONA’S ROUTING IS LEARNING-ADAPTIVE. It replaces Kademlia’s k -bucket array with a dynamic synaptome that adjusts edge weights based on observed latency. The basic shape — iterative lookup, $\log_2(N)$ hop count — is unchanged from Kademlia; what changes is which peer you ask at each hop.

The five wire operations that drive Axona routing are:

lookup_step (*req*). “Here’s a target. What’s your best next hop?”

The receiver evaluates AP across its synaptome’s progress candidates, returns the winning synapse + its own routing state.

lookahead_probe (*req*). “If I forward to you for target T , who would you forward to?” A one-deep peek that lets the caller compare expected hop costs across candidates before committing.

find_closest_set (*req*). “Give me your K peers closest to this target.” The K-closest harvester used by AxonaManager’s pub/sub axon discovery.

route_msg (*req*). Generic routed-message forwarder. Used by AxonaManager for `__tunneled_direct__` delivery when the target isn’t directly bound.

reinforce (*ntf*). “This lookup succeeded; bump the weight on the synapse I used.” Fire-and-forget LTP signal back along the path. Only flows when the lookup hit at-or-below average latency for the path’s region. On identity-binding transports the handler is a no-op unless the *target* synapse’s peer is currently bound (the B-3 gate: an unauthenticated **reinforce** must not be able to refresh the eviction-inertia of a stale or unverified entry and pin it in the table). Weight is independently clamped to ≤ 1.0 , so inertia is the only lever the gate needs to close.

Two structural-plasticity messages let the network rewire itself without explicit instruction: **triadic_introduce** (“introduce me to your friend at X ”) and **hop_cache** (“I’ve forwarded to X several times; install a direct synapse to X ”). Both are fire-and-forget notifications that obey the synaptome’s admission rule (`_addByVitality`).

Residual, by design: the gate keys on the *target* synapse being bound, not on the *sender*’s identity — so any bound peer can refresh the inertia of any other bound peer’s synapse. The blast radius is bounded to “keep an already-verified, live peer pinned slightly longer,” which is benign; binding the signal to the sender would add cost for no attacker-relevant gain. Revisit only if reinforce ever carries weight a sender could bias.

The performance claim of Axona over Kademlia is roughly 40% fewer wall-clock milliseconds per lookup at 25,000 nodes, mostly from latency-weighted hop selection. The figure is documented in the companion paper *Axona: A Learning-Adaptive Distributed Hash Table and Axonal Publish/Subscribe* (2021-08-26, 2022-07-27).

X Pub/Sub: K-Closest

THE PUB/SUB OVERLAY IS ITS OWN STRUCTURE. It sits on top of Axona routing but doesn't reuse the synaptome for state — each peer maintains a separate set of *axon roles*, one per topic it's chosen as an axon for. A role holds a topic ID, a set of children (subscribers), a replay cache, and a K -replicated peer set (the other root axons for the same topic).⁵

Subscriber-side flow:

1. Compute `topicId` from the topic *descriptor*: an 8-bit region prefix (the descriptor's *region*, or the publisher's own node region when region is omitted) followed by `SHA256` of the canonical `{owner, name, write}`. No synthetic publisher is involved.
2. Find the K peers closest to `topicId` in the local synaptome. (No network walk — the local reachable peer set is what matters; remote-walk K -closest returns ghost IDs that route fan-out into dead ends.)
3. Send `pubsub:subscribe-k` to each, carrying the subscriber's `nodeId` and an optional `lastSeenTs` for replay.
4. Each receiver runs `_addOrRecruitChild`, stores the subscriber under `role.children`, and optionally serves a `pubsub:replay-batch` of cached messages newer than `lastSeenTs`.

Publisher-side flow:

1. Compute the same `topicId`.
2. Build a signed envelope.
3. Find the same K peers closest to `topicId` locally.
4. Send `pubsub:publish-k` to each.
5. Each receiver runs `_onPublishDirect`: dedup against `_seenPublishes`, add to replay cache, fan `pubsub:deliver` to every child in `role.children`.
6. Each child receives `pubsub:deliver`, dedups, fires the local `_deliveryCallback` which routes to the application's subscription handler.

When a publisher's K -closest peers also happen to be in the subscribers' K -closest sets, the tree converges and delivery is single-hop from publisher to axon to subscriber. When they don't, the system relies on the redundancy: a missing axon's subtree is covered by the four other axons. The Activation Potential weighting makes high-quality axons dominate the K -closest set over time.

⁵ Notation: this replication factor is the *root-set size* $R = 5$ (`AxonaManager.rootSetSize`)—distinct from the Kademlia *routing* parameter $k = 20$ (`AxonaDomain._k`), which sizes the closest-set used for *lookups*. “ K -closest” throughout names the generic find-closest- N operation both use: $N = k$ for routing, $N = R$ for replication.

Per-hop `_alreadySeenPublish` dedup is keyed on `publishId` (an opaque, collision-safe routing token — decoupled from any node ID, unlike the old `nodeId:counter` form), an internal hop-dedup token separate from the content-hash `msgId` the application sees. Multiple arrivals of the same `publishId` at one peer collapse to a single `_deliveryCallback` firing.

K-CLOSEST IS COMPUTED AGAINST THE LOCAL SYNAPTOME, so in a browser mesh that fills in gradually the answer is a moving target: a peer that subscribed while it only knew the bridge picked a narrow axon set, and a publisher that joined later computes a wider one — they can fail to overlap, and the publish misses that subscriber. The result is cached per topic under an epoch counter to avoid recomputing on every `pub/sub`; the fix is to *invalidate* that cache as the mesh changes. `AxonaPeer` bumps the epoch on every `onPeerBound`, and `AxonaManager.refreshTick` (a 10s timer the application arms after subscribing) flushes it and re-issues every live subscription against the now-converged set. So a subscription anchored early re-anchors as the mesh widens, rather than staying pinned to a stale, narrow axon set.

Topic Metrics: Scatter-Down, Gather-Direct

`PEER.METRICS(TOPIC)` reports a topic's live state — `publishes`, `current_count` (events retained in the tree right now), `subscribers`, `deliveries`, `pulls`, `reshares` — by querying the same axon tree that carries delivery. It is a *scatter-down, gather-direct* traversal, not a recursive sum back up the tree.

1. **Route to the tree.** `requestMetrics` issues a routed `pubsub:metricsReq` toward the `topicId`; a node that doesn't host the topic returns `'forward'`, so the request DHT-routes until it reaches a root axon.
2. **Scatter down.** On receipt a role replies with its *own* local view, then forwards the request as `pubsub:metricsBroadcast` to each of its children, which recurse — the probe walks the tree along the same parent→child edges that fan messages out. A bounded seen-request set (keyed on `requestId`) makes each relay process a given probe once, so redundant edges neither loop nor double-count.
3. **Gather direct.** Every responding relay sends its partial result *straight back to the original requester* (`pubsub:metricsResp`), not up through its parent. Scattering down but gathering directly keeps the root from becoming an aggregation hotspot.
4. **Bounded collection.** The requester accumulates responses for a fixed window (`timeoutMs`, default 500 ms) — there is no global barrier — then folds them. `relayCount` (distinct responders) is the honest coverage signal; the result is a best-effort snapshot, not a ledger.

WHY THE FOLD DIFFERS PER FIELD. Each relay holds a local slice: its replay cache, its `children` map, and per-post counters. The

tree replicates the message queue across the K root axons but *partitions subscribers* across the tree, so a uniform sum would be wrong. The client combines each quantity according to its nature.

Field	Fold	Why
publishes	max	distinct posts are replicated across roots; summing would multiply by K
current_count	max	replicas hold the same retained set; max also tolerates a lagging replica
deliveries, pulls, reshares	sum	partitioned work counters — each relay tallies only the actions it performed
subscribers	max	membership is partitioned; exact for an un-split (single-root) topic, a lower bound once the tree splits
relayCount	count	distinct responders — a coverage indicator

Replicated set-cardinality folds with max; partitioned per-action counters fold with sum; partitioned membership folds with max (approximate).

Subscribable Metrics: A Derived Topic

THE SCATTER-GATHER ABOVE is right for an occasional probe and wrong as an app primitive: per-user polling makes cost scale with audience $\times$ poll-rate, hammering the roots forever. The fix needs no new wire and no kernel state — it is a *convention* over the existing pub/sub: a topic's metrics are themselves **published**, to a second topic *derived* from the first, which clients **sub** instead of polling.

THE DERIVATION. For a data topic with resolved ID T , $\text{metricTopic}(T) = \{\text{region} = \text{the region byte of } T, \text{ name} = \text{"axona:metric:"} \parallel T\}$ — an open topic anyone who knows T can recompute and subscribe. The reserved **axona:metric:** namespace (rather than a bare hash) is load-bearing: it makes a metric topic *self-identifying*, so a root recognises one and refuses to compute metrics-of-metrics — the recursion guard. Because clients and infrastructure must derive it and apply the guard *identically*, both live in the kernel core beside **deriveTopicId**, not in the optional **std** helpers (relays vendor **src/** only).

THE LOOP. The primary root of each live *open* topic recomputes its local snapshot — via **peer.rootedTopics()**, the read side, which reads replay-cache state with *no* network — and pubs a *signed* record to $\text{metricTopic}(T)$ on a ~ 5 -minute cadence, self-triggered on its own liveness for T . A subscriber's **sub(metricTopic(T), {since:'all'})** then gets the latest snapshot via replay-on-subscribe plus every later update — one subscription replacing an unbounded stream of point

A self-authenticating ownership gate means a relay answers only when the topic's publisher matches the requester, so today the owner alone gets non-empty results. **current_count** and live **subscribers** were added in kernel v2.11.0.

queries. The expensive part (counting) happens once per cadence at one node, no matter how many are watching, and distribution rides the delivery path subscribers already pay for.

HISTORY FOR FREE. Snapshots are ordinary retained messages, so they accumulate in the replay cache and age out at the 48h hold ceiling. A subscriber therefore gets a *rolling ~48h trend* at no extra cost — the cache TTL *is* the retention window; no overwrite, no compaction.

TWO DELIBERATE LIMITS. (i) *Open topics only*: an owner-only topic's subscriber count is owner-private (the gate above), and republishing it to an open metric topic would leak exactly that — so the loop skips owned topics; their metrics stay behind `peer.metrics()` for the owner. (ii) *Advisory, not authoritative*: the metric topic is open, hence spoofable; the snapshot is signed so a reader may pin trust to a known relay's key, but the protocol cannot *prove* a count. Metrics are decoration, never a security input. `peer.metrics()` survives as an **owner-only** on-demand reader (v3.5.0): it answers only the owner of an owned topic and refuses open/public topics outright, so a non-owner can no longer aim a K-root fan-out at any topic — open topics are read only through their metric topic.

XI The Application API

EVERY AXONA APPLICATION TALKS TO THE PROTOCOL THROUGH ONE IMPORT. The barrel export at `@axona/protocol` is the entire contract surface — a small set of factory functions, one peer class, and a handful of helpers. This section is the field-by-field reference.

Identity

`createNodeIdentity({ lat, lng })` *mint the CONNECTION identity*

lat number. Latitude in decimal degrees.

lng number. Longitude in decimal degrees.

Returns a Promise resolving to `{ id, pubkey, pubkeyHex, privkey, privateKey, region }`. `id` is the 66-char hex `nodeId` (8-bit S2 prefix from `(lat,lng)` + 256-bit SHA-256 of the public key). `privateKey` is the in-memory `CryptoKey`; `privkey` is the base64 export for persistence. This is the **connection/transport** keypair — it forms the

Shipped kernel v3.4.0
(`metricTopic/isMetricTopic`,
`rootedTopics()`) and relay v0.16.0
(the publish loop) — a convention,
no flag day. On testnet one keyspace-
hosting relay republished ~300 open
topics per cycle with the recursion
guard converging: metric topics
rooted in cycle n are recognised and
skipped in cycle $n+1$.

`nodeId`, drives the handshake and routing, and is passed to `AxonaPeer` as `nodeIdentity` (and to the transport factories as `identity`). It *never* signs a publish.

`createAuthorIdentity({ persistAs? })` *mint an AUTHORSHIP identity (Author ID)*

`persistAs` (optional) string storage key. Omitted, the author identity is ephemeral; given, it is load-or-created and saved (browser `localStorage` by default, or a custom `store`).

Returns a Promise resolving to a location-free keypair { `authorId`, `pubkey`, `pubkeyHex`, `privateKey`, ... } — no `nodeId`, no region. `authorId` is the public key, the `signerPubkey` that appears on a signed publish. This is the key you pass per-publish as { `signWith` }; authorship is not a place, and one peer can run several personas by varying `signWith`.

`dumpIdentity(identity)` *serialize for storage*

Returns a Promise resolving to a plain JSON object you can write to disk or IndexedDB. The reciprocal `loadIdentity(envelope)` reconstructs the `CryptoKey`-bearing form on the next launch.

Transports

`webTransport({ ... })` *browser transport: WebSocket to the bridge + WebRTC mesh*

`bridgeUrl` string. e.g. `wss://bridge.axona.net`.

`identity` connection identity envelope from `createNodeIdentity`.

`log` (optional, default () => {}) function (`event`, `data`) that receives transport-internal diagnostic events.

`WebSocketImpl` (optional) constructor for the `WebSocket` class. Defaults to `globalThis.WebSocket`.

`autoHandshake` (optional, default true) when true, `transport.start()` drives the full bridge admission sequence end-to-end — the WebSocket version gate, the application-level `hello/hello-ack` exchange, and binding the bridge as a peer. Set to false only if you're driving the handshake yourself.

`peerVersion` (optional, default `KERNEL_VERSION`) semver string sent in `client-hello`.

`handshakeTimeoutMs` (optional, default 15000) how long to wait for the bridge's `hello` before rejecting the *first* `start()`.

`reconnect` (optional, default true) auto-reconnect with exponential backoff on an unexpected socket close (bounds tunable via

`reconnectInitialMs / reconnectMaxMs`). A 4426 version-gate rejection is the one close that does *not* reconnect.

Returns a `CompositeTransport`. After `await transport.start(identity.id)` the bridge is reachable as a peer (`transport.bridgeNodeId`) and the WebRTC mesh is ready to accept channels. Beyond the `Transport` contract the returned object carries the connection-observability surface the factory now owns:

`onBridgeState(cb)` subscribe to state transitions `connecting | open | stale | disconnected | upgrade-required`; `bridgeState` is the current value, `upgradeReason` the 4426 detail.

`onWelcome(cb)` the bridge's `welcome` (`{connId, version, kernelVersion, turn}`), also cached on `bridgeInfo` and replayed to late subscribers.

`bridgeRtt / bridgeRttAvg` last and windowed ping→pong round-trip in ms.

`onPingTraffic(cb)` per-frame (`nodeId, 'sent' | 'recv'`) pulses for live link indicators.

`reconnectNow()` force an immediate reconnect with the backoff reset (tab-resume / network-online).

`nodeTransport.server({ ... })` *server-side WebSocket transport (Node 20+)*

`nodeTransport.client({ ... })` *Node WebSocket client*

Used by `axona-bridge`'s embedded peer and by Node applications.

Same `Transport` contract as the browser transport; no WebRTC.

`simTransport({ network, identity })` *in-process simulation transport*

`network` `SimNetwork` shared across peers.

`identity` identity envelope.

For tests and `dht-sim`. Messages are scheduled in the event loop with synthetic latency from a real internet topology dataset.

AxonaPeer

`new AxonaPeer({ ... })` *the per-node DHT instance*

`node` `NeuronNode` instance for this peer (carries id, lat, lng, the synaptome).

`domain` `AxonaDomain` the peer joins. Construct one per network with `new AxonaDomain({ k: 20 })`, where `k` (*default 20*) is the routing closest-set size (the root-set replication factor *R* is separate; see § X).

transport the transport instance from `webTransport`, `nodeTransport`, or `simTransport`.

nodeIdentity (*optional*) the connection identity envelope from `createNodeIdentity`. Used for the handshake, routing, and signing the node's own `kill/unpub`. It *never* signs a `publish` — authorship is supplied per-publish via `{ signWith }` (an author identity), so a peer holds no default author.

axonaManager (*optional*) explicit `AxonaManager`. When omitted, one is built automatically on `peer.start()`.

persist (*optional*) `PersistenceAdapter` for identity / synaptome / subscription checkpoints (not replay caches). Default: no persistence.

Lifecycle

`await peer.start()` *wire handlers, install delivery hook*

`await peer.stop()` *detach handlers and event listeners*

`await peer.join(sponsor?)` *production bootstrap path*

sponsor (*optional*) 66-char hex node ID. If given, opens a transport channel to that peer and seeds the synaptome from it. If omitted, returns immediately — the peer is “joined” but isolated; inbound connections from other peers populate the synaptome.

`await peer.leave(opts?)` *leave the network gracefully*

opts.drain (*default true*) wait for in-flight publishes to settle.

opts.notify (*default true*) send a `peer-leaving` notification to every peer in the synaptome.

opts.timeoutMs (*default 5000*) maximum drain window in ms.

Pub / Sub

`await peer.pub(topic, message, { signWith })` *publish to a topic*

topic structured descriptor `{ region?, owner?, name, write? }` (*not* a string). Three canonical shapes: an open lobby `{ region, name }` (anyone may publish); an author feed `{ region, owner, name, write: 'owner' }` (owner-only writes; if `region` is omitted it defaults to the publisher's own node region, so subscribers must name the region to read it from elsewhere); an author inbox `{ region, owner, name }` (open write under an owner namespace).

message JSON-serializable payload.

signWith **required**. An *author* identity from `createAuthorIdentity`; the publish is signed with its Ed25519 key. Pass the `ANONYMOUS` sentinel (`{ signWith: ANONYMOUS }`) to publish unsigned. Omitting `signWith` throws `PUBLISH_NO_PUBLISH_IDENTITY` — anonymity is explicit, never the silent default. There is no `{ sign: false }` and no `{ publisher }` option.

On an owner-only topic the signer must be the owner key, or `pub` throws `WRITE_POLICY_VIOLATION` (the topic's roots enforce the same at ingress). Returns a Promise resolving to the `msgId` — the 64-char hex content hash `SHA256(signerPubkey || message)`. Throws `PublishError` on invalid inputs or signing failure.

`await peer.sub(topic, handler, opts?)` *subscribe to a topic*

topic structured descriptor `{ region?, owner?, name, write? }` — same addressing as `pub`. To read another author's feed, pass their `owner` Author ID together with the feed's `region` — an Author ID has no region of its own, so discovery needs both the owner key and the region the feed was published in.

handler function (`envelope`) => void invoked for each delivery.

The envelope is `{ msgId, seq, ts, topic, message, signerPubkey? }` (`signerPubkey` and `signature` present only when the publish was signed; `topic` is the descriptor). A retraction (see `peer.kill`) arrives on the same handler as a marker, `{ deleted: true, msgId, topic }` — branch on `env.deleted` to drop your local copy.

opts.since (default: *live tail*) replay control. Pass `'latest'` for the most-recent cached message plus future deliveries, `'all'` for the full replay cache plus future, or a Unix-ms `number` for everything newer than that timestamp.

Returns a Promise resolving to a `Subscription` handle. Call `sub.stop()` to cancel that handle, or `peer.unsub(topic, opts)` (kernel v2.10.0) to leave the topic by name — stops every local subscription for it and, when the last goes, sends the network unsubscribe (self-only by construction).

`await peer.pull(msgId, opts)` *fetch a message from the relay's replay cache — by ID, or the latest*

msgId 64-char hex content hash, or null/omitted to fetch the topic's most-recent message (highest `seq`; kernel v2.10.0).

opts.topic the topic *descriptor* (required) — same shape as `sub`, used to derive the topic ID and locate the relay.

opts.timeoutMs (default 1000).

Returns the envelope or `null` if not found (which includes a message that has *expired* past its hold time). A successful pull slides the message's hold forward, bounded by its 48 h ceiling.

`await peer.metrics(topic, opts)` *aggregate counters across the relay tree*

topic the topic descriptor.

opts.timeoutMs (default 500).

Returns { publishes, current_count, subscribers, deliveries, pulls, reshares, relayCount }.

`await peer.host(topic?, opts?)` *store + serve a topic for others, without subscribing (kernel v2.40.0)*

topic string, or omitted. With a topic, the node hosts that one topic (given as a descriptor, same shape as `sub`). Omitted, the node hosts its **own keyspace neighborhood** — it is recruited as a root for whatever topics fall near its node ID (“host whatever lands near me”), the natural mode for a relay.

Hosting is **decoupled from subscribing**: a host stores and serves a topic for other peers but is *not* a consumer of it — it registers no delivery handler and is never added to the node's own subscription set. Mechanically it announces itself with the *same* `pubsub:subscribe-k` a subscriber already sends, so existing roots recruit it with no wire change (no flag day), then caches and serves replays exactly as a root does. It respects the proximity gate (B-2): a host only ends up routing topics it is genuinely among the K-closest to, so the primitive cannot force a node into the serving set for arbitrary or distant topics. `peer.unhost(topic?)` reverses it; `health().hosting` reports { `keyspace`, `topics` }. This is the participation primitive for relays and other infrastructure that should *carry* a topic without consuming it — before it, a node entered a topic's serving fabric only by subscribing, so a relay that meshed but never subscribed carried nothing.

Message lifecycle (kernel v2.10.0)

Three lifecycle controls, each **self-authenticating** — the authority to act is proven by the same keypair that proved authorship/ownership, with no registry or central gatekeeper.

`await peer.kill(topic, msgId, { signWith })` *retract a message you published (creator-only)*

The kill must be signed (`signWith`) with the *same author key that signed the original message*, else `KILL_SIGN_FAILED`. A topic's root axons accept it only on that proof; they drop the message from their

replay caches, record a short-lived tombstone (so a lagging replica can't resurrect it), and forward a delete marker to subscribers. It is best-effort redaction, *not* a cryptographic un-send — a subscriber that already received the plaintext can keep it — and an anonymous (`{ signWith: ANONYMOUS }`) message has no provable creator and so cannot be killed. Throws `KillError`.

`await peer.touch(topic, msgId, { signWith })` *keep-alive; reset a message's hold time*

Slides a message's expiry forward (bounded by the 48 h ceiling) so an actively-wanted message isn't swept. Self-authenticating like `kill`: on an open topic anyone may touch a message; on an owned topic only the owner key may. Throws `TouchError`.

`await peer.unpub(topic, { signWith, destroy? })` *remove a topic's whole queue (owner-only)*

The *owner* is the Author ID named in the topic descriptor (`owner`); `unpub` must be signed (`signWith`) with that author key, and roots verify the signer's pubkey equals the descriptor's `owner` self-authenticatingly. `opts.destroy` (default `false`) escalates from “clear the messages” to total removal (messages *and* config/ownership state). Open (owner-less) topics cannot be unpublished (`UNPUB_PUBLIC_TOPIC`). Throws `UnpubError`.

Topic limits

Each topic's replay queue is bounded (kernel v2.10.0): **max messages** (default 100, ceiling 256) with deterministic eviction of the lowest-ordered message by signed (`seq`, `ts`, `msgId`) — a max of 1 is a retained / latest-value slot; a **hold time** (default 24 h, hard ceiling 48 h) after which a message expires and is swept; and, on *open* (anyone-may-publish) topics, a **per-publisher quota** (one-quarter of the queue) so a single publisher can't flood out the rest. Owner-gated topics are governed by their publish rules instead.

Direct messaging

`await peer.send(targetId, message)` *RPC-style; awaits the remote handler's return value*

`peer.notify(targetId, message)` *fire-and-forget*

`peer.onMessage(handler)` *register a handler for incoming direct messages*

targetId 66-char hex node ID. The peer must already be in the synaptome.

message JSON-serializable payload.

handler function (*fromId*, *message*) => *result*. The return value is sent back to the caller on *send*; ignored on *notify*. Async handlers are awaited.

Mesh introspection

peer.peers() returns an array of 66-char hex node IDs currently in the synaptome.

peer.onPeerJoin(handler) registers (*nodeId*) => void for new arrivals.

peer.onPeerLeave(handler) registers (*nodeId*) => void for departures.

peer.health() returns { *nodeId*, *synaptomeSize*, *peers*, *subscriptions*, *axonRoles*, *hosting*, *wireVersion*, *started*, *transport*, *meshDegraded* }.

peer.onLog(level, handler) *level* is 'debug', 'info', 'warn', or 'error'.

peer.onError(handler) registers (*err*) => void for unhandled errors.

peer.onUpgradeRequired(handler) registers () => void, fired when the bridge closes the connection with UPGRADE_REQUIRED (WS 4426) because the peer's wire version is below the network floor.

Lookup

await peer.lookup(targetKey) *iterative K-closest walk*

targetKey 66-char hex 264-bit key. Usually a topic ID or a node ID.

Returns { *found*, *path*, *hops*, *time* }. *found* is a boolean — whether the walk reached a peer at the target key; *path* is the list of nodes traversed; *hops* is the count; *time* is wall-clock ms.

Author class (agent legibility)

await peer.setAuthorClass(cls, { signWith, operator?, operatorSignWith?, label? }) publishes a voluntary, signed *author-class attestation* — *cls* is 'agent' or 'human' — to the author's own owner-only profile topic, declaring what kind of actor is behind the key.⁶

await peer.getAuthorClass(authorId) resolves any author's declared class from its Author ID alone, returning { *class*, *operator*, *operatorVerified*, *label*, *ts* }. *class* is 'agent', 'human', or 'unstated' — absence is *unstated*, never silently 'human'. It is *provenance*, not *detection*: a deceptive actor can decline to flag or flag falsely; what it shapes is the honest default (see the agent-legibility item under Future Directions).

⁶ The attestation is signed by the author key and carried on a pinned-region, owner-only topic derived from the Author ID; a declaring peer should *host()* that topic so the claim stays resolvable after the roots evict it. An optional *operatorSignWith* attaches a counter-signature, letting a named operator vouch for the key (*operatorVerified* below).

Snapshots

`await peer.snapshot()` returns a JSON object capturing identity, synaptome, and subscriptions (plus the wire version). Symmetric with `AxonaPeer.fromSnapshot(snap, deps)`, which reconstructs a peer from one of these dumps without touching the network.⁷

⁷ It is a routing/identity checkpoint, not a message store — axon roles and replay caches are *not* captured, so a restored peer rebuilds the topics it hosts or serves from the live mesh.

Helpers

`await deriveTopicId(descriptor)` returns the 66-char hex topic ID for a topic descriptor { `region?`, `owner?`, `name`, `write?` } — the same derivation `pub/sub` use internally. Useful for displaying the publish location in a UI.

`geoCellId(lat, lng, bits)` returns the integer S2 cell at the given precision. `bits = 8` matches the top 8 bits of every Axona `nodeId`.

XII Identity and Security

THERE ARE TWO KINDS OF IDENTITY, AND THEY ARE SEPARATE. A *node* (connection) identity is an Ed25519 keypair whose public key is hashed with SHA-256 to form the bottom 256 bits of the `nodeId`, with the top 8 bits set to the S2 cell of the user's (`lat`, `lng`) (`createNodeId`). An *author* identity is a bare Ed25519 keypair with no `nodeId`, no region, and no coordinates (`createAuthorIdentity`) — authorship is not a place. Private keys never leave the host; the node public key is broadcast on every connection (hello/hello-ack) so peers verify each other, while the author public key travels only as the `signerPubkey` on a signed publish.

Every signed publish carries the *author's* public key (`signerPubkey`) inline; the kernel verifies the `ed25519:<hex>` signature against it with `crypto.subtle.verify`. The node/transport key **never** signs a publish — signing is `signWith`-only with an author identity, and an unsigned publish must opt in explicitly via the `ANONYMOUS` sentinel; there is no default signer. Verification is per-envelope, not per-peer — an untrusted relay can carry a signed publish from a trusted author unchanged, and any receiver along the way can verify the chain of custody.

WRITE POLICY IS ENFORCED AT THE STORING NODE, not merely claimed by the sender. A topic descriptor declares its policy — `open` (anyone may publish, self-signed) or `owner` (only the named author key) — and the author signs the *whole* descriptor. The node that stores and serves the topic independently recomputes the topic

ID from that signed descriptor and, for an owner topic, requires `signerPubkey === owner` before accepting a message — at *every* ingress path: direct publish, K-closest relay-accept, and anti-entropy sync. A publish naming someone else’s owner-only topic is refused at the publisher (`WRITE_POLICY_VIOLATION`) and again at the root, so a feed or profile cannot be spoofed by a third party even reaching a storing node directly. Authorship itself is location-free: the signed envelope reveals *who* signed (the Author ID) and *which* topic, but the Author ID carries no region and never appears in a node address. The one place location can surface is the topic’s region byte. When a publisher names an explicit `region`, that byte is the region they chose, not where they are — an author publishing into `useast` from anywhere discloses nothing about their own location. When `region` is *omitted*, it defaults to the publisher’s own node region, which does disclose that coarse cell; a location-conscious author therefore names a region explicitly rather than leaning on the omitted-region default.⁸

Persistence is opt-in. The kernel’s `PersistenceAdapter` contract exposes `write(namespace, value)` and `read(namespace)`; concrete adapters target IndexedDB (browser) and the filesystem (Node). When wired, AxonaPeer checkpoints identity, synaptome seeds, and pending subscriptions on a 5-second debounce. On restart, the identity envelope is reloaded and validated (the stored `nodeId` must match the freshly-derived one); subscriptions are reinstalled as soon as the mesh has stabilised.

The `crypto.subtle` dependency is a hard requirement: the demo and the peer fail to boot in insecure contexts (HTTP pages outside localhost) because the Web Crypto API simply isn’t available. The bridge serves over HTTPS; the peer’s GitHub Pages deployment enforces HTTPS at the Pages level; the demo at `demo.axona.net` required Let’s Encrypt cert provisioning before identity derivation would run at all.

Proof of work, decoupled from the address

IDENTITY MINTING IS GATED BY A PROOF OF WORK (the E-1 placement defense), shipped *inert* at difficulty 0 in kernel v2.34 so that difficulty is a tunable protocol parameter rather than a format change.⁹ The work is a search for a nonce such that $H(\text{domain} \parallel \text{role} \parallel \text{pubkey} \parallel \text{nonce})$ carries at least N leading zero bits; the winning nonce rides as a sibling wire field (`pow` on the hello, `signerPow` on the envelope). Per-role difficulty is intentional — a transport identity can be ground into a root position (eclipse) and warrants a higher cost; a publish key can only flood, so its cost is calibrated to anti-flood, not anti-eclipse.

⁸ A scope caveat, not an over-claim: the peer a publisher hands a message to over its own connection can still locally correlate that connection with the author key at send time, and the protocol carries no defence against network-level traffic analysis (out of scope for v1). At-rest and onward-hop data carry no location link beyond the chosen region byte.

Because the *node* identity is bound to a region prefix, moving geographic location changes the `nodeId`; axona-peer’s `identity.js` preserves it across reloads unless the user re-selects a region, so a roaming user keeps the same DHT (node) identity for the life of the browser profile. Authorship is unaffected either way: an Author ID has no region, so the same author key signs from any location or device.

⁹ The work is a search for a nonce such that $H(\text{domain} \parallel \text{role} \parallel \text{pubkey} \parallel \text{nonce})$ carries at least N leading zero bits; the winning nonce rides as a sibling wire field (`pow` on the hello, `signerPow` on the envelope). Per-role difficulty is intentional — a transport identity can be ground into a root position (eclipse) and warrants a higher cost; a publish key can only flood, so its cost is calibrated to anti-flood, not anti-eclipse. The work is a search for a nonce such that $H(\text{domain} \parallel \text{role} \parallel \text{pubkey} \parallel \text{nonce})$ carries at least N leading zero bits; the winning nonce rides as a sibling wire field (`pow` on the hello, `signerPow` on the envelope). Per-role difficulty is intentional — a transport identity can be ground into a root position (eclipse) and warrants a higher cost; a publish key can only flood, so its cost is calibrated to anti-flood, not anti-eclipse. new crypto. A. Back, *Hashcash — A Denial of Service Counter-Measure* (2002; orig. proposed 1997) — the $H(\text{input} \parallel \text{nonce})$ leading-zero-bits primitive; I. Baumgart & S. Mies, *S/Kademlia: A Practicable Approach Towards Secure Key-Based Routing* (ICPADS 2007) — crypto-puzzle Sybil/eclipse resistance for DHTs, and the *static* (ID-constraining) vs *dynamic* (separate-nonce) puzzle distinction this design turns on; J. R. Douceur, *The Sybil Attack* (IPTPS 2002) — the upstream result that a self-authenticating network needs a resource proof at all.

CRUCIALLY, THE WORK IS NEVER ON THE ADDRESS. The `nodeId` remains `[prefix]||SHA-256(pubkey)` and is unconstrained by the PoW at any difficulty. Encoding the work *into* the address — the “vanity address” form, S/Kademlia’s static puzzle — would be catastrophic for a heterogeneous keyspace: forcing leading-zero address bits clusters every PoW’d publisher (and the owned topics it anchors) onto `0x000...`, the keyspace edge, collapsing the DHT’s uniform load distribution — and only a *subset* of identities (publish keys) is PoW’d, while transport IDs and topics stay uniform. A trailing-zero variant would preserve placement (XOR closeness is dominated by the high bits) but re-couple difficulty to the address format, reviving the flag-day. The separate nonce avoids both: zero address bits constrained, full address entropy, difficulty tunable in place.

XIII Production Deployment

THE PUBLIC STACK RUNS AT FOUR URLS.

bridge.axona.net. The east (primary) bridge. Runs on a DigitalOcean droplet as a **Docker Compose stack** of three containers: the bridge (Node), **Caddy** (automatic Let’s Encrypt TLS plus reverse proxy), and **coturn** (a TURN relay for the peers that can’t negotiate a direct WebRTC link). Kept current via `git pull && docker compose build && docker compose up -d`.¹⁰ Caddy terminates `wss://` and proxies to the bridge on loopback; the bridge itself binds `127.0.0.1` only, never world-facing. `/healthz` is exposed for monitoring — an unauthenticated caller gets just `{status, version, kernelVersion}`, while the full synaptome/directory/uplink body and `/diag` require a `HEALTHZ_TOKEN` header.

bridge-west.axona.net. A second federated bridge in a western region, running the same Docker stack. It holds an *uplink* to the east bridge so the two embedded peers share one connectome and the bridge directory federates across both — a peer dialing either one reaches the same mesh. It also advertises its own **coturn** credentials into the welcome frame and its directory entry.

axona.net. The reference peer application, served via GitHub Pages from the **axona-peer** repo’s main branch. Every push redeploys. The vendored **axona-protocol** kernel ships with the app; the version row in the UI confirms which kernel an open tab is running.

demo.axona.net. The minimal pub/sub demo, served via GitHub Pages from the **axona-protocol** repo’s main branch (where **examples/minimal-pubsub-browser/** lives). Same Pages deployment pattern; HTTPS-enforced because **crypto.subtle** requires a secure context.

Before any difficulty > 0 ships, the scaffolding work hash (currently SHA-256) must be replaced with a memory-hard function so phone/browser minters aren’t out-classed by ASIC grinders.

¹⁰ The `systemd + nginx + certbot` path the bridge originally shipped with still exists in `deploy/` and is what **testnet.axona.net** runs; prod cut over to the Docker stack with the kernel 3.1.0 promotion.

THE v0.3 WIRE FLAG-DAY IS COMPLETE. The whole fleet — prod `bridge.axona.net` and `bridge-west.axona.net`, plus **testnet** (`testnet.axona.net` + `demo-testnet.axona.net`, served from a droplet) — runs the **3.x** kernel line: `WIRE_VERSION 3.0`, `AUTH_PROTO axona/5`, bridge gate `REQUIRED_WIRE_MAJOR=3` with both version floors at 3.0.0, so a pre-v3 client is closed with `UPGRADE_REQUIRED` (WS code 4426). The wire-2→wire-3 cutover was a hermetic flag-day — a wire-2 and a wire-3 peer reject each other’s envelopes — and was promoted to prod behind a gate rather than dragged along by a push; *within* the 3.x line, additive point releases are wire-compatible and roll with no flag-day. As of this writing the deployed bridges report kernel 3.6.0 and the kernel source sits a few point releases ahead at 3.8.0. The testnet static hosts serve their assets with `Cache-Control: no-cache` so a kernel module update is revalidated rather than served stale from a browser’s ES-module cache after a deploy.

A coordinated release works as follows: a fix lands in `axona-protocol`, bumps `KERNEL_VERSION` and the package version, tags the commit. `axona-peer` resyncs its `vendor/axona-protocol/` from the new kernel, bumps `PEER_VERSION`, and pushes (Pages picks up). The bridge updates its `@axona/protocol` dependency in `package.json`, bumps `VERSION` in `server.js`, pushes, and the operator runs the deploy script. Every component ends up reporting the same `kernelVersion` in its visible metadata.

XIV Future Directions

THE ARCHITECTURE IS STABLE. The shape of the layering, the wire format, the BigInt-canonical address discipline, the K-closest pub/sub semantics — these are not expected to change. What’s coming sits inside the existing seams:

Multi-bridge gossip. Production *already* runs federated bridges — an east `bridge.axona.net` and a west `bridge-west.axona.net`, joined by an outbound *uplink* so the pair share one connectome: each bridge’s embedded peer runs a `CompositeTransport` (inbound browser WebSocket *plus* an uplink client dialing the peer bridge), and the bridge directory federates across both. What remains future is full *N*-way peer-list gossip, so a peer connected to one bridge is *directly* reachable from any other beyond today’s uplink bootstrap chain. Either way it is a deployment story, not a protocol constraint — routing already treats each bridge as an ordinary `nodeId`.

Demote the bridge to an ordinary peer. Federated bridges re-

Every Axona component is required to surface its `KERNEL_VERSION` at a runtime-inspectable point: the bridge in `/healthz` and the welcome frame, the peer in its UI version row, the demo in its footer. The rule exists because before it did, “which kernel is everyone on?” required SSH-ing into the bridge and grepping `node_modules`.

move the single *rendezvous*; this companion change removes the single *super-peer*. The routing and pub/sub layers already treat a bridge as just another `nodeId` — there is no `isBridge` flag, no eviction exemption, no routing privilege.¹¹ The residual special status is purely at the *bootstrap* layer: a peer holds its bridge `WebSocket` open for signaling and *auto-readmits* the bridge to its synaptome on every (re)connect, so the bridge never actually leaves the routing table. Once a peer holds enough healthy mesh peers it should *drop the bridge from routing and root candidacy* — keeping the socket open for signaling only — and stop auto-readmitting its `nodeId`. With several bridges sharing the load, this also dissolves the *correlated* failure mode in which one bridge restart removes a root from nearly every topic at once. The goal state: a bridge is a signaling endpoint that any operator can stand up, indistinguishable from any other peer to the routing fabric. Validate convergence and churn behaviour in `dht-sim` before shipping.

Bridge directory. A well-known open topic — the fixed descriptor `{ region: 'useast', name: 'axona:bridge-directory' }` — where every bridge republishes itself once per hour: URL, region, kernel version, advertised TURN credentials, last-seen timestamp.¹² Peers subscribe at startup and cache the directory; a peer whose primary bridge disappears has a list of alternates to try without out-of-band coordination. To prevent spam and DDoS of the channel, the topic uses an *allowed-signer set*: only Ed25519 keys with a valid attestation chain back to one of the genesis bridge keys (baked into the kernel) can publish. Subscribers verify the chain on every publish and drop unattested envelopes at the `AxonaManager` layer, before they reach the application. Genesis keys can sign attestations admitting new bridges without a kernel release. This turns “federated bridges” from a deployment story into something clients can actually discover at runtime.

Capability hardening (endo / SES). Adopt the Agoric and MetaMask object-capability stack to harden the host-side surface, in three landings: `@endo/marshal` for envelope payloads (so payloads can carry `BigInt`, remotable references, and the slot-and-body wire shape `captp` expects);¹³ an opt-in `AxonaPeer.create({secure: true})` that calls `lockdown()` to freeze JavaScript primordials before kernel init, defending against prototype pollution and supply-chain mutation; and `@endo/captp` layered over `peer.send/peer.onMessage` to upgrade direct messaging from string-handler RPC to capability-style RPC with promise pipelining and revocable handles. A later phase ships `@endo/import-bundle` for content-addressed, signature-verified code distribution — the substrate for running third-party

¹¹ The bridge’s present dominance is *emergent*, not granted: every peer dials it first, so it lands in every synaptome and carries the highest degree; and its `0x89` geo-identity sits XOR-close to every `us-east/*` topic. The ordinary K-closest and vitality rules then keep selecting it as a root. The topology makes it central; the algorithm gives it nothing.

¹² Subscribers keep the newest entry per `signerPubkey` and prune entries older than 24 hours. Every peer resolves the same descriptor to the same topic ID (pinned to one region, since v0.3 has no global bucket), so all clients see the same listing regardless of their own geography.

¹³ Marshal output is JSON-shaped with `@qclass` annotations on typed values, so legacy peers see plain objects and capability-aware peers see typed data. Negotiated per-envelope via a `payloadFormat` tag.

agents and creator-published filters inside peer compartments with only the authority each receiver chose to grant.

Persistent topics. Today’s replay cache survives in-memory only.

A pluggable persistence layer for axon `role.replayCache` (with a TTL) lets late subscribers catch up across bridge restarts.

Topic access control — shipped in v3.0.0. What was a future direction here landed in the kernel: a topic descriptor declares `write: 'open'` (anyone may publish) or `write: 'owner'` (only the named owner Author ID), and storing nodes enforce it at every ingress path (`WRITE_POLICY_VIOLATION`; see § XII). What remains future is richer policy — multi-signer allow-lists and capability-scoped write grants beyond the single-owner case.

Cross-tab state sharing. Multiple tabs on the same browser today derive the same identity and form independent peers; a BroadcastChannel-based shared peer would reduce mesh load proportional to tab count.

Native transport. A C-based or Rust-based embedded kernel for IoT-class devices, sharing the wire format with the JS kernel so existing peers can talk to it without adaptation.

Connection-time optimization. The authenticated-identity handshake costs only sub-millisecond crypto (one Ed25519 sign plus one verify per link); the felt cost of a cold start is dominated instead by WebRTC ICE/DTLS bring-up, un-bundled module loading, and the few added round-trips of the mesh handshake — none of which touches steady-state routing or delivery latency. The recovery levers all sit at the client and transport layer, independent of the wire protocol: bundle the kernel modules, render the UI and authenticate in the background, and fold the mesh handshake into the DTLS/SDP setup so the signed hello rides the first data-channel frame.¹⁴ Not yet mission-critical; deferred until adoption makes first-connect time the binding constraint.

Black-hole detection (forwarding vitality). Every hardening to date answers “*is this peer who it claims to be, and may it do this?*” — identity, channel binding, ingress authorization, verified routing admission. A *black-hole node* passes all of them and then silently drops (or, equivalently, fails to forward) the traffic it should relay. This is a Byzantine behavioural fault, not an identity one. Tampering is already defeated (end-to-end signatures + A-1) and replay/re-order by C-2’s `seq`+freshness; *omission*, though, is only *inferable*, never provable — you cannot prove a negative, and a malicious drop is indistinguishable from a crash or a bad link. The direction is a *local* forwarding-vitality signal — each node’s own first-party

¹⁴ The last lever pairs naturally with binding the channel-binding value to the DTLS certificate fingerprint: once the CBV derives from the fingerprint already exchanged in the SDP, no nonce round-trip is needed and the handshake adds no extra latency — security and speed land together.

measurement of whether traffic routed through a peer reached an acknowledged destination, over the existing K-closest path redundancy — driving route-around, never a gossiped reputation verdict (which the no-central-authority constraint forbids). Equivocation is the one bad-faith behaviour that *is* provable (signed conflicting artifacts) and can land earlier. The gating question — the false-positive rate against honest-but-flaky nodes — must be answered in simulation first; see the black-hole threat note and the heterogeneous-protocol simulation plan in `axona-docs`. This is the behavioural complement to verified admission: that governs *who is in your routing table*; this governs *whether they actually do their job*.

Decoupled publish identity — shipped in v3.0.0. This direction landed: authorship is now a free-standing *author* identity (`createAuthorIdentity`) independent of the node it publishes from, signed in per-call via `{ signWith }`, with the node/transport key barred from signing. One peer can hold several authors; an author survives node churn and publishes from any device.¹⁵ The cutover was the v0.3 flag-day — it was *not* additive after all, because folding authorship out of the node key meant the signed envelope had to carry the topic descriptor too, which moved the wire (v2→v3). Three residual concerns remain genuinely future, none in the wire format: **(1)** it anonymizes the *author field*, not the *act* — a first-hop axon can still time-correlate origination (the § XII traffic-analysis caveat); **(2)** anti-abuse must rebind — a free, infinite supply of author keys defeats per-publisher quota, so rate-limiting must attach to the transport node or the author key must carry proof-of-work (the open E-1 / quota threads); and **(3)** the per-publisher `seq` (C-2) is per-*signer*, so one author key used from two devices needs coordinated sequencing. Authorship trust is now a key-distribution question (which Author ID is whom), as for any pseudonymous-identity system.

Governance as gradient (the Gates to Gradients line). A companion six-note series in `axona-docs` works a different axis from the security hardening above — not *who may act* but *how a network with no chokepoint shapes behaviour without a gate to pull*. One of the six has already shipped: **agent legibility**, a voluntary signed author-class attestation (`class: 'agent' | 'human'`) carried on the decoupled author key and resolvable from the Author ID alone — it gates no one, and a subscriber or filter may choose to badge, throttle, or quarantine agent-class traffic.¹⁶ The remaining four sit, like everything here, inside the existing seams; each is a *gradient* (a cost or a piece of context that scales with reach) rather than a *gate* (a binary verdict someone owns and can be captured).

¹⁵ The Nostr `npub` model: a portable author identity, the same key signing from several transport keys.

¹⁶ Kernel v3.7–v3.8:
`buildAuthorClass / verifyAuthorClass`
`+`
`peer.setAuthorClass / getAuthorClass`,
 an owner-only profile topic the declarer `host()`s for durability, and an optional operator counter-signature. The reference web peer is the first consumer. Absence is *unstated*, never silently “human.”

Cascade telemetry — a ruler for the gradients. The highest-leverage, lowest-risk move, and the one to land first: every other gradient is shipped on faith until we can observe whether it actually changed how things propagate. The instrument is *aggregate, privacy-preserving* measurement of cascade *dynamics* — spread velocity, fan-out shape, agent-to-human ratio, and how often a retraction loses its race against the bytes — assembled the same way as everything else: first-party local stats, published to a topic, served by infra nodes that `host()` it, *never* a central collector. The privacy bar is an engineering property, not a promise — aggregation, sampling, *k*-anonymity, or differential-privacy noise applied *above* the shipped derived-metric-topic path so no individual message or participant is reconstructable. Explicitly not content surveillance; it ships a gauge, not a feature. Validate the noise/utility trade-off in `dht-sim` before it grades anything.

Soft retraction and annotation envelopes. The *hard* primitive shipped — owner-only `kill` + signed tombstones + `unpub` (Phase A), best-effort cooperative removal. The *soft* gradient layers on top: a signed *retracted-by-author* signal that removes nothing, references the original `msgId`, and rides the *same* mesh *alongside* the message it concerns — and its generalization, a signed *annotation envelope* that references a `msgId` so any party can attach a correction or dispute.¹⁷ No wire change: both are ordinary signed pub/sub messages keyed to an existing `msgId`.

Forkable filter sets. Of the four levers a platform moderates with — access, distribution, retention, attribution — Axona dissolves three and inverts the fourth, leaving *subscription* as the one surface a community still owns. The direction makes that surface plural and machine-scale: a **filter set** becomes a signed, versioned, *published* object — allow/deny rules over `signerPubkey`, topic, or content predicate — that anyone can subscribe to, compose locally, and *fork* by republishing under their own key. No new kernel primitive: app-layer convention plus a reference implementation, riding the same discover/rank/compose pattern the bridge directory already proves. The reference peer’s agent-class badge is the first, trivial consumer (it reads `getAuthorClass`); a published, composable filter object is the next step.

Friction scaled to reach. The hardest and least certain move: keep the one defensible lesson of the circuit-breaker era — *runaway amplification should cost something* — without a chokepoint to enforce it. Two already-decentralized mechanisms approximate it as a gradient: **per-relay local damping** (a generalization of the bounded-queue / quota / hold-time machinery the kernel already ships) and

¹⁷ Display and weight are the subscriber’s or filter’s choice. Annotations are *context*, never authorities that remove — the move from deletion (impossible on a substrate where every holder keeps the bytes) to attached, signed context everyone can read and act on.

reach-graded proof-of-work demanded by the K-closest roots that already see a topic's subscriber-set size.¹⁸ Each node decides only what *it* forwards and how fast — the soft cost curve emerges from many independent local policies, with no coordinator and nothing to capture. Honest scope: it prices *casual and scaled* amplification only; a determined adversary evades it by fragmenting or under-declaring reach, so it is a gradient, never a guarantee. Simulate the honest-publisher false-positive cost in `dht-sim` before shipping.

The job of the architecture is to make these additions land cleanly. The job of the existing layers is to keep working while they do.

Companion documents: the Axona Programmer Guide and Axona API Reference, the Axona Explainer, and the source repositories at github.com/axona-net.

A Appendix: A Minimal Example

THIS IS THE COMPLETE SOURCE of `minimal-pubsub` — the smallest end-to-end Axona program: two peers in a single Node process exchanging one signed message over the kernel's in-process simulation transport. About a hundred and sixty lines, no browser, no WebRTC, no bridge — just the kernel primitives wired together so the moving parts stay visible.¹⁹

It maps onto the Application API (§ XI) in order: derive the two identities (the location-bound `createNodeIdentity` and the location-free `createAuthorIdentity`), open a `simTransport` on a shared `SimNetwork`, compose `AxonaPeer` + `AxonaManager` from the kernel primitives, place an open topic with a structured `{ region, name }` descriptor, then run a `sub/pub` roundtrip — the publish signed per call with `{ signWith: author }`, never the connection key. The `SimNetwork` stands in for the real transports' channel-open and `axona:hello` admission, so the example needs no network and runs with `node index.js`.

examples/minimal-pubsub/index.js

```
// =====
// minimal-pubsub -- two Axona peers in one Node process, pub/sub roundtrip.
//
// What this demonstrates (kernel v3.0.0 / identity-authorship v0.3):
// * createNodeIdentity -> 264-bit Ed25519 CONNECTION identity in an S2 cell
// * createAuthorIdentity -> a location-free AUTHORSHIP key (Author ID)
// * Two SimTransports on a shared SimNetwork (kernel's in-process router)
// * Composing AxonaPeer + AxonaManager from the kernel primitives
// * Region-keyed topics via a STRUCTURED topic descriptor { region, name }
// * peer.pub(topic, msg, { signWith }) / peer.sub roundtrip across two
  peers
```

¹⁸ This is the *demand* side of the same proof-of-work the decoupled-author anti-abuse thread needs on the *supply* side (E-1): price the act of fanning out, not just the act of holding a key.

¹⁹ For the full browser wiring — WebRTC mesh, bridge fallback, reconnect, the live metrics readout — see `examples/minimal-pubsub-browser/` (the source behind `demo.axona.net`) and the reference peer at `github.com/axona-net/axona-peer`.

```

//
// Run: node index.js
//
// For real-world browser/Node wiring (WebRTC + bridge fallback, etc.),
// see https://github.com/axona-net/axona-peer/blob/main/src/client.js
// -- this example is intentionally simpler to keep the moving parts visible.
// =====

import {
  AxonaPeer, AxonaDomain, NeuronNode, AxonaManager, Synapse,
  SimNetwork, simTransport,
  createNodeIdentity, createAuthorIdentity,
  regionNameForLatLng, clz264,
} from '@axona/protocol';

// -- 1. Region -----
// us-east (Virginia). Both peers connect from this cell, and the topic
// is placed in the matching kernel region. v0.3 region is a first-class
// field of the topic descriptor -- no synthetic publisher needed.
const US_EAST = { lat: 38.0, lng: -77.0 };
const REGION_NAME = regionNameForLatLng(US_EAST.lat, US_EAST.lng); // e.g. '
useast'

// -- 2. Build a peer -----
// One function, called twice -- once for alice, once for bob -- wires
// identity + transport + node + AxonaPeer + AxonaManager together.

async function makePeer({ network, region }) {
  // 2a. Derive a 264-bit Ed25519 CONNECTION identity in this region's S2 cell
  //
  // plus a separate, location-free AUTHORSHIP key used to sign publishes.
  const identity = await createNodeIdentity(region);
  const author = await createAuthorIdentity();

  // 2b. Open a SimTransport on the shared SimNetwork.
  const transport = simTransport({ network, identity, heartbeatMs: 0 });
  await transport.start(identity.id);

  // 2c. Build the local DHT node. NeuronNode holds the synaptome and
  // routing state; AxonaDomain holds parameters shared across peers.
  // NeuronNode XORs ids as BigInts internally, so convert from
  // identity.id (hex string) to BigInt at construction time.
  const node = new NeuronNode({
    id: BigInt('0x' + identity.id),
    lat: region.lat, lng: region.lng,
  });
  node.transport = transport;
  const domain = new AxonaDomain({ k: 20 });

  // 2d. AxonaPeer is the per-node DHT contract implementation. The
  // connection key is the `nodeIdentity`; authorship is supplied
  // per-publish via `signWith`, never by the constructor.
  const peer = new AxonaPeer({ domain, node, nodeIdentity: identity, transport
  });
  await peer.start();

  // 2e. AxonaManager handles pub/sub. It needs a `dht` adapter that
  // forwards K-closest, sendDirect, routeMessage, and handler
  // registration to our AxonaPeer. Most of these are 1-line
  // delegations; sendDirect special-cases self-target for local
  // dispatch.
  const dht = {
    getSelfId: () => peer.getNodeId(),
    findKClosest: (...args) => peer.findKClosest(...args),
    routeMessage: (...args) => peer.routeMessage(...args),
    sendDirect: async (peerId, type, payload) => {
      if (peerId === peer.getNodeId()) {
        // Local-loopback: dispatch into our own direct handler table.

```

```

        const h = peer._directHandlers?.get(type);
        if (!h) return false;
        try { await h(payload, { fromId: peer.getNodeId(), type }); return
true; }
        catch (err) { console.error('self-sendDirect threw:', err); return
false; }
    }
    return peer.sendDirect(peerId, type, payload);
},
onRoutedMessage: (type, h) => peer.onRoutedMessage(type, h),
onDirectMessage: (type, h) => peer.onDirectMessage(type, h),
};
const axonaManager = new AxonaManager({ dht });
peer._axonaManager = axonaManager; // hand the AM directly to the peer
return { peer, identity, author };
}

// -- 3. Wire two peers + connect them -----

const network = new SimNetwork();

const { peer: alice, identity: aliceId, author: aliceAuthor } = await makePeer
({ network, region: US_EAST });
const { peer: bob, identity: bobId } = await makePeer({ network, region:
US_EAST });

// Open a SimNetwork channel between alice and bob so they're directly
// reachable, then admit each other to their synaptomes. Real transports
// (WebRTC mesh, WebSocket bridge) do this admission via the axona:hello
// / hello-ack handshake at channel-open time -- see axona-peer's
// axona_node.js for the production wiring.
await alice._transport.openConnection(bobId.id);

function admitSynapse(localPeer, remoteBigInt) {
    const localId = localPeer._node.id;
    const stratum = clz264(localId ^ remoteBigInt);
    const syn = new Synapse({ peerId: remoteBigInt, latencyMs: 1, stratum });
    syn.weight = 0.5;
    syn.inertia = 0;
    syn._addedBy = 'demo';
    localPeer._node.synaptome.set(remoteBigInt, syn);
}
admitSynapse(alice, BigInt('0x' + bobId.id));
admitSynapse(bob, BigInt('0x' + aliceId.id));

console.log('[alice] nodeId:', aliceId.id);
console.log('[bob] nodeId:', bobId.id);

// Give the kernel a tick to admit each other to their synaptomes.
await new Promise(r => setTimeout(r, 50));

// -- 4. Pub/sub roundtrip -----

// v0.3: a topic is a structured descriptor. An open lobby in a real region
// is just { region, name } -- anyone may publish (self-signed); the region
// places it in the keyspace. No synthetic publisher.
const TOPIC = { region: REGION_NAME, name: 'hello-world' };

const received = [];
const sub = await bob.sub(TOPIC, (envelope) => {
    received.push(envelope);
    console.log('[bob] received:', {
        msgId: envelope.msgId,
        message: envelope.message,
        signerPubkey: envelope.signerPubkey?.slice(0, 16) + '...',
    });
}, { since: 'all' });

```

```

console.log('[bob]   subscribed:', sub.topicId);

// Wait for the subscribe-k frame to reach alice's role.
await new Promise(r => setTimeout(r, 100));

// Alice signs the publish with her AUTHOR key (signWith) -- the connection
// key never signs content.
const msgId = await alice.pub(TOPIC, 'hello from alice', { signWith:
    aliceAuthor });
console.log('[alice] published msgId=' + msgId);

// Let the publish-k -> cache -> fan-out cycle complete.
await new Promise(r => setTimeout(r, 200));

console.log();
console.log(received.length === 1
    ? '[OK] roundtrip ok -- bob received alice\'s envelope'
    : `[X] roundtrip failed -- expected 1 envelope, got ${received.length}`);

await sub.stop();
process.exit(received.length === 1 ? 0 : 1);

```