

AXONA

Machine-Oriented Grounding for AI-Built Apps

David A. Smith

Axona.net

The Axona AI Grounding v4.27.1 2026-07-17

Axona AI Grounding — kernel 4.27.1

This file is the complete, self-contained grounding for an AI system building an application on the Axona protocol. It matches the network it targets: **kernel 4.27.1 / wire 4.0, deployed on testnet (wss://testnet.axona.net)**. Everything below is exact and current; nothing outside this file is required. If this version does not match the bridge you are connecting to, request the matching grounding file.

Axona is a serverless peer-to-peer pub/sub network. Every participant is a peer. Topics are addresses derived from descriptors; messages route to an emergent per-topic root and fan out through a distribution tree. There is no central server, message broker, or database.

HARD RULES — violations cause silent failure or thrown errors

1. **A topic is a descriptor object, never a string.** { region?, owner?, name, write? }. Publisher and subscriber must use IDENTICAL fields — the topic ID is a pure hash of them. A mismatch (one names region, the other omits it) means two different topics and **no delivery, with no error**.
2. **Every publish names its signer.** peer.pub(topic, msg, { signWith }) — signWith is an author identity or the ANONYMOUS sentinel. Omitting it throws PUBLISH_NO_PUBLISH_IDENTITY. There is no default signer.
3. **Two identities, never interchangeable.** The node identity (createNodeIdentity) is the connection — it has a location and forms the node ID; it NEVER signs content. The author identity (createAuthorIdentity) signs content — it has NO location and NO node ID.
4. **await peer.ready() before first pub/sub.** Publishing on a cold mesh is survivable (the kernel retries) but unpredictable. ready() resolves when the routing mesh is warm.
5. **A publish returns no delivery acknowledgment — ever.** This is a privacy invariant (author and transport identities must stay uncorrelated). To confirm delivery, subscribe to the topic yourself and watch for your own msgId to arrive.
6. **Message payloads: 15 KB is the reliable cross-browser floor.** Hard ceiling 256 KB, but anything over 15 KB may fail on real WebRTC channels. For larger payloads use std/chunk (below), never a bigger single publish.
7. **Messages expire.** Default hold ~24 h (48 h ceiling). Re-publishing the identical payload from the same author yields the SAME msgId and refreshes the hold (an upsert; subscribers are not re-notified). Axona is a live messaging fabric, not permanent storage.
8. **kill is the only retraction** and its signature is peer.kill(topic, msgId, { signWith }) — the **descriptor comes first**, exactly like pub. peer.kill(msgId) (omitting the topic) is the most common mistake observed in the field and throws. The kill must be signed by the same author key that signed the original, and is best-effort (receivers already holding the plaintext keep it). Anonymous messages cannot be killed. unpub and touch do not exist in 4.x — do not call them.
9. **Persist the author key if user identity should survive restarts.** createAuthorIdentity({ persistAs: 'myapp:author' }) (browser localStorage). Everything else (node identity, subscriptions) can be ephemeral and re-created each run.

10. **Node 20+ or a browser over HTTPS.** The kernel needs Web Crypto (`crypto.subtle`) and ES modules. No build step, no bundler required.

Install

`npm install github:axona-net/axona-protocol#v4.27.1`

package.json must contain "type": "module".

```
import { connect } from '@axona/protocol/connect.js'; // the one-call bootstrap
import {
  AxonaPeer, AxonaDomain, NeuronNode, ANONYMOUS,
  createNodeIdentity, createAuthorIdentity,
  deriveTopicId, metricTopic, KERNEL_VERSION,
} from '@axona/protocol';
import { webTransport } from '@axona/protocol/transport/web/index.js';
import { makeMessage, readMessage } from '@axona/protocol/std/message.js';
import { publishChunkedBytes, receiveChunkedBytes } from '@axona/protocol/std/chunk.js';
```

Canonical peer bootstrap (copy exactly)

```
const { peer, author, status, disconnect } = await connect({
  bridge: 'wss://testnet.axona.net',
  location: { lat: 38.0, lng: -77.0 }, // the user's real location
  author: 'myapp:author', // string = durable (persistAs); true = ephemeral;
  ↪ false = none
});
// status: { ready, peers, ms, reason } - never rejects; ready:false on timeout

Teardown: await disconnect();
```

Options: `k` (routing set size, default 20), `ready` (forwarded to `peer.ready()`; `false` skips the wait), `transport/nodeIdentity` (injection for tests or custom stacks), `web` (extra `webTransport` options).

Alternative: manual assembly (custom transports / multiple peers)

```
const nodeIdentity = await createNodeIdentity({ lat: 38.0, lng: -77.0 });
const author = await createAuthorIdentity({ persistAs: 'myapp:author' });
const transport = webTransport({ bridgeUrl: 'wss://testnet.axona.net', identity: nodeIdentity
  ↪ });
const node = new NeuronNode({ id: nodeIdentity.id, lat: 38.0, lng: -77.0 }); // hex id
  ↪ accepted (4.14+)
node.transport = transport;
const peer = new AxonaPeer({ domain: new AxonaDomain({ k: 20 }),
  node, nodeIdentity, transport });
await transport.start(nodeIdentity.id);
```

```
await peer.start();
const status = await peer.ready();
```

Topics

A topic descriptor: { region?, owner?, name, write? }

Field	Type	Meaning
region	'useast' '0x89' 137	Geographic cell the topic lives in. All three forms are equivalent. Omitted → defaults to the publisher's own region (co-located peers converge; cross-region readers must then know that region). Prefer explicit.
owner	64-hex Author ID absent	Namespaces the topic under an author.
name	non-empty string	The human label. Required.
write	'open' 'owner'	Who may publish. Defaults: no owner → 'open'; owner present → 'owner'.

The three canonical shapes:

```
const lobby = { region: 'useast', name: 'lobby' }; // open room: anyone
↳ reads/writes
const feed = { region: 'useast', owner: me.authorId, name: 'posts' }; // owner-only feed
↳ (write defaults 'owner')
const inbox = { region: 'useast', owner: me.authorId, name: 'inbox', write: 'open' }; //
↳ anyone writes TO me
```

- await deriveTopicId(descriptor) → 66-hex topic ID (a shareable read handle). sub/pull/metrics accept a descriptor **or** this ID; pub always needs the full descriptor.
- Publishing to someone else's write:'owner' topic throws / is rejected at the network with WRITE_POLICY_VIOLATION.

Publish / subscribe

```
// Publish (signed). Returns the msgId (64-hex content hash). NO delivery ack.
const msgId = await peer.pub(topic, message, { signWith: author });
// message: any JSON-serializable value. Anonymous: { signWith: ANONYMOUS }.
```

```

// Subscribe. handler receives an envelope for EVERY delivery, exactly once per msgId.
const sub = await peer.sub(topic, (env) => {
  if (env.deleted) { /* a retraction: drop env.msgId locally */ return; }
  // env = { msgId, seq, ts, topic, message, signerPubkey?, signature? }
  // env.signerPubkey is the author's ID (absent when anonymous)
}, { since: 'all' });

// since options: omit = live only; 'latest' = newest cached + live;
//                  'all' = full cached history + live; <unix-ms number> = newer than.

await sub.stop();           // stop this handle
await peer.unsub(topic);    // stop ALL local subs for the topic + network unsubscribe

// Retract (author-only, best-effort):
await peer.kill(topic, msgId, { signWith: author });
// Subscribers receive { deleted: true, msgId, topic } on their normal handler.

```

Delivery semantics an app can rely on: - Exactly-once per `msgId` to each subscriber handler. - Publishing again with the identical payload + author = same `msgId` = refresh, not a duplicate delivery. - Late subscribers with `since: 'all'` receive cached history (bounded queue, ~24 h) then the live tail. - The kernel retries a publish internally until it confirms placement; a freshly-joined publisher gets an extra rapid burst. Your code does not retry.

Reads without subscribing

```

const env = await peer.pull(msgId, { topic });           // exact message by content hash, or
↳ null
const last = await peer.pull(null, { topic });           // topic's newest (may be slightly
↳ stale – served by the nearest replica ON PURPOSE)

```

If the strict newest matters, subscribe instead of polling `pull(null)`. A successful pull extends the message's hold a little.

Topic activity (metrics)

```

const m = await peer.metrics(topic); // one-shot: latest published snapshot
// Live: subscribe to the derived metric topic
const mt = metricTopic(await deriveTopicId(topic)); // -> a topic descriptor
await peer.sub(mt, (env) => {
  const snap = JSON.parse(env.message); // { topic, ts, subscribers, current_count, seq, ...
  ↳ }
}, { since: 'latest' });

```

TIMING RULES (metrics are published ON DEMAND — the subscribe itself turns them on; there is no always-on publisher):

- First snapshot: at ROUTING LATENCY (~0.3 s typical; allow a few seconds under churn) —

the root answers the moment the lease arms, independent of other users' publish/watch activity. Cadence: every ~20 s while subscribed. Publishing stops ~70 s after the last metric subscriber leaves.

- `peer.metrics()` (one-shot) normally succeeds within its 1.5 s default window; `stale: true` means the answer didn't arrive in time (churn) — retry or prefer the standing `sub()` form.
- Render silence as UNKNOWN, never as zero: a `current_count: 0` snapshot is the real “no activity” answer. Keep the subscription open in tests.

Metrics are advisory (unauthenticated counts) — never a security input.

Direct messages (peer-to-peer, not pub/sub)

```
const reply = await peer.send(targetNodeIdHex, { any: 'json' }); // RPC: awaits remote
↪ handler's return
peer.notify(targetNodeIdHex, { any: 'json' }); // fire-and-forget
peer.onMessage(async (fromNodeIdHex, message) => { return { ok: true }; });
```

Target must be a currently-known peer (`peer.peers()` lists node IDs). Direct messages address a *node* (device/session), not an author.

Presence / introspection

```
peer.peers() // -> ['<66-hex nodeId>', ...] currently in the routing table
peer.onPeerJoin(id => {}); peer.onPeerLeave(id => {});
peer.health() // { nodeId, synaptomeSize, subscriptions, hosting, wireVersion,
↪ ... }
peer.onLog('warn', (evt, data) => {}); peer.onError(err => {});
peer.onUpgradeRequired(() => { /* show "please update" */ });
```

Large payloads (files, images) — std/chunk

```
// Sender: splits into signed sub-15KB chunks, publishes each.
await publishChunkedBytes(peer, bytes /* Uint8Array */, {
  topic, signWith: author,
  meta: { filename: 'photo.jpg', mime: 'image/jpeg' },
});
// Receiver: reassembles complete payloads.
await receiveChunkedBytes(peer, topic, {
  onComplete: ({ bytes, meta }) => { /* whole file */ },
  onProgress: ({ received, total }) => {},
});
```

Interoperable message bodies — std/message

If other Axona apps may read your topic, wrap text bodies:

```
await peer.pub(topic, makeMessage('hello', { customField: 1 }), { signWith: author });
const { text } = readMessage(env.message); // tolerant reader
```

Serving without consuming — host

```
await peer.host(); // host this node's keyspace neighborhood (relay mode)
await peer.host(topic); // host one topic: store + serve it without subscribing
await peer.unhost(topic);
```

Use in long-running Node processes to keep topic history alive; browser apps rarely need it.

Agent legibility (recommended for AI-operated authors)

```
await peer.setAuthorClass('agent', { signWith: author, label: 'my-bot' });
const cls = await peer.getAuthorClass(someAuthorId); // { class:'agent'|'human'|'unstated',
  ↪ ... }
```

An AI publishing as an autonomous agent SHOULD declare 'agent'. It is voluntary provenance; absence reads as 'unstated', never 'human'.

Envelope shape (what handlers receive)

```
{
  msgId:      "<64-hex sha256(canonical({publisher, message}))>",
  seq:        1716638400123, // per-author monotonic
  ts:         1716638400000, // author's clock (ms)
  topic:      { region, owner, name, write }, // the SIGNED descriptor
  message:    <your payload>,
  signerPubkey: "<64-hex author id>", // absent if anonymous
  signature:   "ed25519:<128-hex>" // absent if anonymous
}
```

Signatures are verified by the kernel at network ingress AND on delivery; `env.signerPubkey` present means verified. The signature binds the whole descriptor, so a message cannot be replayed onto a different topic.

Limits

Limit	Value
Publish size (hard)	256 KB
Publish size (reliable floor — use std/chunk above this)	15 KB

Limit	Value
Message hold	~24 h default, 48 h ceiling; re-publish refreshes
Per-topic cached history	bounded queue (late joiners see recent history, not everything ever)
Per-publisher share of an open topic's queue	quota-bounded (one author cannot flood)
Live-publish freshness window	signed <code>ts</code> within ± 5 min of the root's clock

Error codes (err.code on thrown `AxonaError`)

Code	Thrown by	Meaning / fix
PUBLISH_NO_PUBLISH_IDENTITY	pub	Missing <code>signWith</code> . Pass an author or <code>ANONYMOUS</code> .
WRITE_POLICY_VIOLATION	pub	Signer $\neq$ owner on a <code>write:'owner'</code> topic.
TOPIC_REGION_REQUIRED	topic derivation	Open topic with no derivable region. Name a <code>region</code> .
PUBLISH_PAYLOAD_TOO_LARGE	pub	Over 256 KB. Chunk it.
PUBLISH_INVALID_TOPIC / SUBSCRIBE_INVALID_TOPIC	pub/sub	Malformed descriptor (e.g. empty <code>name</code>).
KILL_SIGN_FAILED	kill	<code>signWith</code> is not the original author key.
UPGRADE_REQUIRED (WS close 4426)	connect	Kernel/wire mismatch with the bridge. Install the tag matching this file.
TRANSPORT_PEER_UNREACHABLE	send	Target node not connected. Check <code>peer.peers()</code> .

`pull` returning `null` and `kill` resolving `{ ok:false }` are normal outcomes (missing/expired message; nothing to retract) — not errors.

Field-observed mistakes — every one of these has shipped in a real app

Check generated code against ALL of these before calling it done.

1. **`peer.kill(msgId)` — `topic` omitted.** Every content operation takes the **descriptor first**: `kill(topic, msgId, { signWith })`, exactly like `pub(topic, ...)` / `sub(topic, ...)`. There is no id-only form of anything.
2. **Reconnecting the peer to switch users/personas.** WRONG: tearing down the connection when the active author changes (it churns the mesh and drops live WebRTC links mid-ICE). One peer serves ANY number of authors — authorship is chosen per call via `{ signWith }`. Reconnect only when the bridge URL changes.

3. **Sharing a topic without its full descriptor.** An invite, directory ad, QR code, or link that transmits only name (or name + region) sends the recipient to a DIFFERENT topic whenever the original has an owner — {region, owner, name, write} ALL fold into the topic id (HARD RULE 1). Always transmit the complete descriptor.
4. **Deriving presence/recency from arrival time.** A `since:'all'` (or `'latest'`) subscription replays history: a message that arrives NOW may have been published hours ago. Recency logic (online lists, “last seen”) must read `env.ts` (publish time), never the local clock at delivery. For ephemeral signals like heartbeats, subscribe with `since` omitted (live-only) — replaying stale heartbeats is pure noise.
5. **Faking encryption with public inputs.** The protocol signs messages; it does NOT encrypt them, and it provides NO encrypt-to-author primitive. An author ID is **public** (it is `env.signerPubkey` on every signed message) — any scheme that derives a decryption key from it is plaintext with extra steps. Confidentiality needs a real app-layer key exchange (e.g. keys delivered out-of-band, or a proper ECDH scheme); if you cannot do that, say so in the UI rather than shipping theater.
6. **Inventing a fallback when derivation fails.** If `deriveTopicId` (or any kernel call) throws, surface the error. Substituting a locally-made-up id “so the UI keeps working” silently diverges that client from every other peer on the topic — the failure mode is invisible no-delivery (HARD RULE 1).

Troubleshooting rules

- **Connected but nothing arrives** → publisher and subscriber derived different topic IDs. Diff the descriptors field-by-field (including an omitted vs named `region`), or → published before `ready()`.
- **Cannot mix BigInt and other types** → pre-4.14 kernel with a hex id passed to `NeuronNode`; on this version the hex id is accepted directly (and `connect()` avoids the question entirely).
- **Works on localhost, fails deployed** → page not HTTPS (`crypto.subtle` unavailable).
- **Large message silently missing for some receivers** → over the 15 KB reliable floor. Use `std/chunk`.
- **User identity resets every reload** → author not persisted. Use `createAuthorIdentity({ persistAs })`.
- **WS close 4426** → version mismatch; reinstall the pinned tag.

Environment notes

- Browser: HTTPS only. Node: v20+; the same `webTransport` connects to the bridge over WSS (the WebRTC mesh is browser-side; Node peers converse via the bridge and routing).
- The testnet bridge is `wss://testnet.axona.net` (kernel 4.27.1, wire 4.0) — the network this grounding targets. Production (`wss://bridge.axona.net`) runs the same wire-4 line, typically one release behind; the two are wire-compatible but SEPARATE networks (a peer joins one or the other).
- Multiple tabs = independent peers (fine, but each is a separate node).

End of grounding. Human-oriented companions: the Axona Quick Start, Programmer Guide, and

API Reference (same version).