

Axona API Reference

Reference for every public symbol exported from `@axona/protocol` v2.40.0. Organized by what application developers actually reach for; deep-customization surfaces are at the end.

Companion documents:

- Quick Start — 5-minute working roundtrip.
- Programmer Guide — mental model + worked example + pitfalls.
- Security changelog — what each kernel version protects (authenticated handshake, channel binding, pub/sub trust boundary, verified routing admission). See §17 below for the developer-facing summary of the security model.

Imports throughout assume:

```
import { /* ... */ } from '@axona/protocol';
```

Sub-path imports (e.g. `@axona/protocol/contracts/Transport.js`) work for symbols that need it but most apps only need the main barrel.

Table of contents

Application surface (what you'll use first)

1. Identity
2. AxonaPeer construction + lifecycle
3. Pub/sub
4. Direct messaging
5. Mesh introspection
6. Topic IDs
7. Envelopes + verification
8. Errors
9. Persistence

Transport + protocol surface

10. Contracts
11. Sim transport
12. Handshake

Low-level utilities

13. Ed25519 helpers
14. Hex / ID math
15. Geo utilities
16. Constants

Types referenced throughout

Identity, Envelope, Subscription, Peer, TopicId.

Application surface

Types

Identity (type)

```

{
  id:          string;           // 66-char lowercase hex (264-bit nodeId)
  pubkey:      Uint8Array;       // 32-byte Ed25519 public key
  pubkeyHex:   string;           // 64-char lowercase hex
  privateKey:  CryptoKey;        // Web Crypto Ed25519 (browser + Node 20+)
  region:      { lat: number, lng: number };
  createdAt:   number;           // ms epoch
}

```

The top 8 bits of `id` are `geoCellId(lat, lng, 8)` (the S2 cell at level 8); the bottom 256 bits are `sha256(pubkey)`.

Envelope (type)

```

{
  msgId:       string;           // 64-char hex sha256 of canonical inputs
  seq:         number;           // per-publisher monotonic sequence (kernel v2.9.0)
  ts:          number;           // ms epoch when published
  topic:       string;           // the application-level topic name
  message:     any;              // any JSON-serializable payload
  signerPubkey?: string;         // 64-char hex; present iff signed
  signature?:  string;           // "ed25519:<128 hex>"; present iff signed
}

```

`seq` (envelope format **v2**, kernel v2.9.0) is folded under the signature alongside `ts`. The kernel uses it for **replay/freshness** (a captured envelope can't be re-injected as live once its signed `ts` ages out of the ~5-min window) and for **deterministic ordering** ("latest", bounded-queue eviction). You don't set it — `peer.pub` stamps it.

A **retraction** (see `peer.kill`) is delivered to your `sub` handler as a minimal marker instead of a full envelope:

```
{ deleted: true, msgId: string, topic: string | null }
```

Branch on `env.deleted === true` to drop your local copy of `msgId`.

Subscription (type) Opaque handle returned by `peer.sub()`. Has:

```

sub.id         // string - unique per subscription
sub.topicId    // string - 66-char hex topic ID
sub.topicName  // string - your original topic name
sub.stop()     // Promise<void> - cancel; idempotent

```

PeerId (type) In the kernel's public API, peer IDs are **66-char lowercase hex strings**. The synaptome stores them as `bigint` internally; conversion happens at the API boundary. Pass strings unless a function's signature specifically says `BigInt` (`peer.lookup`, the synapse admit example).

TopicId (type) Same shape as `PeerId` — 66-char lowercase hex. Top 8 bits encode the addressing-mode S2 prefix, bottom 256 bits are sha256 of the canonical input.

1. Identity

`deriveIdentity({ lat, lng, extractable? })` → `Promise<Identity>` Generate a fresh 264-bit Ed25519 identity. Top 8 bits of the resulting `id` encode the S2 cell containing (`lat`, `lng`); the bottom 256 bits are `SHA-256(pubkey)`, so the `nodeId` is **self-authenticating** — a peer can only claim an `id` it holds the private key for (this is what the `axona/5` handshake checks; see §17).

```
const id = await deriveIdentity({ lat: 38.0, lng: -77.0 });
console.log(id.id); // 'df7c5e...' (66 hex chars, us-east prefix)
console.log(id.id.slice(0, 2)); // 'df' (us-east S2)
```

Optional extractable (default true): pass false to hold the signing key as a **non-extractable** `CryptoKey` — it can sign but cannot be read back out by page script, so an XSS incident or a malicious dependency cannot exfiltrate the long-term identity. Recommended for ephemeral browser identities. Leave true only when you need to `dumpIdentity` it for persistence.

Returned `Identity`: { id (66-hex), pubkey (Uint8Array), pubkeyHex (64-hex), privateKey (CryptoKey), region, createdAt, sign(bytes), verify(bytes, sig) }.

Throws `IdentityError` if Web Crypto isn't available or lat/lng are out of range.

`dumpIdentity(identity) → Promise<envelope>` Serialize an `Identity` to a JSON-safe envelope (the `privateKey` is exported as base64 PKCS#8).

```
const env = await dumpIdentity(identity);
// {
//   id:      '<66 hex>',
//   pubkey:  '<64 hex>',
//   privkey: '<base64 PKCS#8 string>',
//   region:  { lat, lng },
//   createdAt: <ms>,
// }
localStorage.setItem('my-identity', JSON.stringify(env));
```

`loadIdentity(envelope) → Promise<Identity>` Inverse of `dumpIdentity`. Reconstructs the full `Identity` from the JSON envelope by importing the PKCS#8 key.

```
const env = JSON.parse(localStorage.getItem('my-identity'));
const identity = await loadIdentity(env);
```

Throws `IdentityError` (`IDENTITY_LOAD_FAILED`) if the envelope is malformed or the pubkey doesn't match the id's hash.

2. AxonaPeer construction + lifecycle

The peer is the per-node DHT contract implementation.

`new AxonaPeer({ domain, node, identity, transport?, axonaManager?, persist?, engine? })`
Construct a peer. Required:

- `domain` — an `AxonaDomain` instance (shared routing parameters, e.g. { k: 20 }). Mind the two distinct “K”s: the Kademlia **routing** parameter **k = 20** (`AxonaDomain._k` — sizes the closest-set for *lookups*) is **not** the pub/sub **replication factor R = 5** (`AxonaManager.rootSetSize` — the root-set size that *anchors each topic*). “K-closest” elsewhere names the generic find-closest-N operation both use (N = k for routing, N = R for replication).
- `node` — a `NeuronNode` (or any object with {id, alive, transport, synaptome}).
- `identity` — from `deriveIdentity`; needed for signed publishes.

Optional:

- `transport` — overrides `node.transport`. The peer uses it directly for send/notify/lookup round-trips.
- `axonaManager` — pre-built `AxonaManager`. If omitted, the peer resolves one on first pub/sub via `engine.axonaManagerFor(node)`.

- `persist` — a `PersistenceAdapter`; enables auto-checkpointing of identity envelope, subscriptions, and synaptome snapshots.
- `engine` — legacy; pass null for new code.

```
const peer = new AxonaPeer({
  domain: new AxonaDomain({ k: 20 }),
  node: new NeuronNode({ id: BigInt('0x' + id.id), lat, lng }),
  identity: id,
  transport,
});
```

peer.start() → `Promise<void>` Bring the peer up — installs wire-level routing handlers, sets up the delivery hook, and (if `persistence` is wired) hydrates any persisted state.

Idempotent. Returns the peer once started.

peer.stop() → `Promise<void>` Tear down. Drains in-flight pub calls, stops timers, releases listeners. Idempotent.

peer.join(opts?) → `Promise<void>` Production bootstrap: discover the network, populate synaptome to target size, ready for routing. Optional `opts` customize fan-out parameters; defaults are usually fine.

In simple in-process setups (where you admit synapses directly), you don't need `join()` — `start()` is enough.

peer.leave({ drain?, notify?, timeoutMs? }) → `Promise<void>` Graceful shutdown. With `drain: true` (default), waits up to `timeoutMs` (default 5000) for in-flight pubs to complete. With `notify: true` (default), sends `pubsub:unsubscribe-k` for every active subscription so axons can sweep our entries promptly.

```
await peer.leave({ drain: true, notify: true, timeoutMs: 3000 });
```

peer.getNodeId() → `string | bigint` Returns whatever was passed as `node.id` at construction time. If you followed the convention of passing a `BigInt` to `NeuronNode`, this returns `BigInt`; if you passed a string, it returns a string.

3. Pub/sub

peer.pub(topicName, message, opts?) → `Promise<msgId>` Publish a message to a topic.

Arg	Type	Notes
<code>topicName</code>	<code>string</code>	Application-defined name (gets hashed).
<code>message</code>	<code>any</code>	JSON-serializable. Strings, numbers, objects, arrays all OK. Size: see below.
<code>opts.publisher</code> <code>opts.sign</code>	<code>string \ null \ undefined</code> <code>boolean</code>	Addressing mode. See §6. Default <code>true</code> . Set <code>false</code> for anonymous publishes (no <code>signerPubkey/signature</code> in envelope).

Returns: 64-char hex `msgId` — sha256 of the canonical envelope inputs. Subscribers see this same id in `env.msgId`.

Maximum message size: 256 KiB (kernel v2.8.1; was 64 KiB before). The limit is on the *serialized envelope* (`JSON.stringify({ msgId, ts, topic, message, signature, signerPubkey })`), so your usable payload is 256 KiB minus a few hundred bytes of envelope overhead. It’s measured on string length, so heavily multi-byte (CJK/emoji) content can exceed 256 KiB on the wire.

`peer.pub` rejects an oversized message up front (kernel v2.8.2) with `PublishError(PUBLISH_PAYLOAD_TOO_LARGE)` — you get an immediate, catchable error rather than a message that silently fails to arrive. The root axon also enforces the same cap at ingress as defense-in-depth.

Pub/sub is a *broadcast* path (replicated to K root axons, cached, fanned to all subscribers), so it is the wrong place for images/documents even under 256 KiB. For binary content, publish a small **content-reference** (hash + size + mime) and transfer the bytes out-of-band, content-addressed.

```
const synth = '<region S2 hex>' + '0'.repeat(64);
const msgId = await peer.pub('news/world', { title: 'hi' }, {
  publisher: synth,
});
```

Throws `PublishError`:

- `PUBLISH_INVALID_TOPIC` — empty topic name.
- `PUBLISH_SIGN_FAILED` — signing failed (identity missing privateKey, Web Crypto unavailable, etc).
- `PUBLISH_PAYLOAD_TOO_LARGE` — serialized envelope exceeds the 256 KiB cap (kernel v2.8.2; see “Maximum message size” above).
- `PUBLISH_INVALID_MESSAGE` — the message is **not JSON-serializable** (circular reference, `BigInt`, etc.), so the envelope can’t be encoded.

`peer.sub(topicName, handler, opts?)` → `Promise<Subscription>` Subscribe to a topic. `handler` is invoked for every delivered envelope.

Arg	Type	Notes
<code>topicName</code>	<code>string</code>	Must match what publishers use.
<code>handler</code>	<code>(envelope) => void</code>	Called with the full <code>Envelope</code> .
<code>opts.publisher</code>	<code>string</code> \ <code>null</code> \ <code>undefined</code>	Must match publisher’s mode.
<code>opts.since</code>	<code>'all'</code> \ <code>'latest'</code> \ <code>number</code> \ <code>undefined</code>	Replay mode. See below.

since modes:

since	What you get
omitted / <code>undefined</code>	Live tail only — no cached messages.
<code>'all'</code>	Every message in the axon’s replay cache, then live tail.
<code>'latest'</code>	The most recent ~1s of cache, then live tail.
<code><number></code>	Messages with <code>publishTs > since</code> (ms epoch).

Returns: a `Subscription` handle. Call `.stop()` to cancel.

```
const sub = await peer.sub('news/world', (env) => {
  console.log(env.ts, env.message);
}, { publisher: synth, since: 'all' });

// later...
await sub.stop();
```

Your handler also receives **retraction markers** — { deleted: true, msgId, topic } — when a publisher kills a message (see `peer.kill`). Branch on `env.deleted === true` to drop your local copy:

```
const sub = await peer.sub('news/world', (env) => {
  if (env.deleted) { dropFromUI(env.msgId); return; }
  render(env.message);
}, { publisher: synth, since: 'all' });
```

Throws `SubscribeError`:

- `SUBSCRIBE_INVALID_TOPIC`
- `SUBSCRIBE_HANDLER_MISSING`

`peer.unsub(topicName, opts?)` → `Promise<{ ok, removed }>` Unsubscribe from a topic by name (kernel v2.10.0) — the counterpart to `peer.sub`. Stops **every** local subscription you hold for the topic (no need to keep the `Subscription` handle) and, once the last one goes, sends the network unsubscribe so the topic’s root axons drop you. Idempotent: unsubscribing a topic you’re not on returns { ok: true, removed: 0 }.

Arg	Notes
topicName	The topic to leave.
opts.publisher	Must match what you passed to <code>sub</code> (default = your own feed, <code>null</code> = public, <code>hex</code> = another publisher’s feed).

```
await peer.unsub('news/world', { publisher: synth });
```

Self-only by construction — it removes only *your* subscriberId (the network enforces this), so it can’t be used to unsubscribe anyone else. Throws `SubscribeError(SUBSCRIBE_INVALID_TOPIC)` on an empty topic.

`peer.pull(msgId, opts)` → `Promise<Envelope | null>` Fetch an envelope from the K-closest axons. Returns `null` if it’s not in any reachable cache window (which includes a message that has **expired** past its hold time — see Topic limits).

Arg	Notes
msgId	64-char hex from pub / a previous delivery — or null/omitted to fetch the topic’s most-recent message (kernel v2.10.0).
opts.topic	The topic name.
opts.publisher	Same addressing mode used to publish.
opts.timeoutMs	Default 1000.

```
// a specific message...
const env = await peer.pull(msgId, { topic: 'news/world', publisher: synth });
// ...or the latest on the topic (no msgId):
const latest = await peer.pull(null, { topic: 'news/world', publisher: synth });
```

“Latest” is the highest-ordered message by signed `seq` (then `ts`, `msgId`). A successful pull also **slides the message’s hold time** forward (now + hold), bounded by its absolute 48 h ceiling — reads keep a hot message alive, but can’t pin it forever.

Throws `PullError` on transport failure (not on cache miss — that returns `null`). Passing a non-null, non-64-hex `msgId` throws `PULL_INVALID_MSGID`.

`peer.metrics(topicName, opts?)` → `Promise<metricsObj>` Best-effort delivery, retention, and subscriber counts for a topic.

```
const m = await peer.metrics('news/world', {
  publisher: synth,    // null for a public topic
  timeoutMs: 500,
});
// {
//   publishes:    <distinct messages ever seen>,
//   current_count: <events live in the tree right now>, // kernel v2.11.0
//   subscribers: <max direct subscribers at a relay>,  // live, kernel v2.11.0
//   deliveries:   <total fan-out deliveries>,
//   pulls:        <total pull() hits>,
//   reshares:     <total reshare bumps>,
//   relayCount:   <distinct relays that replied>,
// }
```

`current_count` is the number of published events **currently retained** (live — non-expired, non-killed) across the topic’s tree, so it falls as messages are killed or age out under the hold-time TTL; `publishes` is the cumulative count of distinct messages ever counted, so it only rises. `subscribers` is the max direct-subscriber count reported by any single responding relay — exact for an unsplit topic, a lower bound once the tree splits into sub-axons. Values are aggregated from whichever axons reply within `timeoutMs`. Use for UX (“12 people in this room”), not for billing.

Coverage (kernel v2.14.0). The request fans out to the topic’s whole K-closest root set — the same set publishes replicate to — and merges the replies (max for replicated counters like `current_count/subscribers`, sum for partitioned ones like `deliveries`). Before v2.14.0 it queried only the single closest root, so in a churning mesh `current_count` could read 0 even while subscribers were receiving a full replay from a sibling root; if you saw that, the fix is a kernel bump, not a change to your call.

Access (kernel v2.12.0). Metrics for an **unowned** topic — a public topic (`publisher: null`) or a synthetic region-keyed topic (`prefix||0...0`, which no key can own) — are readable by anyone. For an **owner-keyed** topic (`publisher` = a real node ID) only that owner gets a response; everyone else’s request is declined (self-authenticating, no gatekeeper). So `metrics()` on someone else’s owned topic resolves with zeroed/empty counters, not their data.

`peer.host(topicName?, opts?)` → `Promise<{ ok, scope, topicId? }>` Store and serve a topic for other peers **without subscribing to it** (kernel v2.40.0) — the relay / infrastructure primitive. It makes the node a willing root/replica so publishes land on it and subscribers pull replays from it, but it registers **no handler** and delivers nothing to your app.

- `peer.host('news/world', { publisher: synth })` — host one named topic.
- `peer.host()` — no argument: host this node’s **own keyspace neighborhood**, i.e. get recruited as a root for whatever topics fall near its node ID (“host whatever lands near me”). This is the zero-config relay mode.

`opts.publisher` selects the topic-id derivation exactly like `peer.sub` (default = this node’s feed, `null` = public topic, hex = someone else’s feed). `peer.unhost(topicName?)` reverses it (omit the topic to turn off keyspace hosting). `health().hosting` reports `{ keyspace, topics }`.

```
// a headless relay that serves whatever is near it, consuming nothing:
await peer.host();
```

```
// or host specific app topics:
await peer.host('pow-bench/results', { publisher: synth });
```

Decoupled from sub. Hosting answers “I’ll carry this for others”; subscribing answers “I want to receive this.” A host is never added to your subscription set and fires no delivery handler. Under the hood it announces with the same `pubsub:subscribe-k` a subscriber sends, so existing roots recruit it with no wire change, and it respects the proximity gate — a host only roots topics it is genuinely K-closest to. Before v2.40.0 the only way a node joined a topic’s serving fabric was to `sub()`, so a relay that meshed but never subscribed carried nothing.

`peer.kill(topicName, msgId, opts?)` → `Promise<{ ok }>` Retract a message you published (kernel v2.10.0) — “unsend”. The topic’s root axons accept the kill **only if it’s signed by the same key that signed the original message** (creator-only, self-authenticating), drop it from their replay caches, tombstone the `msgId` so a lagging replica can’t resurrect it, and forward a delete marker to current subscribers (which arrives on their `sub` handler as `{ deleted: true, msgId, topic }`).

Arg	Notes
<code>topicName</code>	The topic you published to.
<code>msgId</code>	Required 64-char hex (the value <code>pub</code> returned). No “kill latest”.
<code>opts.publisher</code>	Must match what you passed to <code>pub</code> .

```
const msgId = await peer.pub('news/world', { title: 'oops' }, { publisher: synth });
await peer.kill('news/world', msgId, { publisher: synth });
```

Best-effort redaction, **not** a cryptographic un-send. A subscriber who already received the message can keep it; an offline subscriber may never see the purge. And an anonymous (`sign: false`) message has no provable creator and so **cannot** be killed.

Throws `KillError`: `KILL_INVALID_TOPIC`, `KILL_INVALID_MSGID`, `KILL_SIGN_FAILED` (no identity — kills must be signed).

`peer.touch(topicName, msgId, opts?)` → `Promise<{ ok }>` Keep a message alive (`touch` kernel v2.15.0; ownership gate v2.16.0). A `touch` is signed and routed to the topic’s K-closest roots; each root that holds the message **resets its hold-time expiry to now + hold** (bounded by the message’s absolute 48 h ceiling — exactly the bound a `pull` respects), moves it to the **head of the queue**, and makes it the **last entry evicted**. Use it to keep a still-relevant message (a pinned status, a current value) past its default hold **without re-publishing** it.

Who may touch — by **topic ownership** (the unpub/metrics model, *not* `kill`’s creator-only one):

- **open topic** (public, or a synthetic region-keyed anchor) → **anyone** may;
- **owned topic** (publisher-keyed) → **only the owner** (the touch must be signed by the owner’s key — its `sha256(pubkey)` is the owner `nodeId` suffix).

Arg	Notes
<code>topicName</code>	The topic you published to.
<code>msgId</code>	Required 64-char hex (the value <code>pub</code> returned).
<code>opts.publisher</code>	Must match what you passed to <code>pub</code> (selects the topic, and on an owned topic identifies the owner you must be).

```
const msgId = await peer.pub('status/ops', { state: 'green' }, { publisher: synth });
// ...23 h later, before it ages out, refresh it without re-publishing:
await peer.touch('status/ops', msgId, { publisher: synth });
```


Touch can **extend** a message’s life but never **un-bound** it: the absolute 48 h ceiling from first publish still applies, so a touch (like a pull) can’t pin a message forever — whoever does the touching.

Throws `TouchError`: `TOUCH_INVALID_TOPIC`, `TOUCH_INVALID_MSGID`, `TOUCH_SIGN_FAILED` (no identity — touches must be signed).

`peer.unpub(topicName, opts?) → Promise<{ ok }>` Remove a topic’s whole message queue (kernel v2.10.0) — **owner-only**. The owner is the identity whose `nodeId` seeds the topic id; root axons verify ownership self-authenticatingly (the signer’s pubkey must bind to the owner `nodeId`, and that `nodeId` must derive the topic id).

Arg	Notes
<code>opts.destroy</code>	<code>false</code> (default) drops the messages, keeps the topic config; <code>true</code> is total removal (messages + config/ACL + ownership state).
<code>opts.publisher</code>	Owner selector; default = this peer.

```
await peer.unpub('news/world'); // clear the queue
await peer.unpub('news/world', { destroy: true }); // remove the topic entirely
```

Public (ownerless) topics have no owner to prove ownership and **cannot** be unpublished. **Throws** `UnpubError`: `UNPUB_INVALID_TOPIC`, `UNPUB_PUBLIC_TOPIC`, `UNPUB_SIGN_FAILED`.

Topic limits — queue size & hold time Every topic’s replay queue is bounded (kernel v2.10.0):

- **Max messages** — default **100**, max **256**. When full, the lowest-ordered message (by signed `seq`) is evicted — deterministic, so all replicas converge. A max of **1** is a *retained / latest-value* slot (each publish replaces the prior one); pair it with `peer.pull(null, ...)`.
- **Hold time** — default **24 h**, hard ceiling **48 h**. A message past its hold expires and is swept; it stops being delivered or pulled. A `pull` slides the hold forward, bounded by the ceiling.
- **Per-publisher quota (open topics only)** — on a public/anyone-may-publish topic, one publisher can occupy at most $\frac{1}{4}$ of the queue, so a single flooder can’t evict everyone else. Owner-gated topics aren’t quota’d.

In Phase A these are protocol defaults; owner-set per-topic values arrive with the topic-config object in a later phase.

4. Direct messaging

For 1:1 messages between specific peers, bypassing pub/sub.

`peer.send(peerIdHex, message) → Promise<reply>` Request/response. Resolves to whatever the recipient’s `onMessage` handler returns (or rejects with the remote handler’s thrown error).

```
const reply = await alice.send(bobIdHex, {
  kind: 'ping', body: 'hello',
});
// reply = whatever bob's handler returned
```

Throws `TypeError` if `peerIdHex` isn’t 66-char hex.

`peer.notify(peerIdHex, message) → Promise<boolean>` Fire-and-forget. Resolves once the frame is enqueued (NOT when delivered). Return value is `true` on successful enqueue, `false` if the transport rejected it.

```
await alice.notify(bobIdHex, { kind: 'typing', user: 'alice' });
```

`peer.onMessage(handler) → void` Register the single direct-message handler. Calling `onMessage` twice replaces the previous handler — for multi-kind dispatch, do it inside the handler:

```
peer.onMessage(async (senderIdHex, message) => {
  switch (message?.kind) {
    case 'ping':
      return { reply: 'pong' }; // send() resolves to this
    case 'typing':
      console.log(`${senderIdHex} typing`);
      return; // notify() discards
  }
});
```

`senderIdHex` is set by the transport at the receiving end. It's the authenticated sender (per the transport's binding); for higher assurance verify a signed envelope inside `message`.

5. Mesh introspection

`peer.peers() → string[]` Current synaptome membership as 66-char hex strings.

`peer.onPeerJoin(handler) → () => void` Fires when a peer is admitted to the synaptome. Handler receives (`peerIdHex`, `ctx`). Returns an unsubscribe function.

```
const unsub = peer.onPeerJoin((peerId, ctx) => {
  console.log('admitted', peerId, 'via', ctx?.addedBy);
});
// later: unsub();
```

`peer.onPeerLeave(handler) → () => void` Symmetric. Fires when a peer is evicted (TTL, churn, explicit unbind). Same signature.

`peer.lookup(targetKey) → Promise<lookupResult>` Iterative XOR-routed walk to find `targetKey` (typically a peer ID, but any 264-bit key works).

```
const r = await peer.lookup(BigInt('0x' + targetHex));
// {
//   found: true,
//   hops: 3,
//   time: 47, // ms; live RTT since v1.1.2
//   path: [<bigint>, <bigint>, ...], // the hop sequence
// }
```

`targetKey` may be a `BigInt` or a 66-char hex string — `topicToBigInt` converts internally.

`peer.health() → healthObj` A cheap, synchronous diagnostic snapshot — safe to poll on a UI tick or from a status endpoint:

```
peer.health();
// {
```

```
// nodeId:      '<66 hex>',
// synaptomeSize: 6,           // routing-table entries
// peers:       ['<66 hex>', ...], // same, as a list
// subscriptions: 2,
// axonRoles:   [{ topic, isRoot, children, cacheSize }],
// wireVersion: '1.0',
// started:     true,
// transport: {                // null on sim/node; populated on web
//   boundCount: 9,             // peers authenticated via axona/5
//   meshChannels: 8,           // WebRTC data channels (any state)
//   meshOpen: 8,               // open data channels
//   meshBound: 8,              // open AND axona/5-authenticated
//   bridgeState: 'open',
// },
// meshDegraded: false,        // true  channels open but NOT bound
// }
```

`meshDegraded` is the **routing-truth** signal: `true` means data channels are open but the `axona/5` handshake hasn't authenticated them, so no verified routing is flowing. A single `true` tick can be a normal mid-handshake transient; treat a value that stays `true` across several polls as the real signal. (`boundCount`/`meshBound` are the honest “usable peers” counts — distinct from the raw channel count.)

`peer.onLog(level, handler) → () => void` Subscribe to structured log events at a given level. **The level comes first**; returns an unsubscribe function.

```
const off = peer.onLog('warn', (msg, ctx) => console.warn(msg, ctx));
// level  'debug' | 'info' | 'warn' | 'error'
```

`peer.onError(handler) / peer.onUpgradeRequired(handler) → () => void` `onError` fires on background `AxonaError` emissions; `onUpgradeRequired` fires (and also fan-routes through `onError`) on a wire-version mismatch, carrying the `UpgradeRequiredError` with { `reason`, `serverVersion`, `clientVersion`, `downloadUrl` }.

```
peer.onError((err) => {}); // err = AxonaError
peer.onUpgradeRequired((err) => showUpgradeBanner(err.context.downloadUrl));
```

6. Topic IDs

Topic IDs share the same 264-bit address space as node IDs. The top 8 bits encode an S2 region; the bottom 256 bits are `sha256(canonical input)`.

`deriveTopicId(publisherId, topicName) → Promise<string>` Compute the 66-char hex topic ID.

<code>publisherId</code>	Mode	Top 8 bits	Hash input
<code>null</code> or <code>undefined</code> or <code>''</code>	Public	00 (global bucket)	<code>topicName</code>
<code>'<66 hex>'</code> (real <code>nodeId</code>)	Publisher-keyed	publisher's S2 prefix	<code>publisherId + ':' + topicName</code>
<code>'<s2 hex><64 zeros>'</code>	Region-keyed	region's S2 prefix	<code>synth + ':' + topicName</code>

The kernel **doesn't validate** that a 66-char synthetic ID corresponds to a real peer — it's just used as the hash prefix. Use this to bind topics to a region without tying them to a specific publisher.

```

import { deriveTopicId, geoCellId } from '@axona/protocol';

// Public mode
const pub = await deriveTopicId(null, 'world-news');
// '00' + sha256('world-news')

// Publisher-keyed
const pk = await deriveTopicId(myIdHex, 'posts');
// myS2 + sha256(myIdHex + ':posts')

// Region-keyed (recommended for most chat / forum / feed apps)
const s2 = geoCellId(lat, lng, 8);
const synth = s2.toString(16).padStart(2, '0') + '0'.repeat(64);
const rk = await deriveTopicId(synth, 'world-news');
// regionS2 + sha256(synth + ':world-news')

```

Throws `TypeError` if `topicName` is empty or `publisherId` is non-null and not 66-char hex.

Other pubsub/post.js exports

```

makePost({ topic, content, references?, identity?, sign? }) // → SignedPost
verifyPostHash(post) // → boolean
verifyTopicOwnership(post) // → boolean
verifySignature(post) // → Promise<boolean>
canonical(post) // → string (JSON-canonicalized)
sha256Hex(str) // → Promise<string>

```

These are lower-level helpers for protocol implementors building on top of the post-layer rather than `peer.pub`. Most apps don't need them.

7. Envelopes + verification

`buildEnvelope({ topic, message, ts?, seq?, identity, sign? }) → Promise<Envelope>` Construct an envelope explicitly. Useful when you want to sign a DM payload before sending it via `peer.send`, or to test verification flows.

```

const env = await buildEnvelope({
  topic: 'dm:to-bob',
  message: { text: 'hi' },
  seq: Date.now(), // per-publisher monotonic; 0 if omitted
  identity, sign: true,
});
// env = { msgId, seq, ts, topic, message, signerPubkey, signature }

```

`seq` (envelope format v2) is folded under the signature. `peer.pub` supplies a monotonic value automatically; you only pass it when building envelopes by hand. The signature covers a **domain-tagged** core (axona:pubsub-envelope:v2) so it can't be replayed as another kind of signature.

Throws `TypeError` if `identity` is missing `privateKey` or `pubkeyHex` when `sign: true`.

`verifyEnvelope(envelope) → Promise<{ ok, reason?, signed }>` Verify that `envelope.signature` is valid for the domain-tagged (`seq`, `ts`, `topic`, `message`) canonical inputs against `envelope.signerPubkey`, and that `msgId` matches. Requires `seq` (envelope format v2) — a pre-v2 envelope fails with `missing_seq`.

Returns a result object, not a boolean — check `.ok`:

```
const res = await verifyEnvelope(env);
if (!res.ok) console.warn('forged/invalid envelope', env.msgId, res.reason);
// res.signed === false for an unsigned envelope that is otherwise valid
// (res.ok === true). Your application decides whether to accept unsigned.
```

A common mistake: `if (!(await verifyEnvelope(env)))` is **always false** — the return value is an object (truthy). Always test `.ok`.

`reason` (when `!ok`) is one of `not_an_object` / `missing_*` (incl. `missing_seq`) / `unknown_signature_scheme` / `wrong_key_or_signature_length` / `bad_signature` / `bad_msgid`. As of kernel v2.7.0, root axons also verify the publisher signature at ingress before fan-out (finding B-4); as of v2.9.0 they additionally reject stale replays (freshness window + per-publisher `seq`, finding C-2) — but apps should still verify what they consume.

`verifyEnvelope` checks the signature and `msgId`; it does **not** check freshness (that's intentional — the legitimate replay-to-late-subscribers path serves cached history). Freshness is enforced only at live-publish ingress inside the kernel; the `checkFreshness(envelope, { now?, maxSkewMs? })` helper is exported if you need it.

`computeMsgId({ seq, topic, ts, message, signature? })` → `Promise<string>` Stand-alone `msgId` computation matching what `buildEnvelope` does (the canonical hash now binds `seq`). Useful for verifying a `msgId` server-side.

```
const msgId = await computeMsgId({
  seq: env.seq, topic: env.topic, ts: env.ts, message: env.message,
  signature: env.signature, // include for a signed envelope
});
console.log(msgId === env.msgId); // true
```

8. Errors

All thrown errors inherit from `AxonaError`:

```
class AxonaError extends Error {
  code: string; // stable UPPER_SNAKE identifier
  cause?: Error; // wrapped underlying error
  context?: object; // machine-readable details
}
```

Subclasses:

Class	Example codes
<code>IdentityError</code>	<code>IDENTITY_INVALID</code> , <code>IDENTITY_LOAD_FAILED</code> , <code>IDENTITY_KEY_IMPORT_FAILED</code>
<code>TransportError</code>	<code>TRANSPORT_NOT_STARTED</code> , <code>TRANSPORT_TIMEOUT</code> , <code>TRANSPORT_NOT_BOUND</code> , <code>TRANSPORT_CLOSED</code>
<code>PublishError</code>	<code>PUBLISH_INVALID_TOPIC</code> , <code>PUBLISH_SIGN_FAILED</code> , <code>PUBLISH_PAYLOAD_TOO_LARGE</code> , <code>PUBLISH_INVALID_MESSAGE</code>
<code>SubscribeError</code>	<code>SUBSCRIBE_INVALID_TOPIC</code> , <code>SUBSCRIBE_HANDLER_MISSING</code>
<code>KillError</code>	<code>KILL_INVALID_TOPIC</code> , <code>KILL_INVALID_MSGID</code> , <code>KILL_SIGN_FAILED</code>
<code>UnpubError</code>	<code>UNPUB_INVALID_TOPIC</code> , <code>UNPUB_PUBLIC_TOPIC</code> , <code>UNPUB_SIGN_FAILED</code>

Class	Example codes
PullError	PULL_INVALID_MSGID, PULL_AXONS_UNREACHABLE
MetricsError	METRICS_AXONS_UNREACHABLE
UpgradeRequiredError	UPGRADE_REQUIRED (bridge close code 4426)

The full code list is in `ErrorCodes`:

```
import { ErrorCodes } from '@axona/protocol';
console.log(ErrorCodes.PUBLISH_SIGN_FAILED); // 'PUBLISH_SIGN_FAILED'
```

Switch on `.code`, not `.message`:

```
try { await peer.pub(topic, msg); }
catch (err) {
  if (err.code === ErrorCodes.PUBLISH_SIGN_FAILED) {
    await reload();
  } else { throw err; }
}
```

Wire-format helpers

```
isWireError(obj)      // + boolean - does this look like an over-the-wire AxonaError?
fromWire(obj)         // + AxonaError - reconstruct
err.toWire()          // + { __axonaError: true, class, code, message, context }
```

Errors thrown inside remote handlers (i.e., inside another peer's `onMessage` reached via `peer.send`) survive serialization with their class and code intact.

9. Persistence

PersistenceAdapter (interface)

```
class PersistenceAdapter {
  async load(key)      { /* returns previously-saved string or null */ }
  async save(key, value) { /* persists; value is string */ }
  async clear()        { /* wipe everything; idempotent */ }
}
```

Implement this to plug in your own storage (Redis, S3, postgres, etc.). Two built-in implementations:

Both ship as **classes** behind sub-path imports (so neither pulls `globalThis.indexedDB` into Node nor `node:fs` into browser bundles), and you pass an instance as the constructor's **`persist`** option.

`IndexedDBPersistence({ dbName })` Browser. Uses a single object store inside the named DB.

```
import { IndexedDBPersistence } from '@axona/protocol/persistence/indexeddb.js';

const peer = new AxonaPeer({
  domain, node, identity,
  persist: new IndexedDBPersistence({ dbName: 'my-app' }),
});
```

`FilePersistence({ dir })` Node. Writes JSON files under `dir` (default `./axona`).

```
import { FilePersistence } from '@axona/protocol/persistence/file.js';

const peer = new AxonaPeer({
  domain, node, identity,
  persist: new FilePersistence({ dir: './state' }),
});
```

Manual snapshots

```
const state = await peer.snapshot();
// state = { identity?, subscriptions, synaptome, lookups, ... }
// - any JSON-safe object

// restore via the static factory (fromSnapshot is static, not an instance method):
const peer = await AxonaPeer.fromSnapshot(state, { engine, node, transport });
```

Transport + protocol surface

10. Contracts

Abstract base classes that transports / DHT implementations extend. Application developers rarely use these directly — they're for authoring custom transports or alternate DHT engines.

Transport (abstract class)

```
import { Transport } from '@axona/protocol';

class MyTransport extends Transport {
  async start(localNodeId) { /* ... */ }
  async stop() { /* ... */ }
  async send(peerId, type, body) { /* ... */ } // request/response
  async notify(peerId, type, body) { /* ... */ } // fire-and-forget
  async openConnection(peerId) { /* ... */ }
  async closeConnection(peerId) { /* ... */ }
  isConnected(peerId) { /* ... */ }
  getLatency(peerId) { /* RTT ms or -1 */ }
  onRequest(type, handler) { /* ... */ }
  onNotification(type, handler) { /* ... */ }
  onPeerDied(handler) { /* ... */ }
}
```

The full contract surface is in `@axona/protocol/contracts/Transport.js` (sub-path import).

DHT (abstract class) Used by `AxonaManager`'s dht adapter. Defines `getSelfId`, `findKClosest`, `routeMessage`, `sendDirect`, `onRoutedMessage`, `onDirectMessage`. See `axona-peer/src/browser_engine.js` for a concrete adapter that wraps `AxonaPeer`'s methods.

BootstrapService (abstract class) For implementing peer discovery (signaling brokers, seed lists, DNS-SD, etc). Rarely needed by app code.

11. Sim transport

In-process router for testing + simulators.

```
new SimNetwork({ latencyFn? })

import { SimNetwork } from '@axona/protocol';

// All edges = 0ms
const net = new SimNetwork();

// 50ms each way
const net = new SimNetwork({ latencyFn: () => 50 });

// Custom per-pair
const net = new SimNetwork({
  latencyFn: (fromId, toId) => Math.abs(hash(fromId) - hash(toId)) % 100,
});
```

simTransport({ network, identity, heartbeatMs?, ... }) → **SimTransport** Factory that builds + wires a **SimTransport** connected to a **SimNetwork**. Same constructor options as **new SimTransport**({...}) but with a more concise call site.

```
import { SimNetwork, simTransport } from '@axona/protocol';

const network = new SimNetwork();
const t = simTransport({ network, identity, heartbeatMs: 0 });
await t.start(identity.id);
await t.openConnection(otherIdentity.id);
```

heartbeatMs: 0 disables heartbeats (recommended for tests; the default heartbeat is helpful for production where peers may go away silently, useless for in-process synchronous flows).

11b. Web transport (browser)

The production browser transport: a **WebSocket** to the bridge for bootstrap + signaling, and a **WebRTC** mesh for direct peer-to-peer traffic. It owns the bridge socket lifecycle (reconnect/backoff, ping/pong, stale detection), the version-gate handshake, and the **axona/5** authenticated-identity handshake on both the bridge link and every mesh link.

webTransport(opts) → **CompositeTransport**

```
import { webTransport } from '@axona/protocol/transport/web/index.js';

const transport = webTransport({
  bridgeUrl: 'wss://bridge.axona.net', // required
  identity, // from deriveIdentity() - signs the handshake
  peerVersion: '3.11.0', // your app version (gated by the bridge)
  reconnect: true, // auto-reconnect with backoff (default)
  // log, autoHandshake, handshakeTimeoutMs, pingIntervalMs,
  // reconnectInitialMs, reconnectMaxMs, WebSocketImpl ...
});
await transport.start(); // resolves after the bridge handshake
```


Beyond the Transport contract (send/notify/onRequest/ onNotification/boundPeers/onPeerBound/start/stop), it exposes an observability surface for the UI:

```
transport.bridgeState           // 'connecting'|'open'|'stale'|'disconnected'|'upgrade-required'
transport.bridgeInfo            // { connId, version, kernelVersion, turn } | null
transport.bridgeRtt             // last ping-pong ms (null until first pong)
transport.bridgeRttAvg          // mean of the recent RTT window
transport.bridgeNodeId          // bridge's 66-hex nodeId (null pre-handshake)
transport.onBridgeState((state, detail) => {}); // + unsub
transport.onWelcome((info) => {});           // + unsub (replays last)
transport.onPingTraffic((dir) => {});        // 'sent' | 'recv'
transport.reconnectNow();                // force an immediate reconnect (tab resume)
transport.boundPeers();                 // bigint[] - axona/5-authenticated peers
```

Channel binding (finding A-1): each mesh link’s axona/5 proof is bound to that connection’s DTLS certificate fingerprint, so a fingerprint-rewriting bridge cannot transparently MITM “direct” peer traffic. The bridge is trusted for signaling only, never for the contents of peer-to-peer links. See §17.

12. Handshake

Wire-version handshake helpers used by transports on a fresh signaling channel. Most app code never touches these — they’re plumbed into `simTransport/webTransport/etc.` internally.

```
buildClientHello({ version, nodeId }) // + frame for client + server
buildServerHello({ version, minVersion }) // + frame for server + client
parseHello(frame) // + { version, nodeId? }
parseVersion(str) // + { major, minor, patch }
compareVersions(a, b) // + -1 | 0 | 1
wireCompatible(peerVersion, minRequired) // + boolean
performClientHandshake(channel, opts) // + handshake result
performServerHandshake(channel, opts)
```

The constants `WIRE_VERSION`, `KERNEL_VERSION`, `UPGRADE_CLOSE_CODE` are exported from the same module — see §16.

Low-level utilities

13. Ed25519 helpers

Web Crypto Ed25519 wrappers. Useful when you want to sign / verify outside of `buildEnvelope` — e.g. for signing direct-message payloads manually.

```
const { pubkey, privateKey } = await generateKeyPair();
// pubkey = Uint8Array(32); privateKey = CryptoKey

const pubHex = await exportPublicKey(pubkey);
const pubKey = await importPublicKey(pubHex);

const sig = await sign(privateKey, dataBytes); // Uint8Array(64)
const ok = await verify(pubKey, sig, dataBytes);

const signer = makeSigner(privateKey); // (bytes) => Promise<sig>
const verifier = makeVerifier(pubKey); // (sig, bytes) => Promise<boolean>
```

Implementation-agnostic: substitute @noble/ed25519 for runtimes that don't have Web Crypto Ed25519 (Chrome <110, Safari <17, Firefox <130, Node <20). The `makeSigner` / `makeVerifier` contract is what the rest of the kernel binds against.

14. Hex / ID math

```
ID_BITS           // 264
HASH_BITS         // 256
S2_BITS           // 8
HEX_CHARS         // 66
MAX_ID            // BigInt (2**264 - 1)
MAX_HASH          // BigInt (2**256 - 1)
MAX_S2            // 255

toHex(bigint)      // → 66-char lowercase hex
fromHex(hex)       // → BigInt
isHexId(s)         // → boolean (66-char lowercase hex)
assembleId(s2Prefix, hash) // → BigInt - concat 8-bit prefix + 256-bit hash
extractS2Prefix(idBig) // → number 0..255
extractHash(idBig) // → BigInt (bottom 256 bits)
s2PrefixOfHex(hex) // → number 0..255

xorDistance(a, b)  // → BigInt
stratumOf(a, b)    // → 0..264 - count of leading zero bits
clz264(bigint)     // → 0..264 - count of leading zero bits in a 264-bit value
randomU256()       // → BigInt - cryptographically random
```

These are useful when implementing custom routing logic, K-closest overrides, or analytics on top of the address space.

15. Geo utilities

```
geoCellId(lat, lng, bits) // → integer - S2 cell ID at the given bit depth
geoCellSubCenters(code)  // → [center0, center1] - the cell's two level-3 sub-cell centers
geoCellHalf(lat, lng)    // → 0 | 1 - which half (sub-cell) a point falls in
haversine(lat1, lng1, lat2, lng2) // → km between two points
roundTripLatency(lat1, lng1, lat2, lng2) // → ms - rough fiber estimate
randomU32()              // → 0..2**32 - 1
```

Region names. Each of the 192 region codes carries exactly **one** human-readable name, so a region always presents the same label (no location-dependent flip-flop). Where a cell straddles an ocean and a landmass the land name wins; a multi-country cell takes its dominant city; homogeneous cells (a single country, open ocean) keep their name. A name is usually unique to one code, but an area larger than one cell may span adjacent codes (e.g. `centrlam`) — `regionCode` then returns the canonical (lowest) code.

```
regionName(code) // → string | null           e.g. regionName(0x89) → 'useast'
regionCode(name) // → integer code | null      (inverse; canonical code for a multi-cell name)
resolveRegion(nameOrCode) // → integer code | null accepts a name OR a numeric/hex code, normalized
regionNameForLatLng(lat, lng) // → string      region name for a coordinate
REGION_NAMES     // → frozen string[] indexed by code [0,192)
regionNames(code) // → [name] | null          deprecated one-element back-compat shim
```

Names match `/^[a-z0-9]{1,8}$/`; open-ocean cells are `<ocean3>_<hex>` (`pac_68`, `atl_0a`, ...). The interactive `examples/s2-region-visualizer/` renders all 192 — one name + the code each.

Sub-path import for the remaining geo helpers (@axona/protocol/utils/geo.js): `getRegion`, `getContinent`, several XOR routing-table builders. Rarely needed by app code.

16. Constants

```
WIRE_VERSION      // '2.0'          - wire format major.minor
KERNEL_VERSION    // '2.40.0'       - kernel semver
AUTH_PROTO        // 'axona/5'      - authenticated-identity handshake tag
UPGRADE_CLOSE_CODE // 4426          - WebSocket close code for version mismatch
ID_BITS           // 264
HEX_CHARS         // 66
```

WIRE_VERSION is what bridges enforce gates against; KERNEL_VERSION is informational and corresponds to the npm release tag.

17. Security model (what the protocol guarantees)

The full, versioned record is the security changelog; this is the developer-facing summary of what's enforced as of kernel v2.16.0. Everything here is **self-authenticating** — no certificate authority, no central trust server, no reputation service.

- **Authenticated identity (axona/5).** A nodeId's bottom 256 bits are SHA-256(pubkey), and every connection runs a handshake proving the peer holds that key (BIND: pubkey hashes to the id · POSSESS: Ed25519 signature · CHANNEL: the signature is bound to this live connection so a captured proof can't be replayed onto another link). A peer cannot claim an id it doesn't own.
- **Channel binding (untrusted bridge).** Mesh links fold each side's DTLS certificate fingerprint into the signed handshake, so a bridge that terminates DTLS to interpose on a "direct" link produces divergent bindings and the handshake fails. The bridge is trusted for *signaling only*.
- **Pub/sub trust boundary.** A peer can only subscribe *itself* (every path checks the authenticated sender, not a payload-named id) — no reflection/amplification at a victim. A node only hosts topics it's among the K-closest to — no memory-exhaustion via random-topic floods. Publisher signatures are verified at root-axon ingress before fan-out.
- **Verified routing admission (eclipse resistance).** Routing-table entries are *earned*: a peer named in routing gossip is a candidate, admitted only after this node connects and the peer authenticates via axona/5. Forged gossip cannot inject hops. (Observe this via `health().meshDegraded` and `boundCount`.)
- **Key hygiene.** Browser identities can hold a **non-extractable** signing key (`deriveIdentity({ extractable: false })`); `loadIdentity` verifies the private key matches its public key on load.

What the protocol does **not** do: encrypt your application payload (it's opaque bytes — encrypt inside `message` if you need confidentiality), or vouch for what an authenticated peer *says* beyond what cryptography proves. Open security work is tracked privately; the public changelog lists only resolved items.

A word on stability

@axona/protocol v2.x is stable for the application surface documented in §§1–9. The transport + protocol surface (§§10–12, incl. the browser `webTransport` and the axona/5 security model in §17) is also stable but expect occasional additions (new transport factories, new bootstrap services). Low-level utilities (§§13–16) may grow but won't change incompatibly.

The `AxonaPeer` constructor still accepts a legacy `engine` argument for backward compatibility with pre-Phase-5 callers; new code should pass `domain` instead.

For changelog and migration notes between minor versions, see <https://github.com/axona-net/axona-protocol/releases>.